

# **Lecture 20:**

## **Drawing Fast: The Graphics Pipeline**

### **Explaining the Pipeline**

# Review: The “abstractions” of 3D Drawing

Triangles as “the” primitive

- each triangle is independent

Vertices are independent (except they make triangles)

- information at vertices, not triangles

Triangles cover pixels

- pixels are independent (each with a coordinate)

- local lighting/shading (texture lookups, light references)

Z-Buffer sorts it all out (order independence)

# What happens when we draw

We draw **triangles**

They start out as positions (local coordinates)

They end up as colored pixels

Today... The triangle's Journey

# Why do you care how the hardware works?

Because it is cool!

# Why do you care how the hardware works?

The limitations of the hardware explain why we do things.

Why prefer texture hacks to more principled lighting?

It will help us understand how to make things fast.

Next week, we will program the hardware directly.

It will not make sense unless without understanding

# The steps of 3D graphics

Model objects (make triangles)

Transform (find point positions)

Shade (lighting– per tri / vertex)

Transform (projection)

Rasterize (figure out pixels)

Shade (per-pixel coloring)

Write pixels (with Z-Buffer test)

1. Put a 3D primitive in the World  
**Modeling (model transformations)**
2. Figure out what color it should be  
**Shading**
3. Position relative to the Eye  
**Viewing** / Camera Transformation
4. Get rid of stuff behind you/offscreen  
**Clipping**
5. Figure out where it goes on screen  
**Projection** (sometimes called Viewing)
6. Figure out if something else blocks it  
**Visibility** / Occlusion
7. Draw the 2D primitive  
**Rasterization** (convert to Pixels)

**Practical Caveats**

**Colors might be determined later**

**Often we combine transformations**

**Z-Buffer reverses these two steps**

# A Pipeline

Triangles are independent

Vertices are independent

Pixels (within triangles) are independent

(caveats about sharing for efficiency)

Don't need to finish 1 before start 2

(might want to preserve finishing order)



# Is it really a pipeline?

In theory, yes – good way to think about

Historically – it is how the hardware works


In practice with modern hardware?  
who knows

# Pipelining in conventional processors

Start step 2 before step 1 completes

Unless step 2 depends on step 1

Pipe Stall




# Triangles are independent!

## No stalls! (no complexity of handling stalls)


Start step 2 before step 1 completes

$C = A * B$   
 $F = D * E$   
 $J = G * H$




Unless step 2 depends on step 1  
Pipe Stall

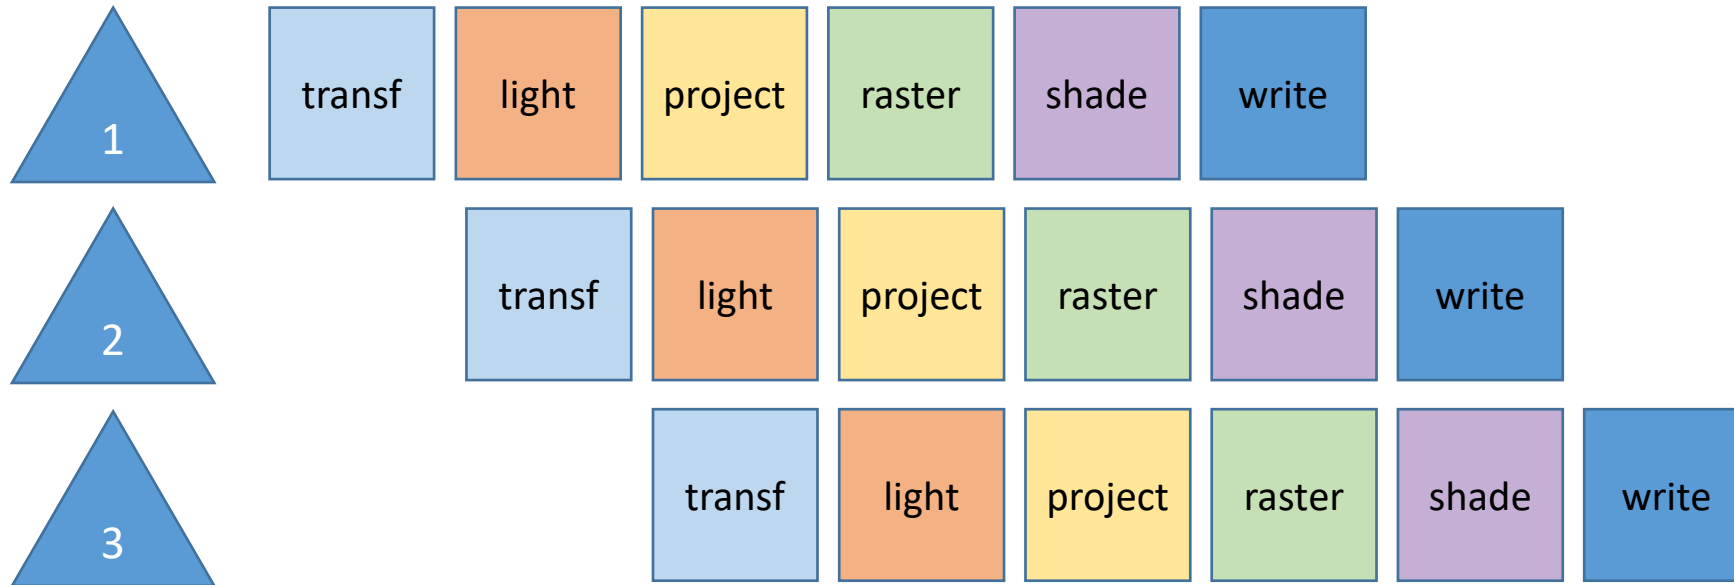
$C = A * B$   
 $F = D * C$   
 $J = G * H$



# A Pipeline




# A Pipeline



# Vertices are independent

## Parallelize!



# Parallelization

Vertex operations

- split triangles / re-assemble

- compute **per-vertex** not per-triangle

Pixel (fragment) operations

- lots of potential parallelism

- less predictable

Use queues and caches

# Why do we care?

This is why the hardware can be fast

It requires a specific model

Hardware implements this model

The programming interface is designed for this model.  
You need to understand it.



# Some History...

Custom Hardware (pre-1980)

rare, each different

Workstation Hardware (early 80s-early 90s)

increasing features, common methods

Consumer Graphics Hardware (mid 90s-)

cheap, eventually feature complete

Programmable Graphics Hardware (2002-)

# Graphics Workstations 1982-199X

Implemented graphics in hardware

Providing a common abstraction set

Fixed function –  
it was built into the hardware

# Silicon Graphics (SGI)

Stanford Research Project 1980

Spun-off to SGI (company) 1982

The Geometry Engine

- 4x4 matrix multiply chip

- approximate division

Raster engine (Z-buffer)

# 1988: The Personal Iris



The 4D-2X0 series

4 processors (240)

Different graphics

1988 - GT/GTX

1990 - VGX



# Why do you care?

It's the first time the abstractions were right  
later stuff adds to it

It's where the programming model is from  
it was IrisGL before OpenGL (before WebGL)

It's the pipeline at its essence  
we'll add to it, not take away

# The Abstractions

Points / Lines / Triangles

Vertices in 4D

Color in 4D (RGBA = transparency)

Per-Vertex transform ( $4 \times 4$  + divide by  $w$ )

Per-Vertex lighting

Color interpolation

Fill Triangle

Z-test (and other tests)

Double buffer (and other buffers)



# What's left to add?

All of this was in hardware in the 80s

1990 – texture

1992 – multi-texture (don't really need)

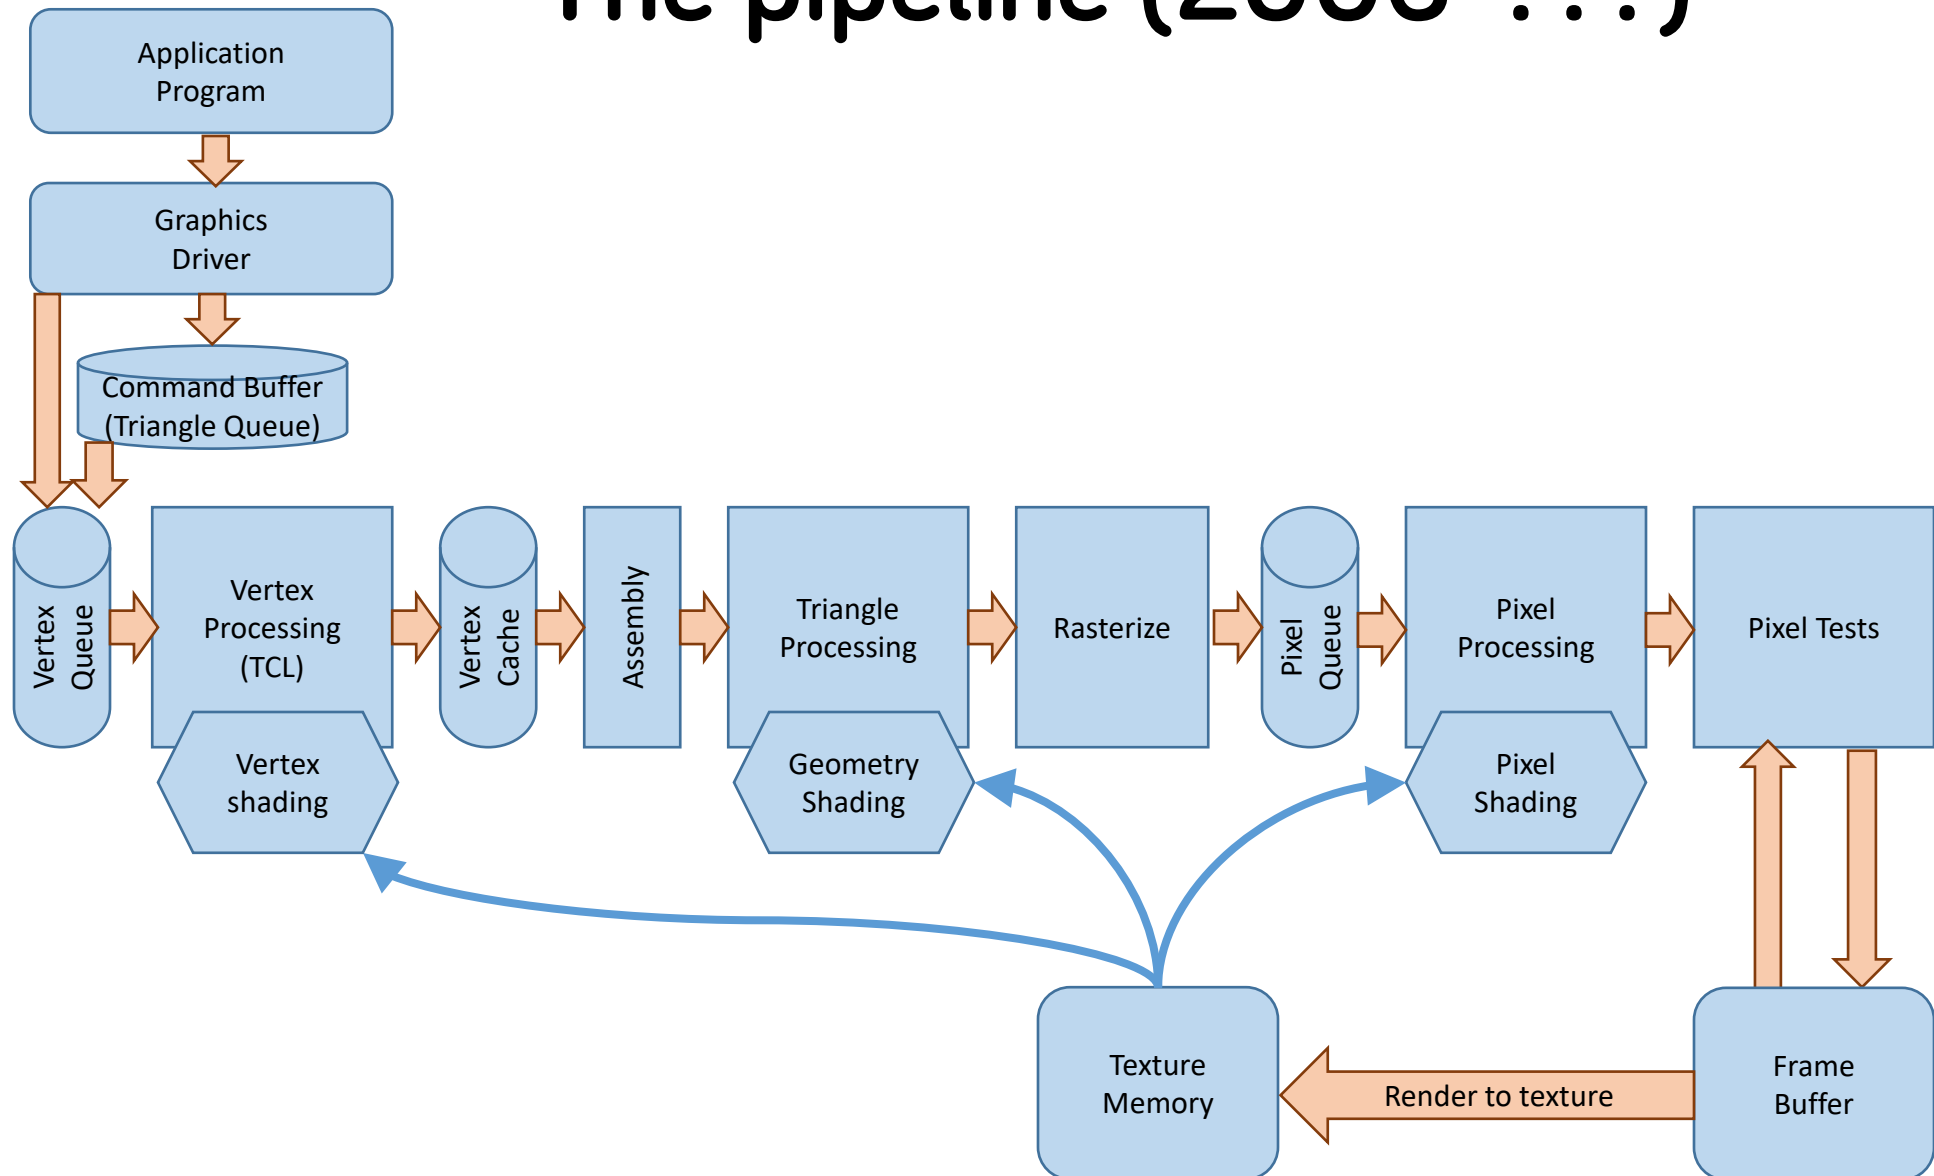
1998 (2000) – programmable shading

2002 – programmable pipelines

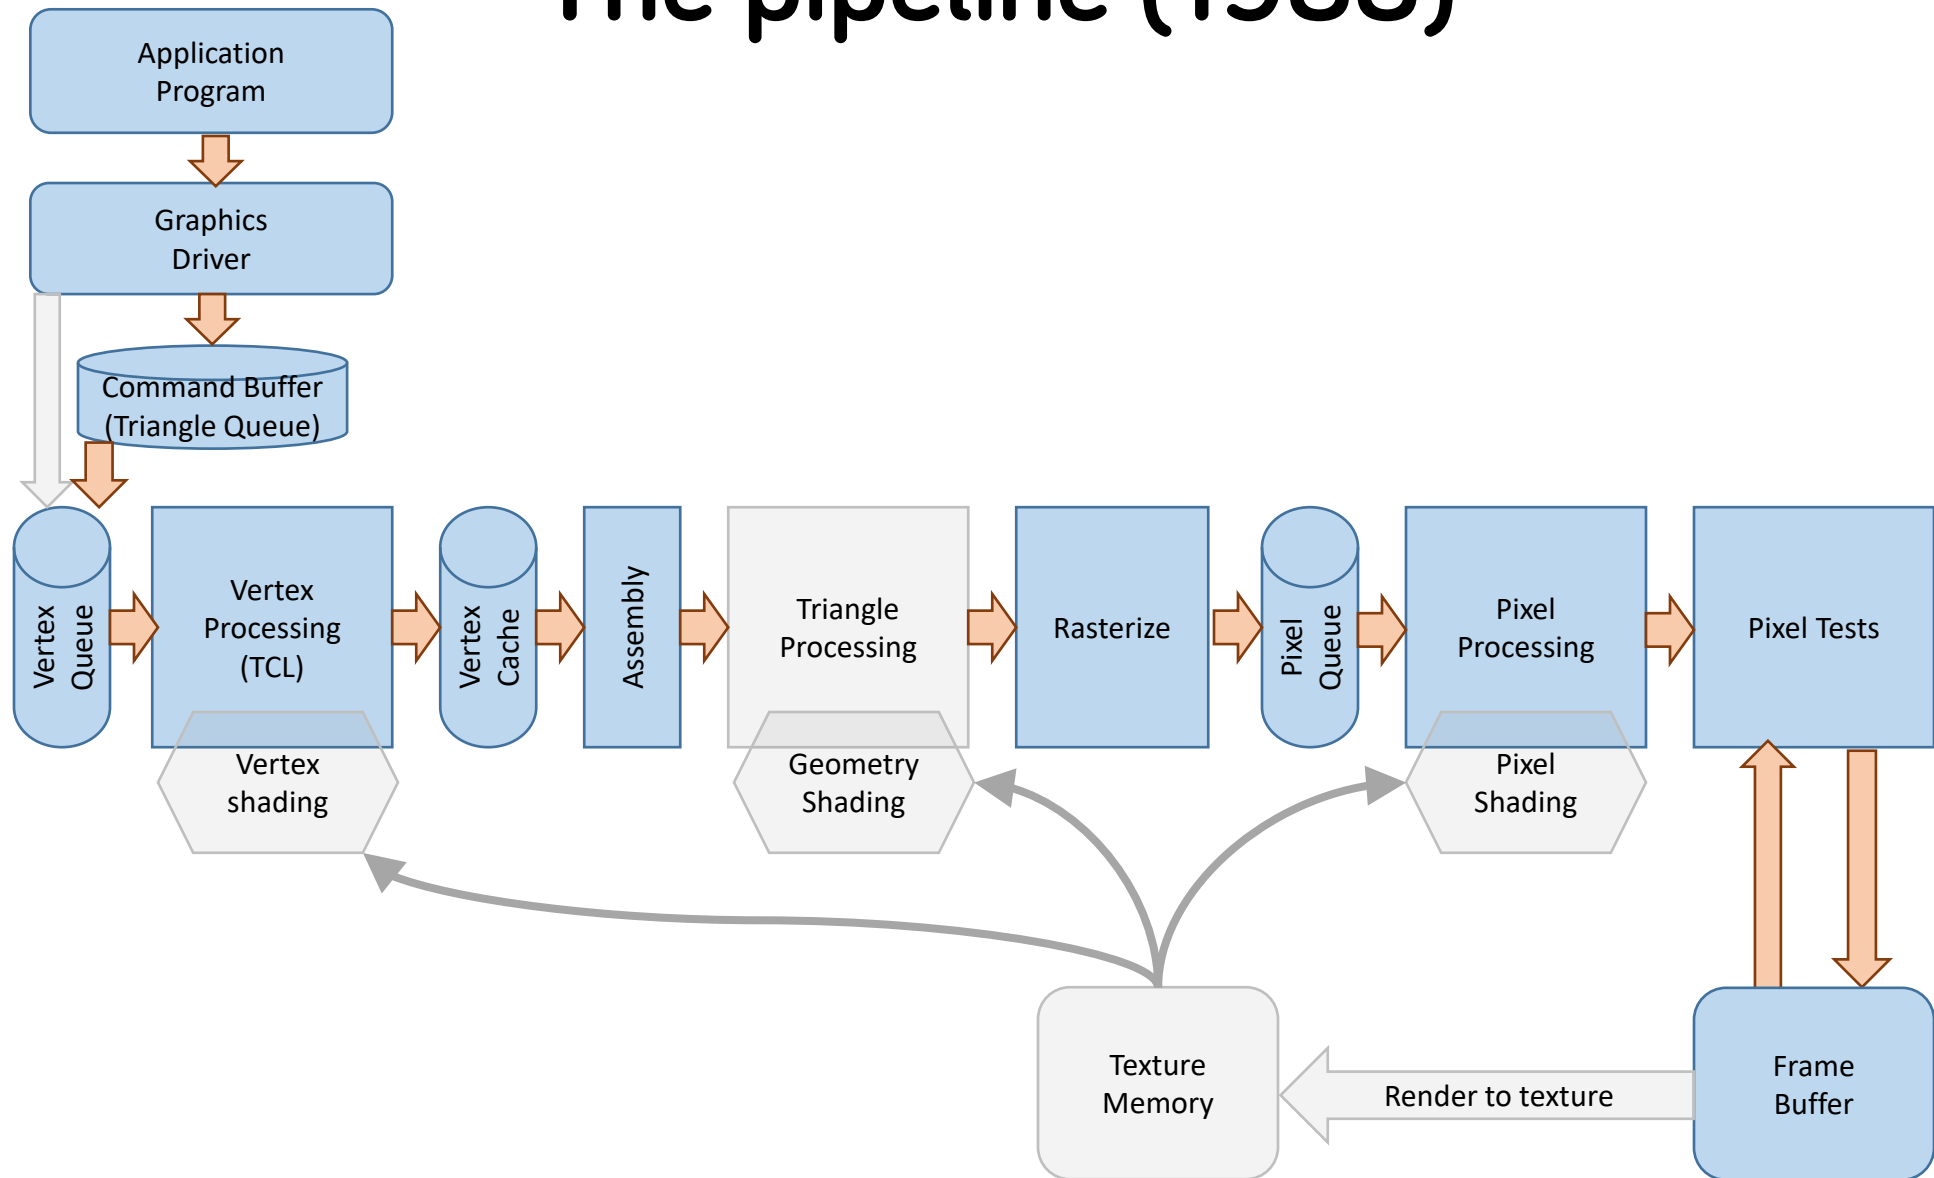
2005 – more programmability



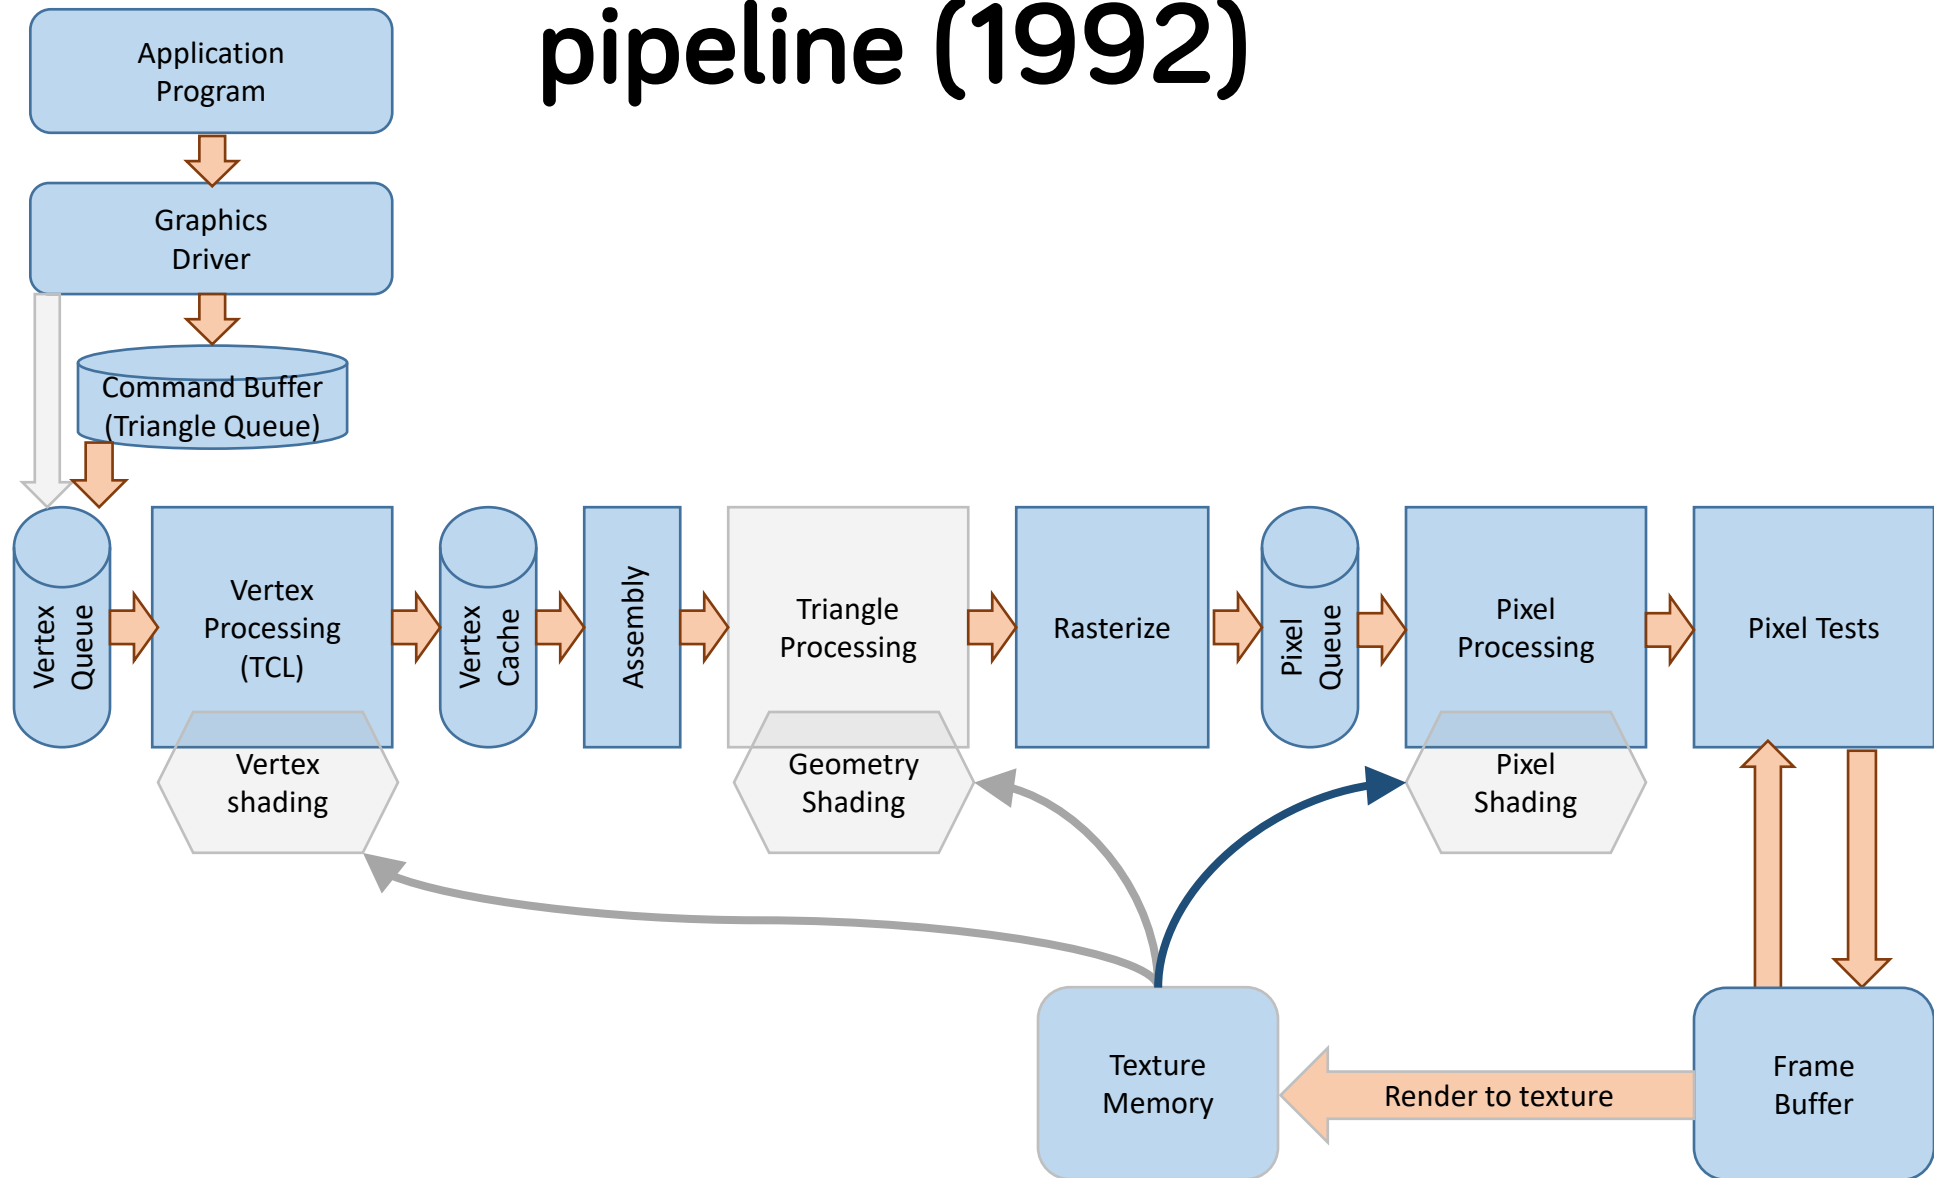
# The pipeline (2006-???)



# The pipeline (1988)



# The full fixed-function pipeline (1992)



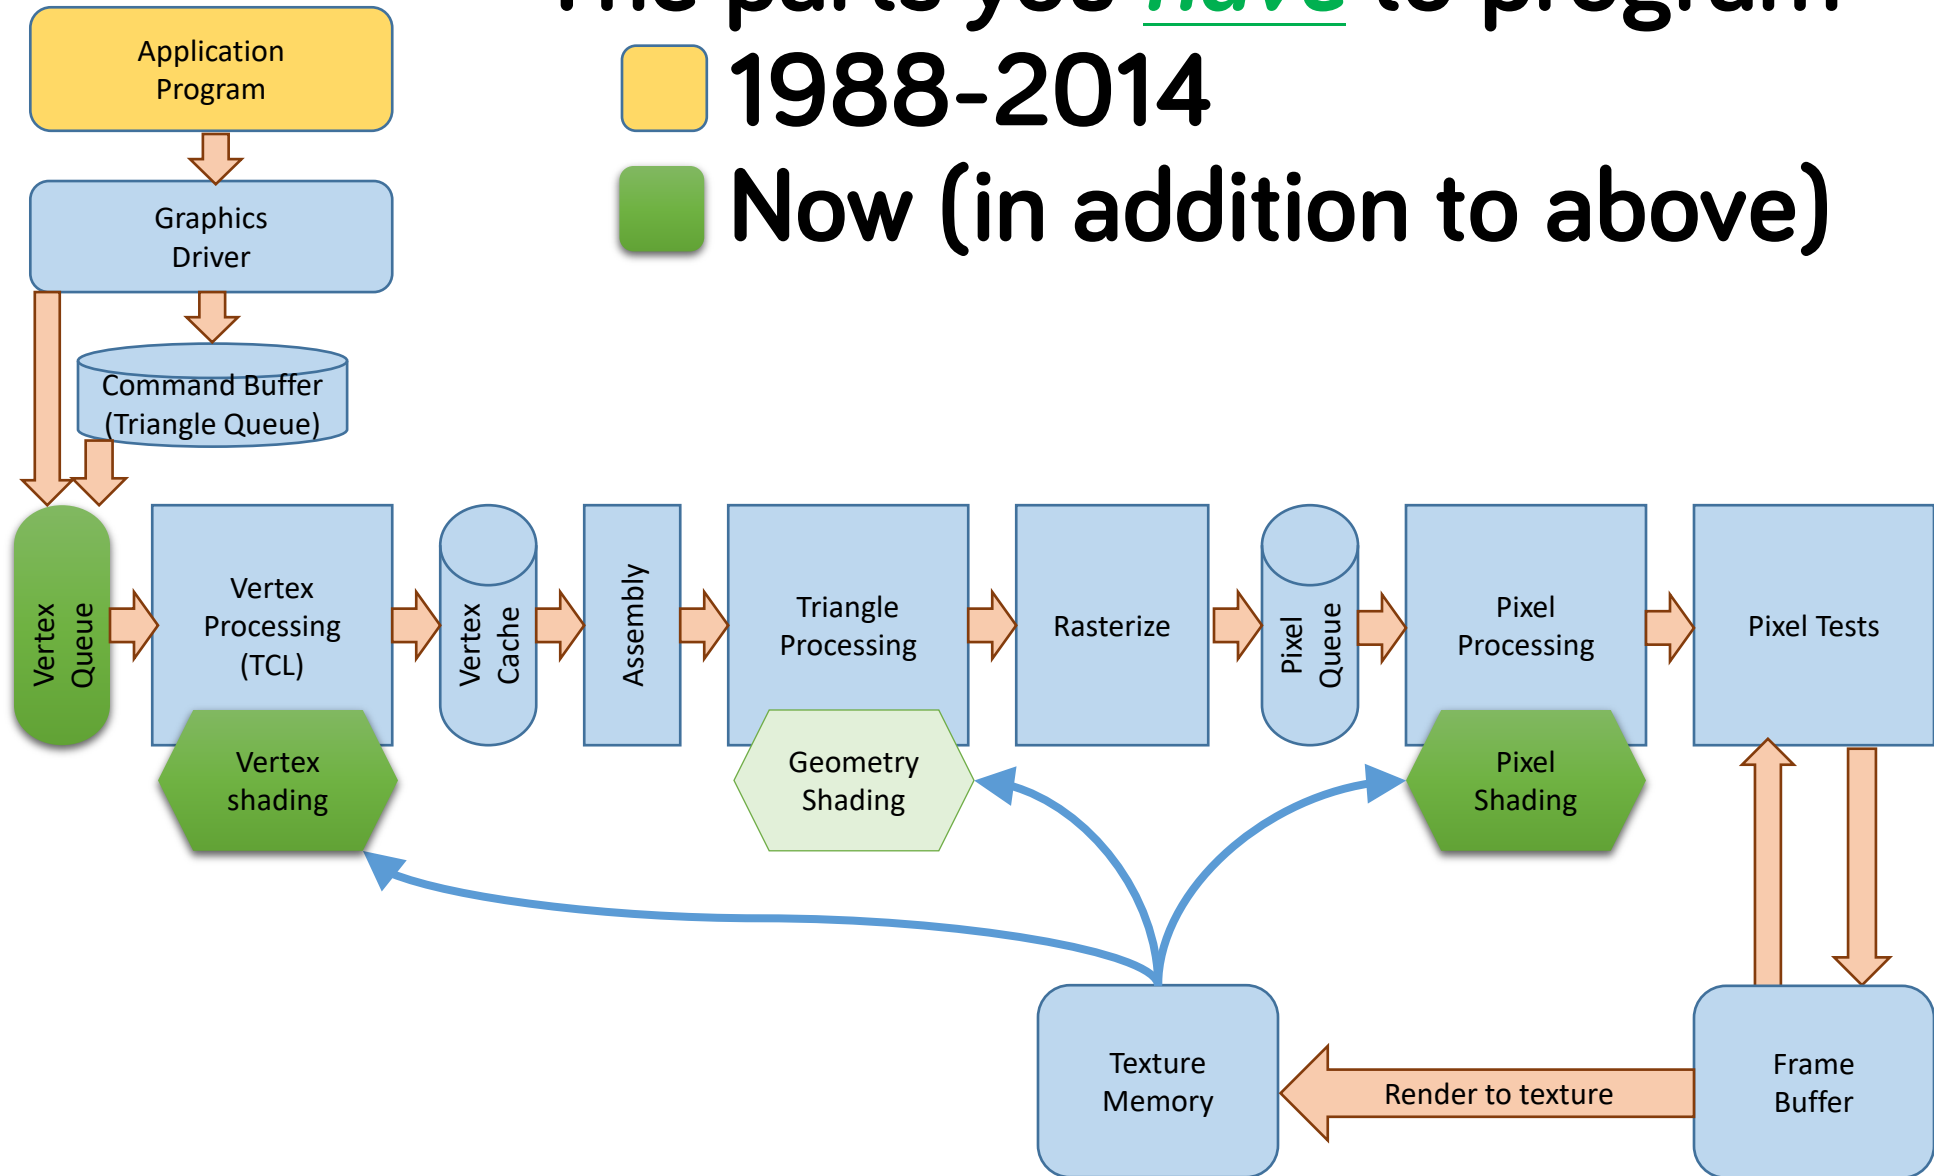
The parts you have to program



1988-2014



Now (in addition to above)



# A Triangle's Journey

# Things to observe as we travel through the pipeline...

What does each stage do?

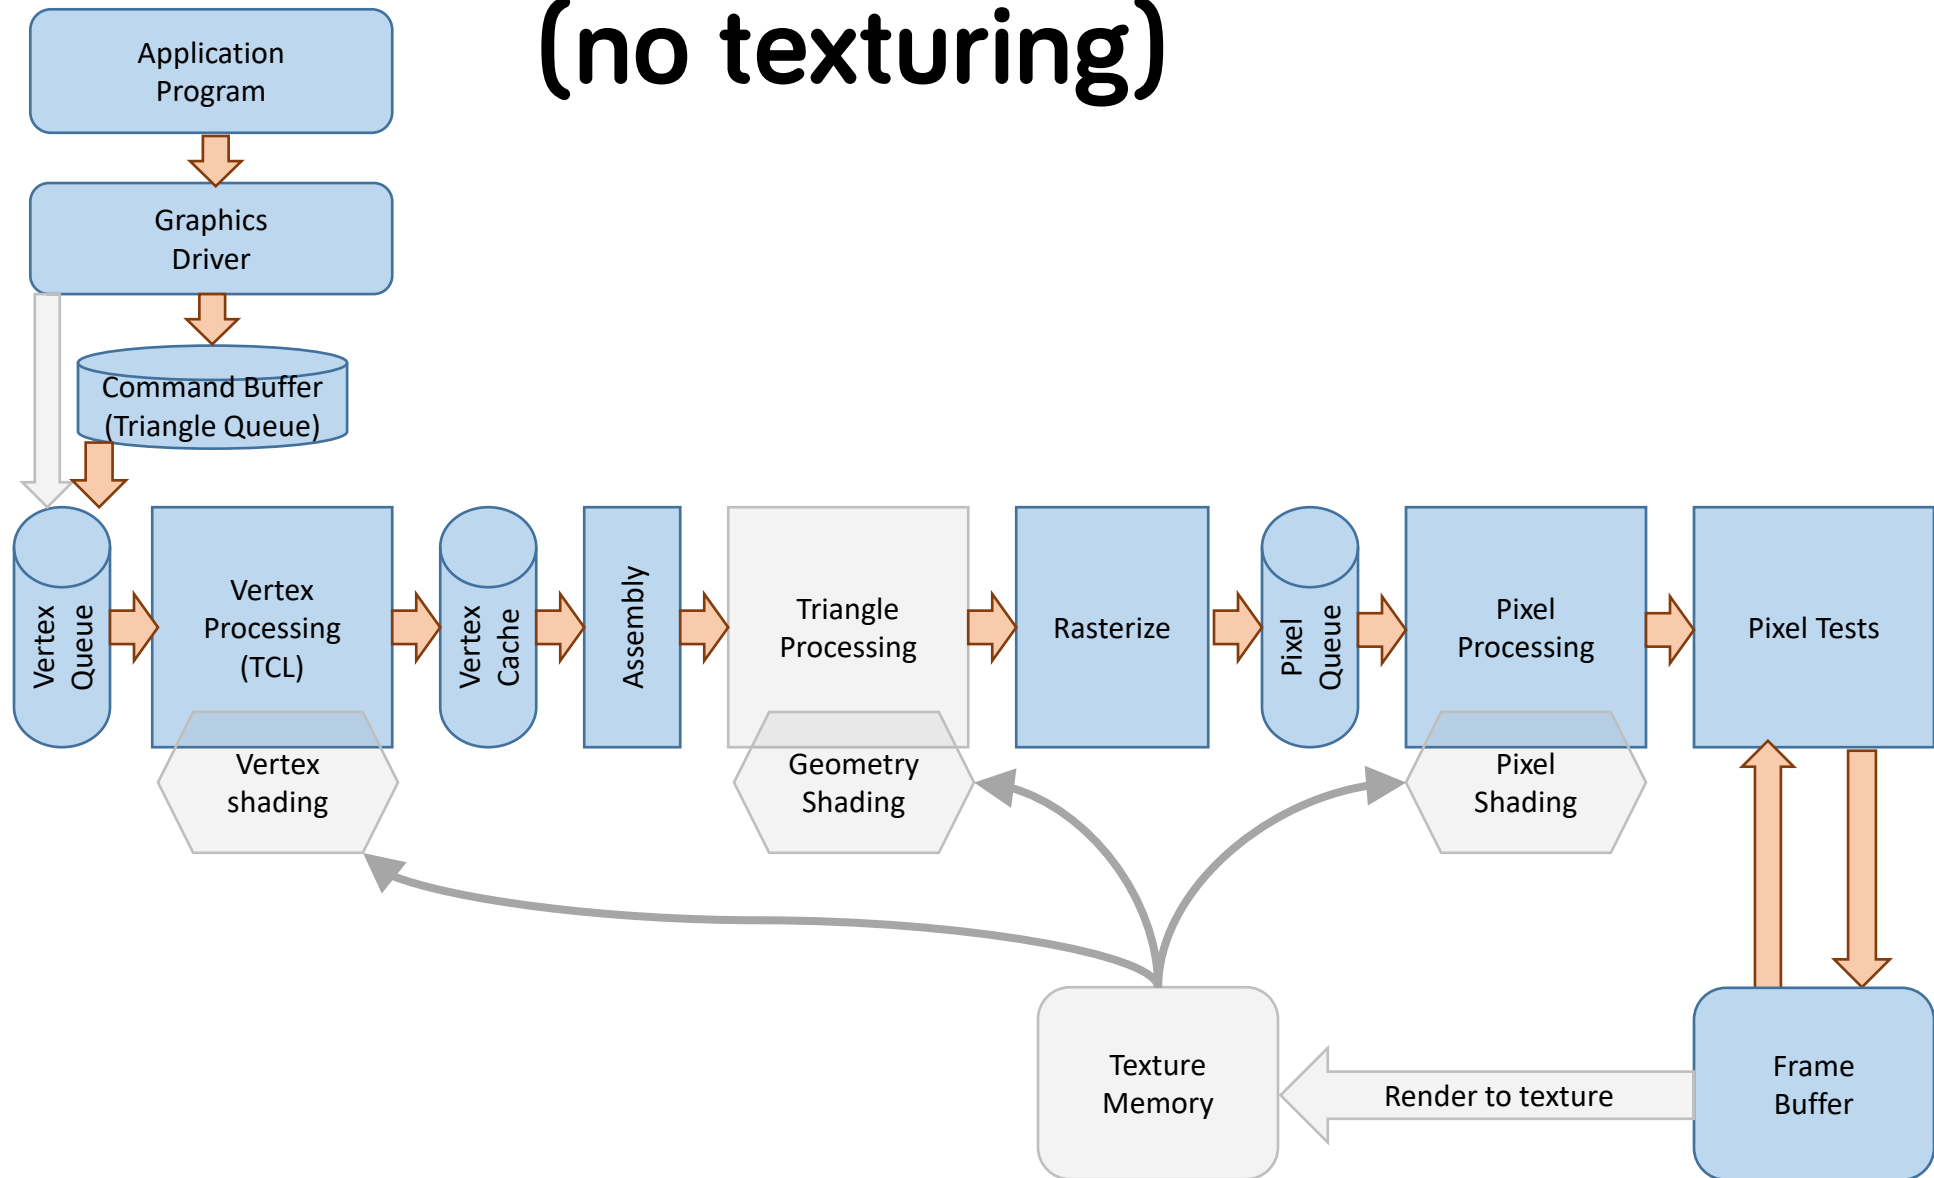
What are its inputs and output?

important for programmability

Why would it be a bottleneck?

and what could we do to avoid it

# The pipeline (1988) (no texturing)



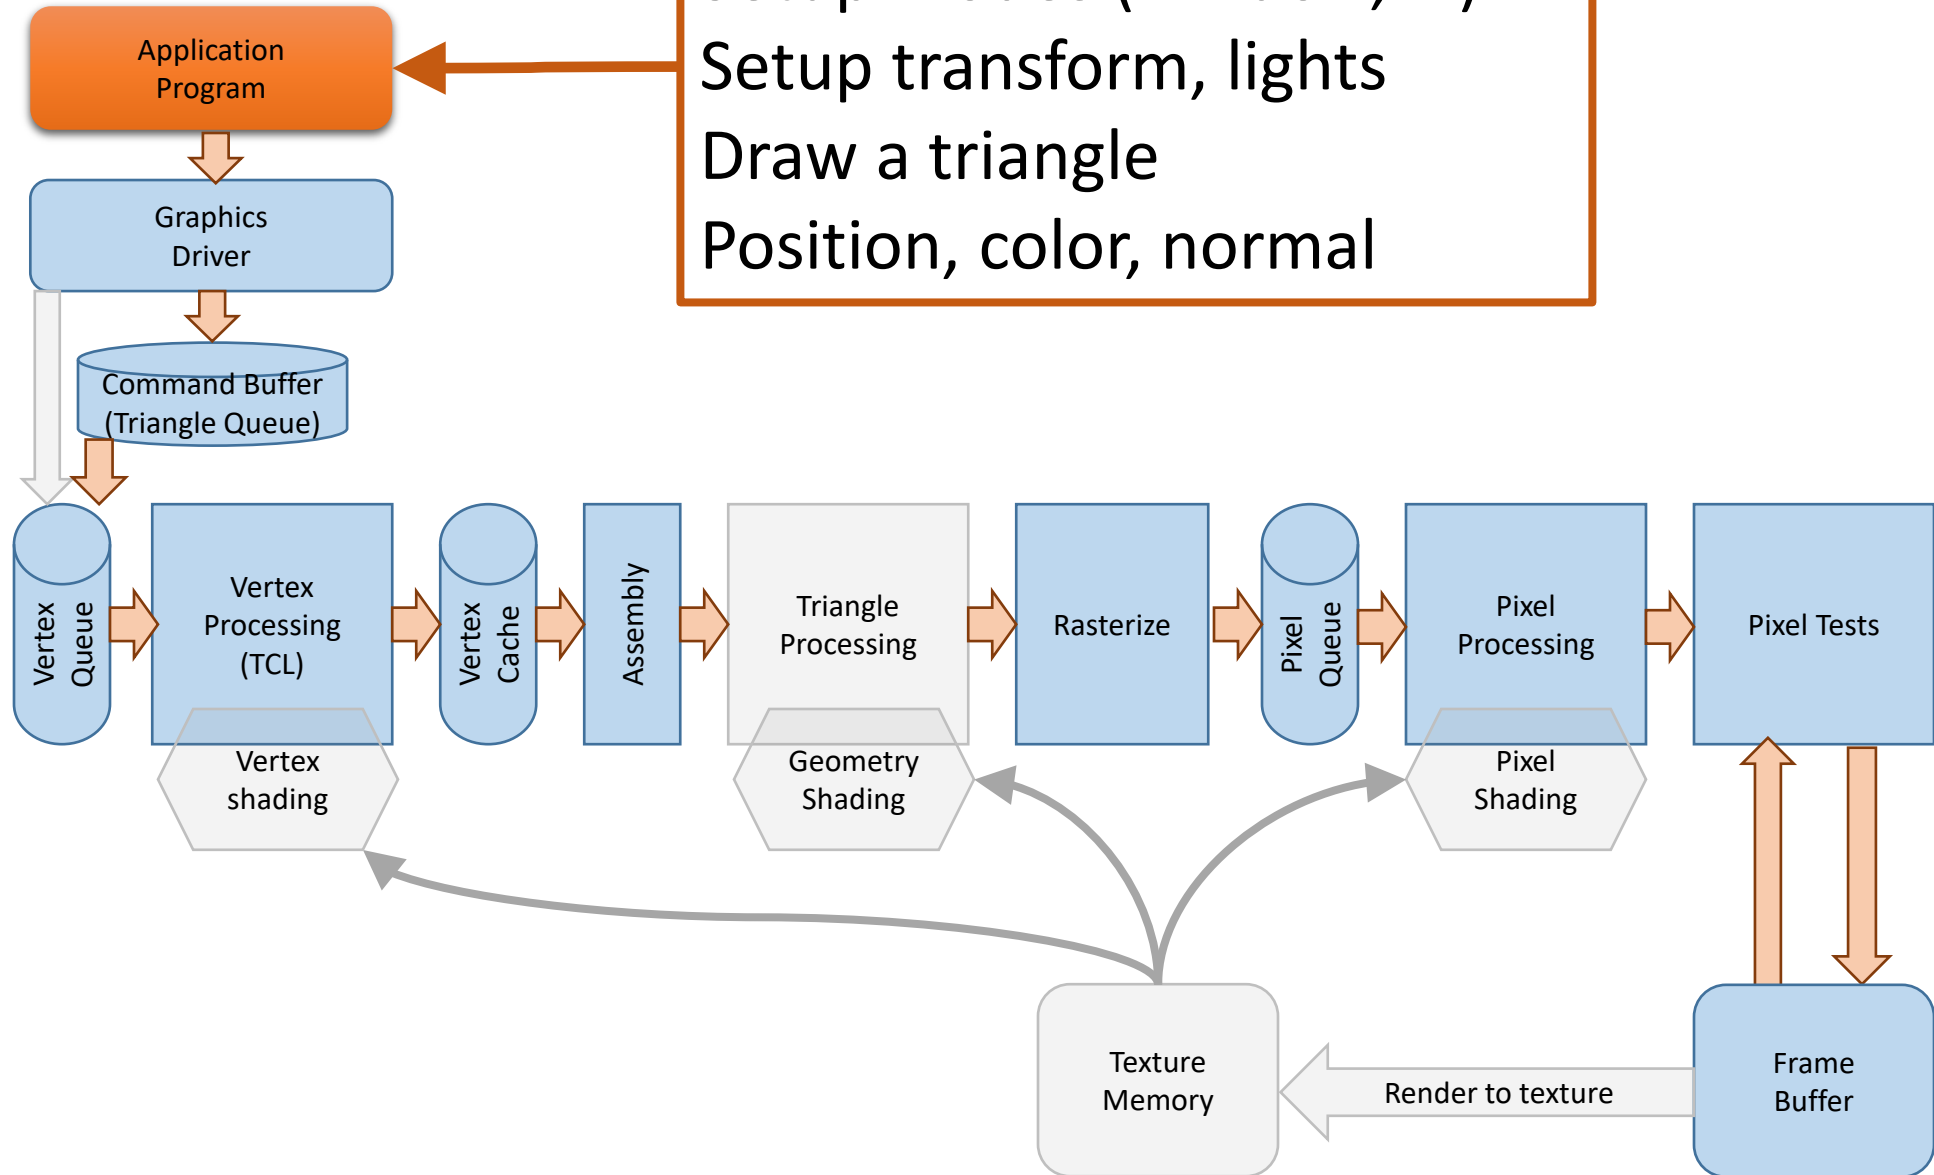
## Start here

Setup modes (window, ...)

Setup transform, lights

Draw a triangle

Position, color, normal





# Drawing a triangle

**Modes** per triangle (group)

which window, how to fill, use z-buffer, ...

**Data** per-vertex

position

normal

color

other things (texture coords)

# Per Vertex?

Modes per triangle

which window, how to fill, use z-buffer, ...

Data per-vertex

position

**normal ← allow us to make non-flat**

**color ← allows us to interpolate**

other things (texture coords)

# Per-Vertex not Per-Triangle

Allows sharing vertices between triangles

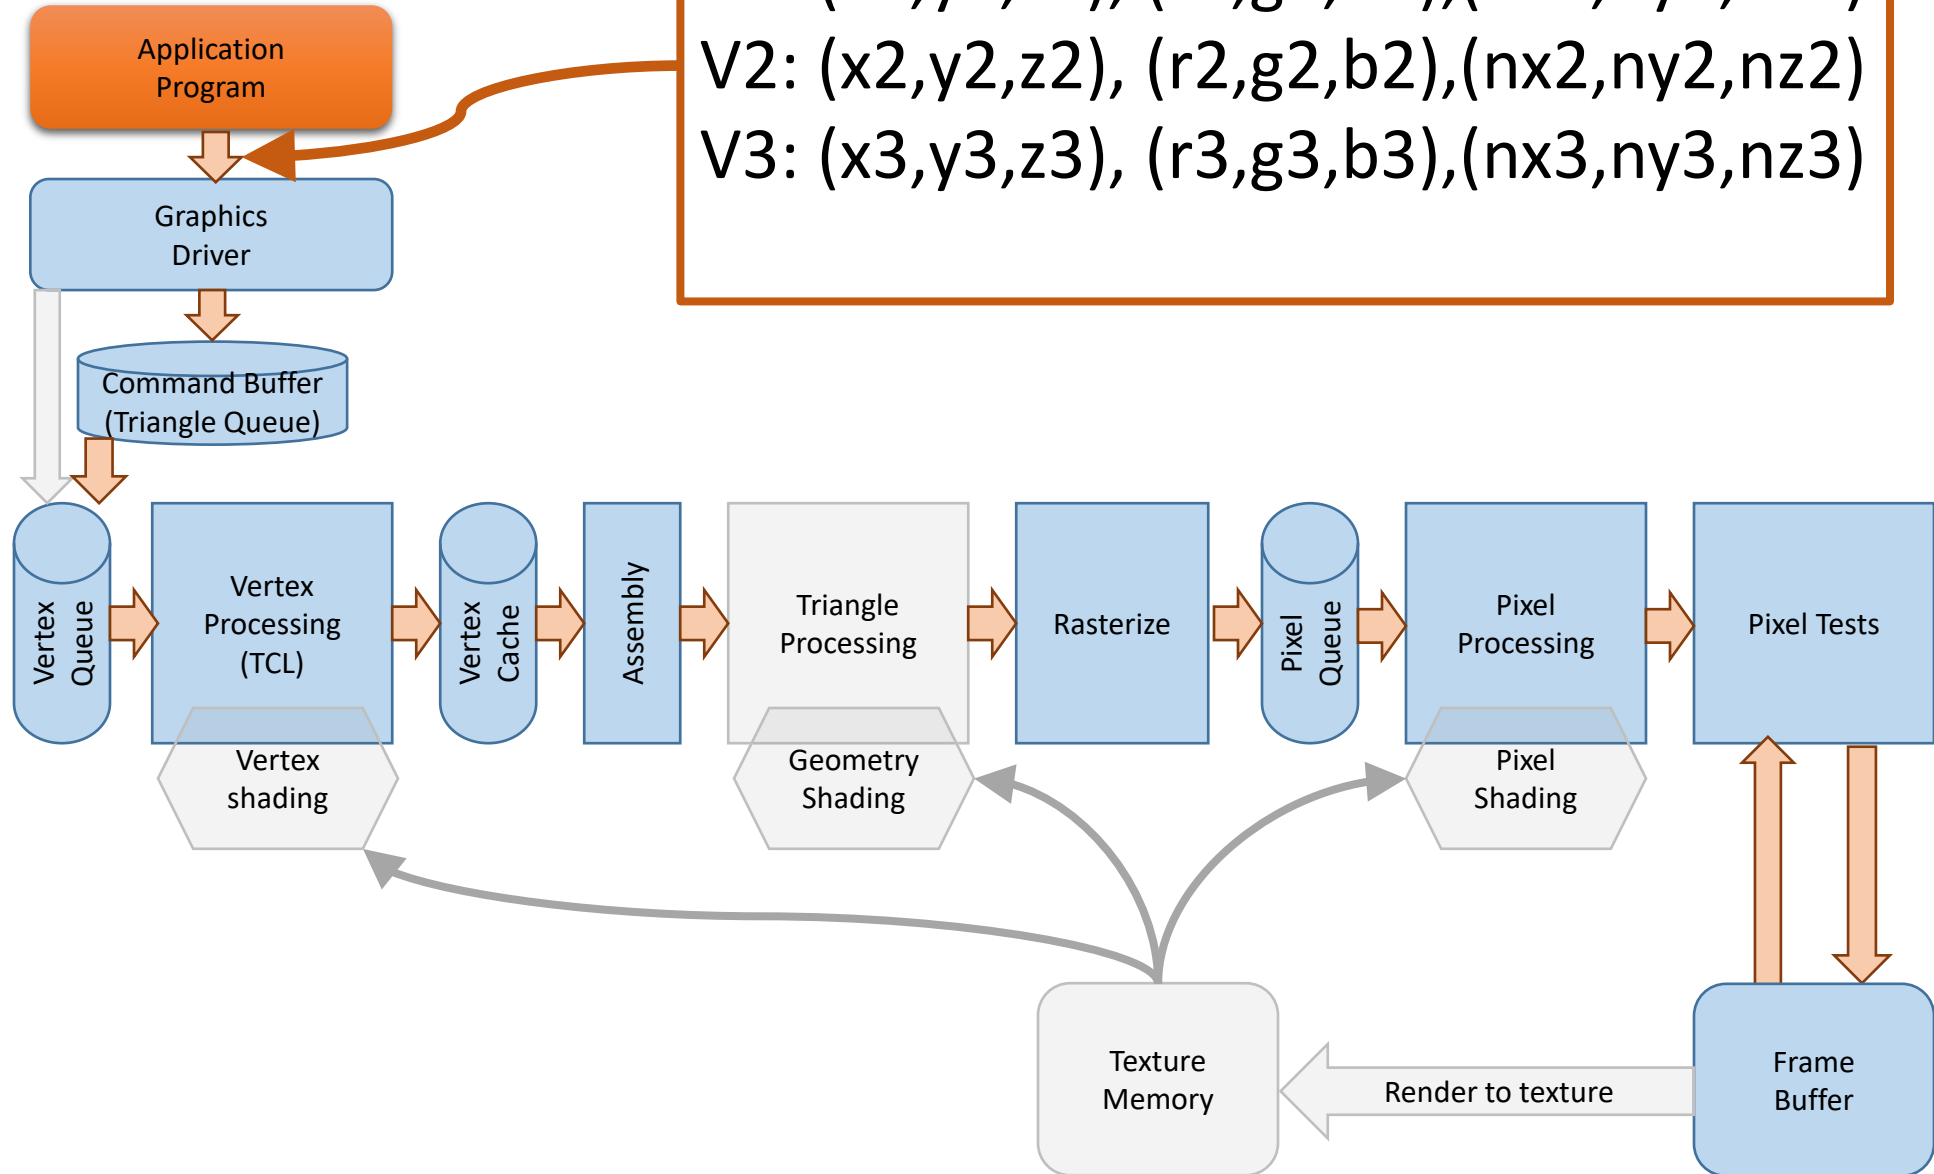
Or make all the vertices the same  
(color, normal, ...) to get truly flat

# Triangle

V1: (x1,y1,z1), (r1,g1,b1),(nx1,ny1,nz1)

V2: (x2,y2,z2), (r2,g2,b2),(nx2,ny2,nz2)

V3: (x3,y3,z3), (r3,g3,b3),(nx3,ny3,nz3)

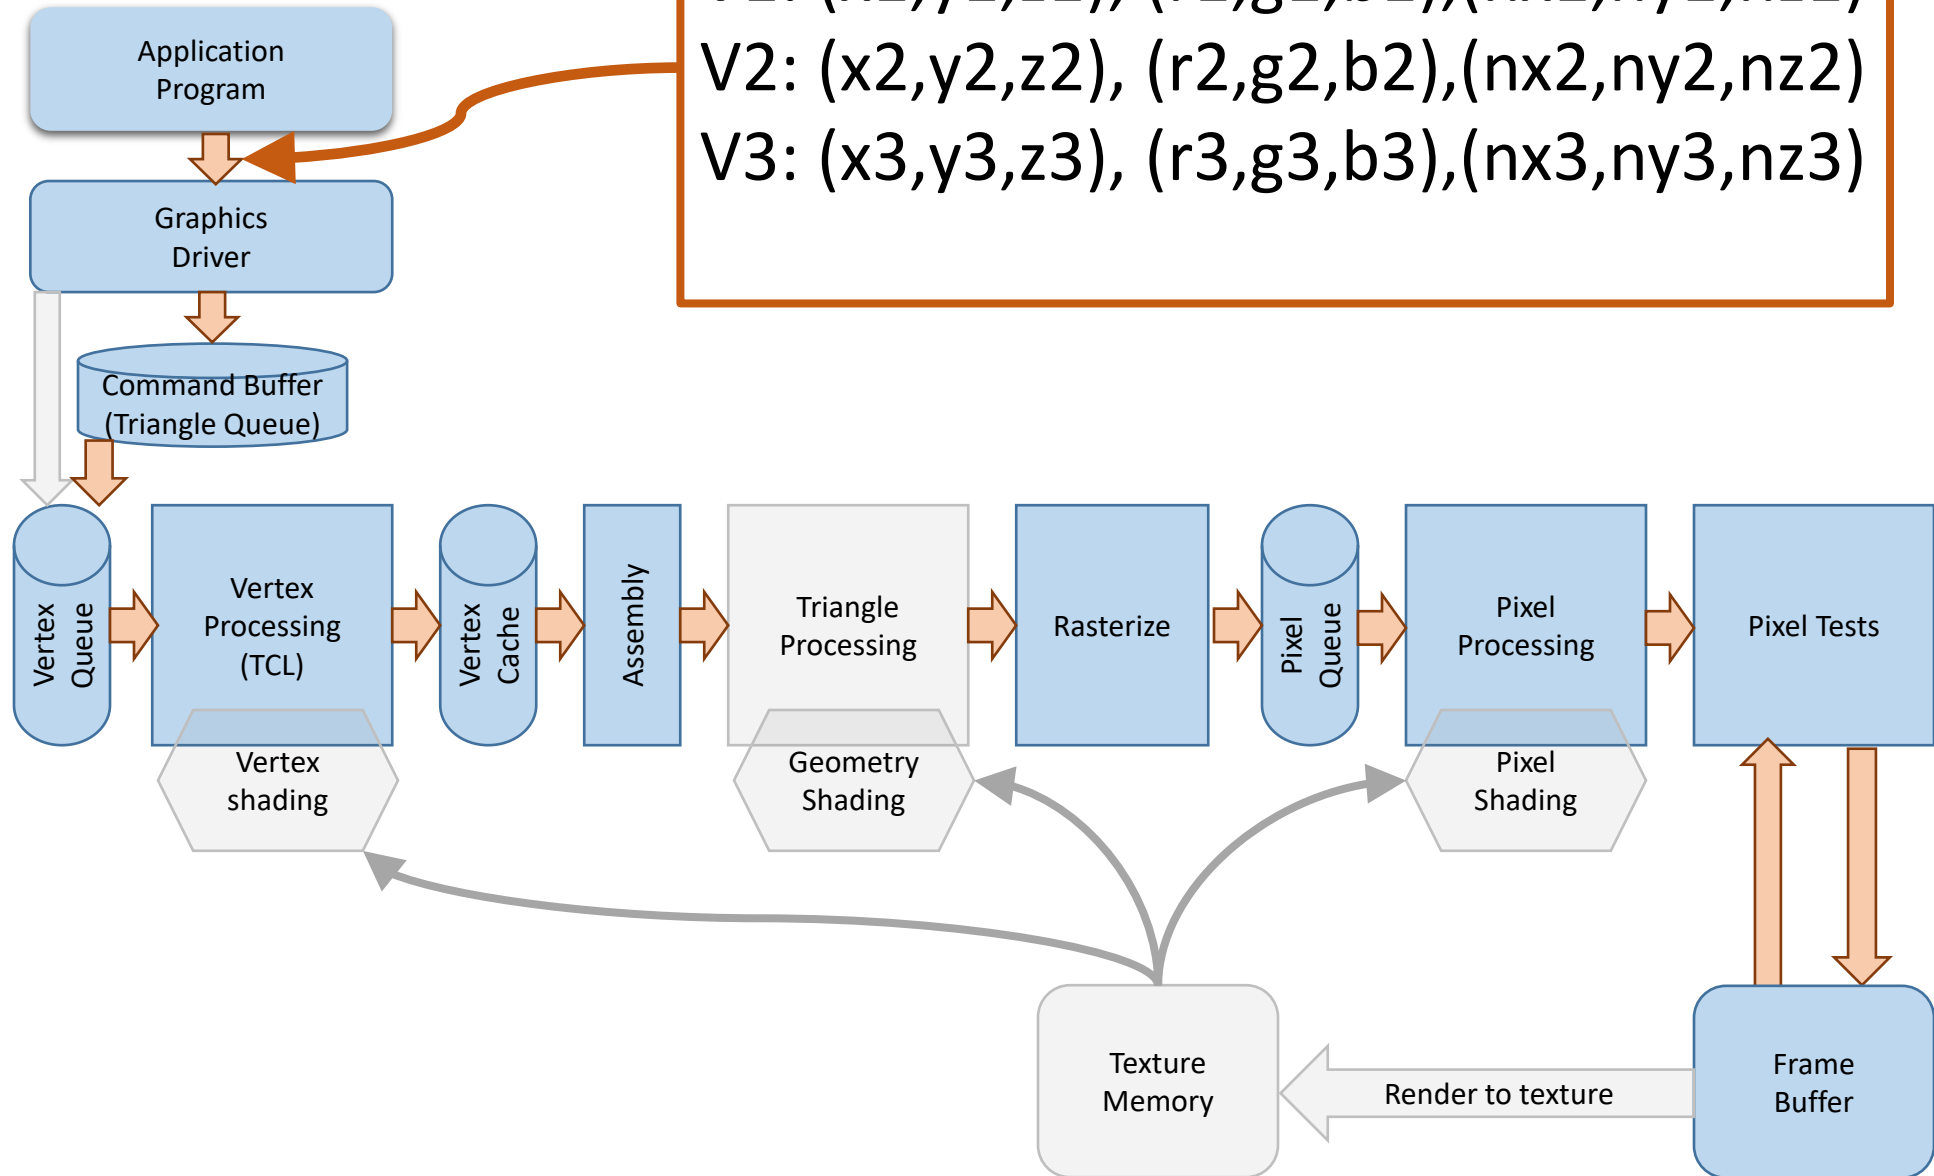


# Triangle

V1: (x1,y1,z1), (r1,g1,b1),(nx1,ny1,nz1)

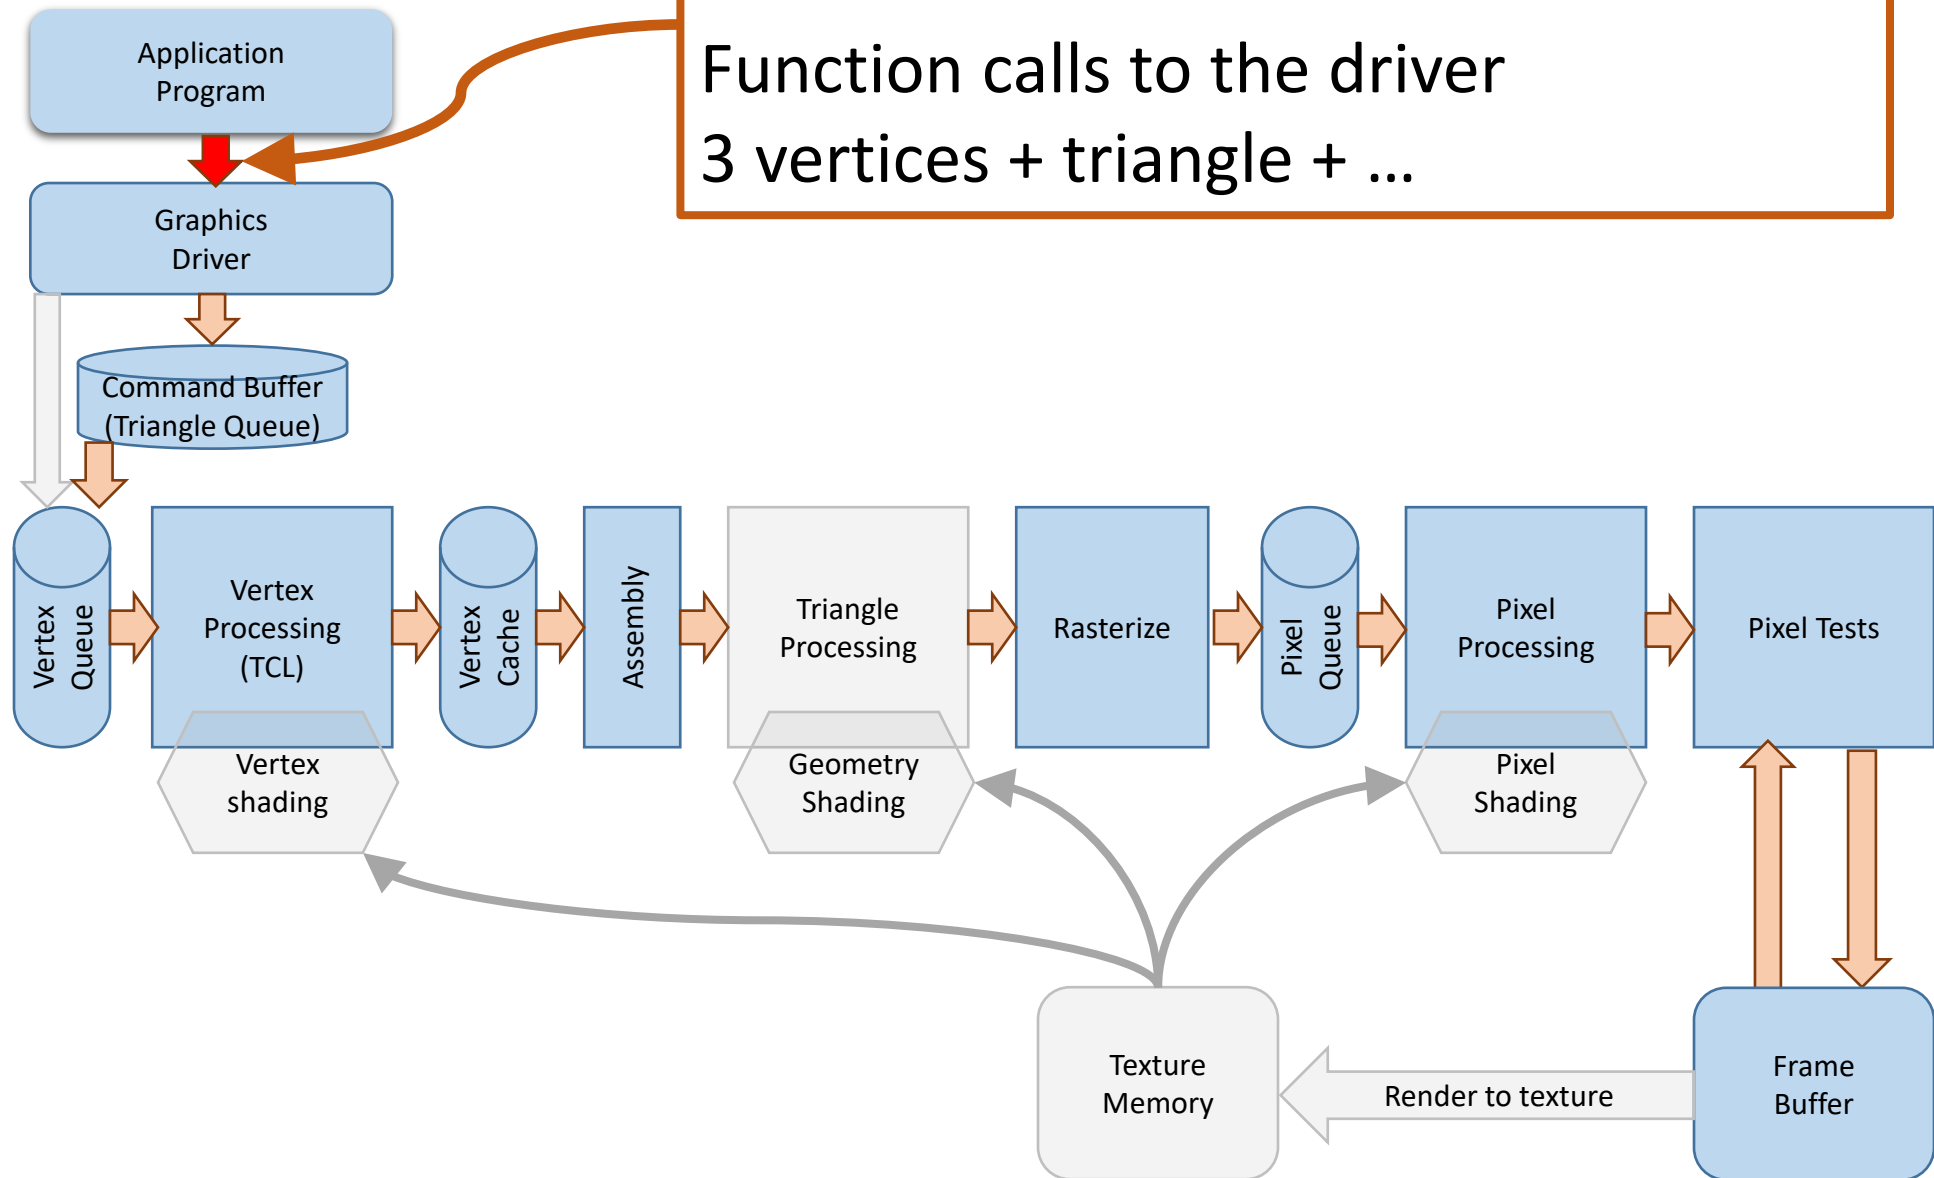
V2: (x2,y2,z2), (r2,g2,b2),(nx2,ny2,nz2)

V3: (x3,y3,z3), (r3,g3,b3),(nx3,ny3,nz3)



Is this a potential bottleneck?

Function calls to the driver  
3 vertices + triangle + ...



# Old style OpenGL

```
begin (TRIANGLE) ;  
  c3f (r1, g1, b1) ;  
  n3f (nx1, ny1, nz1) ;  
  v3f (x1, y1, z1) ;  
  c3f (r2, g2, b2) ;  
  n3f (nx2, ny2, nz2) ;  
  v3f (x2, y2, z2) ;  
  c3f (r3, g3, b3) ;  
  n3f (nx3, ny3, nz3) ;  
  v3f (x3, y3, z3) ;  
end (TRIANGLE) ;
```

11 function calls  
35 arguments pushed

Old days:  
This is a lot less than the  
number of pixels!

Nowadays:  
Just the memory access  
swamps the process

# Attribute buffers!

Block transfers of data

Data for lots of triangles moved as a block

Try to draw groups of triangles

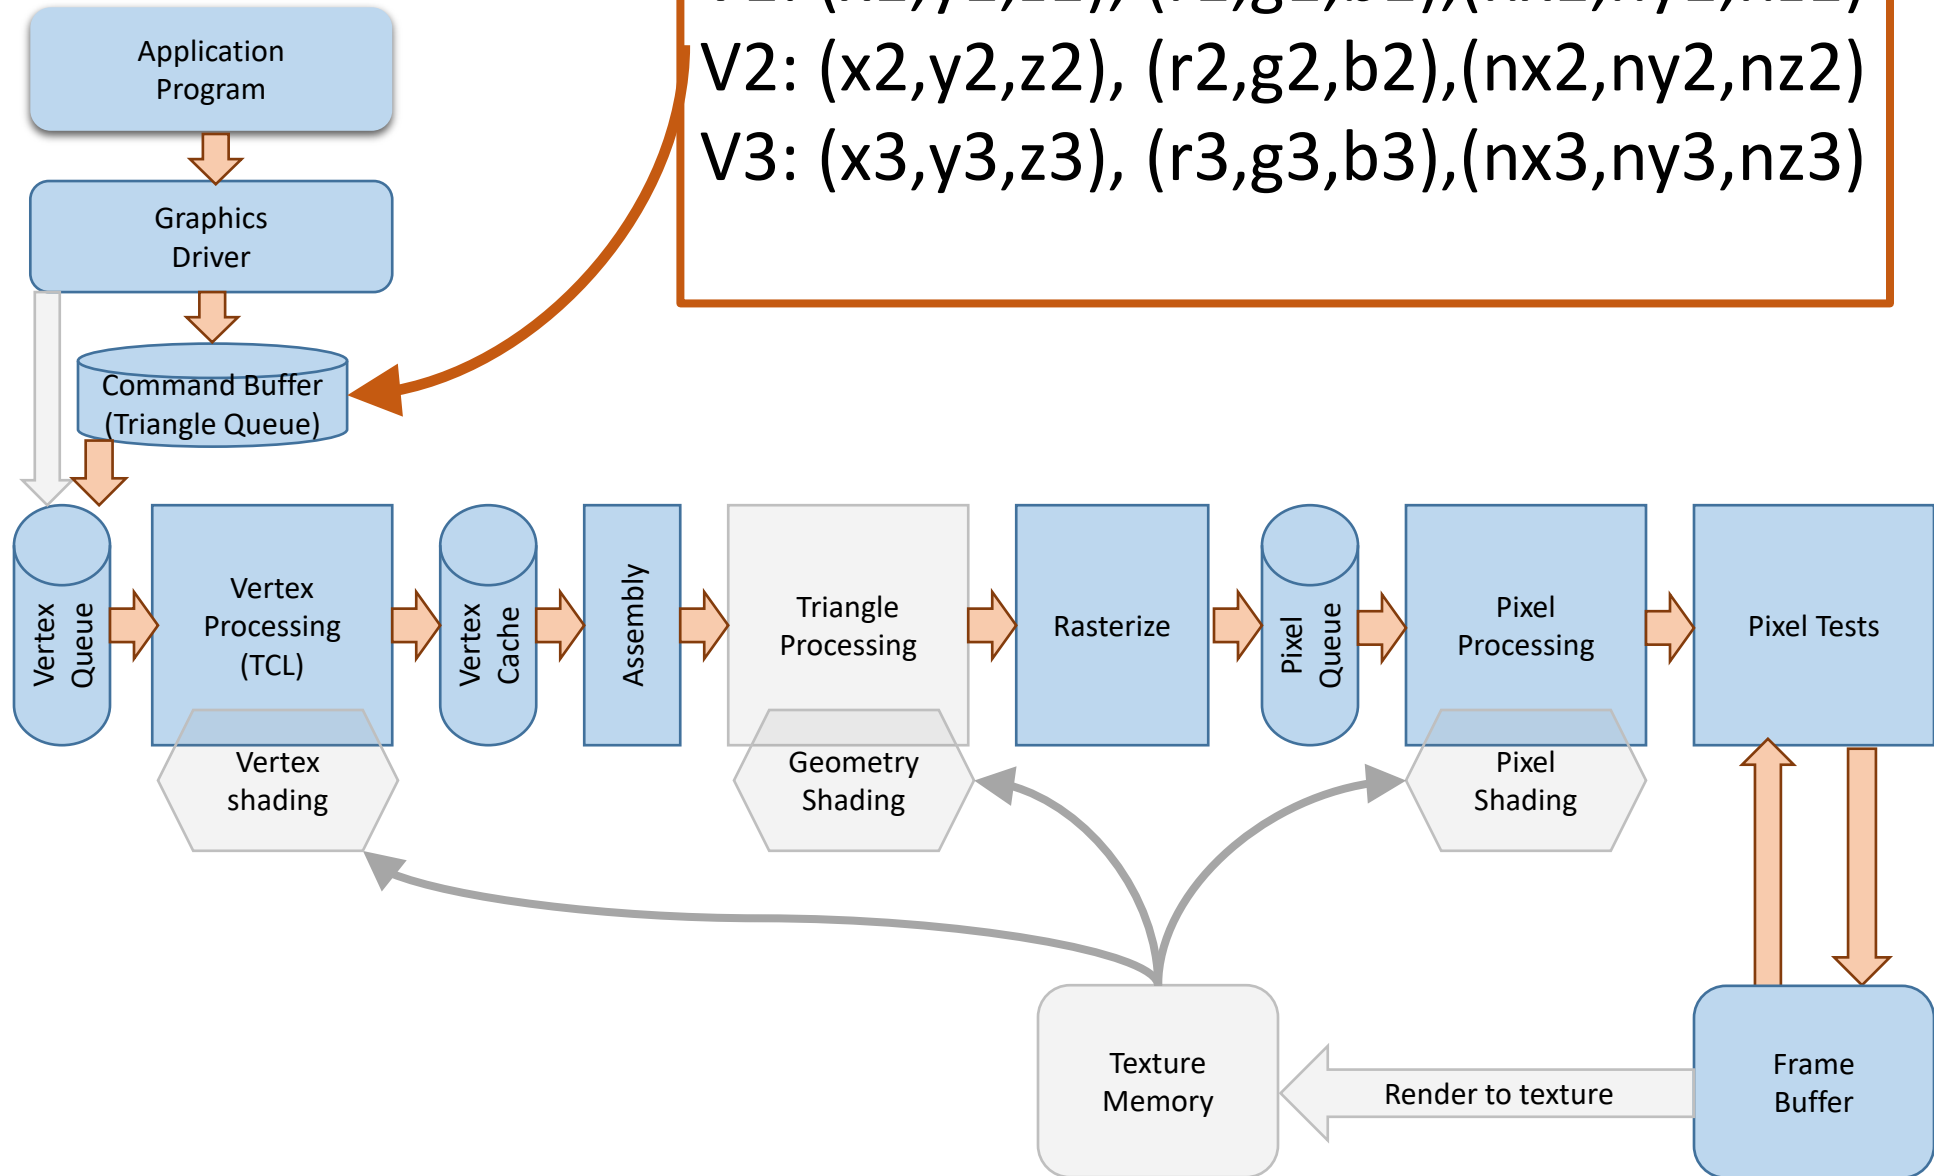


# Triangle

V1: (x1,y1,z1), (r1,g1,b1),(nx1,ny1,nz1)

V2: (x2,y2,z2), (r2,g2,b2),(nx2,ny2,nz2)

V3: (x3,y3,z3), (r3,g3,b3),(nx3,ny3,nz3)

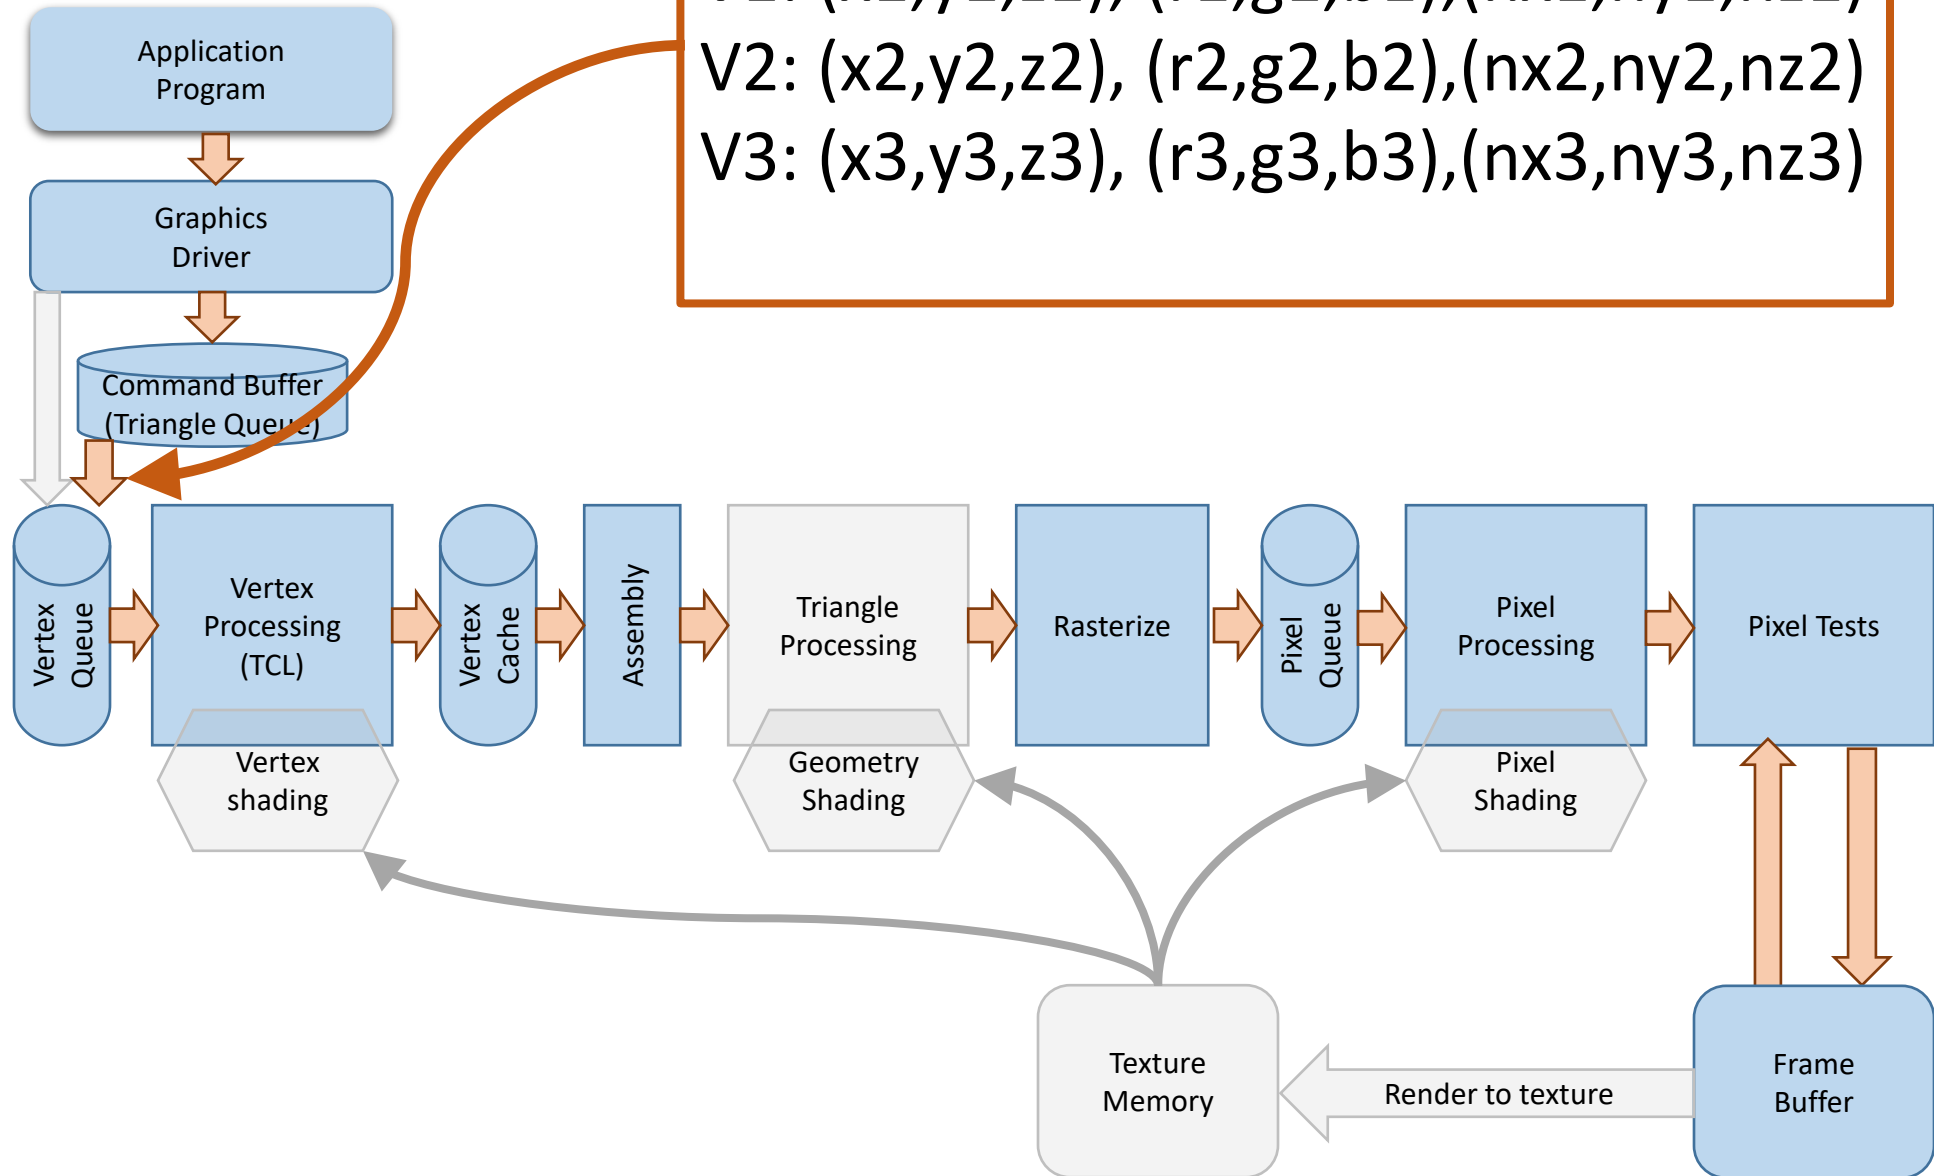


## Split up triangles into **vertices**

V1: (x1,y1,z1), (r1,g1,b1),(nx1,ny1,nz1)

V2:  $(x_2, y_2, z_2), (r_2, g_2, b_2), (nx_2, ny_2, nz_2)$

V3: (x3,y3,z3), (r3,g3,b3),(nx3,ny3,nz3)

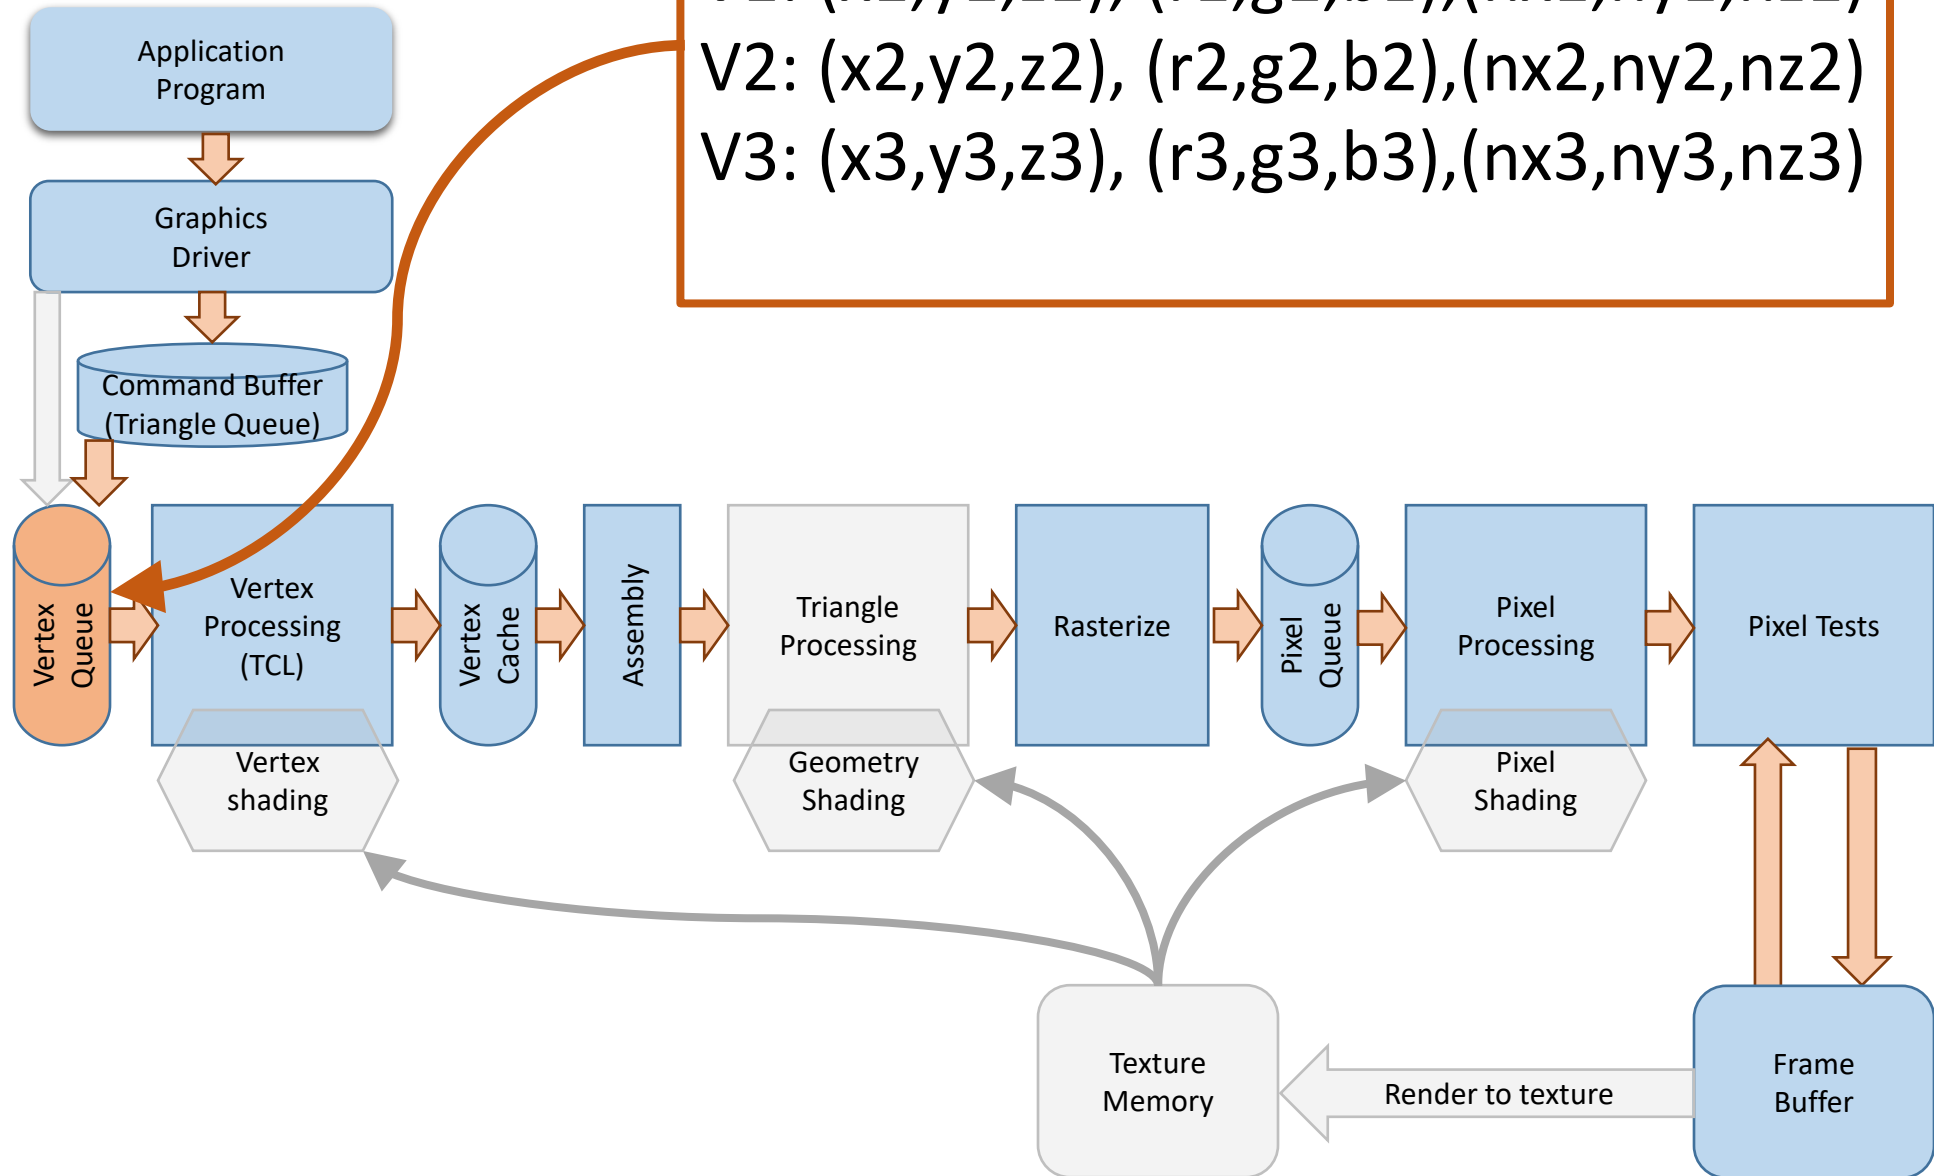


## Buffer / Queue the **vertices**

V1: (x1,y1,z1), (r1,g1,b1),(nx1,ny1,nz1)

V2: (x2,y2,z2), (r2,g2,b2),(nx2,ny2,nz2)

V3: (x3,y3,z3), (r3,g3,b3),(nx3,ny3,nz3)



# Buffering Vertices

Old Days:

- Vertex processing expensive

- Try to maximize re-use

- Process once and use for many triangles

Nowadays

- Getting vertex to hardware is expensive

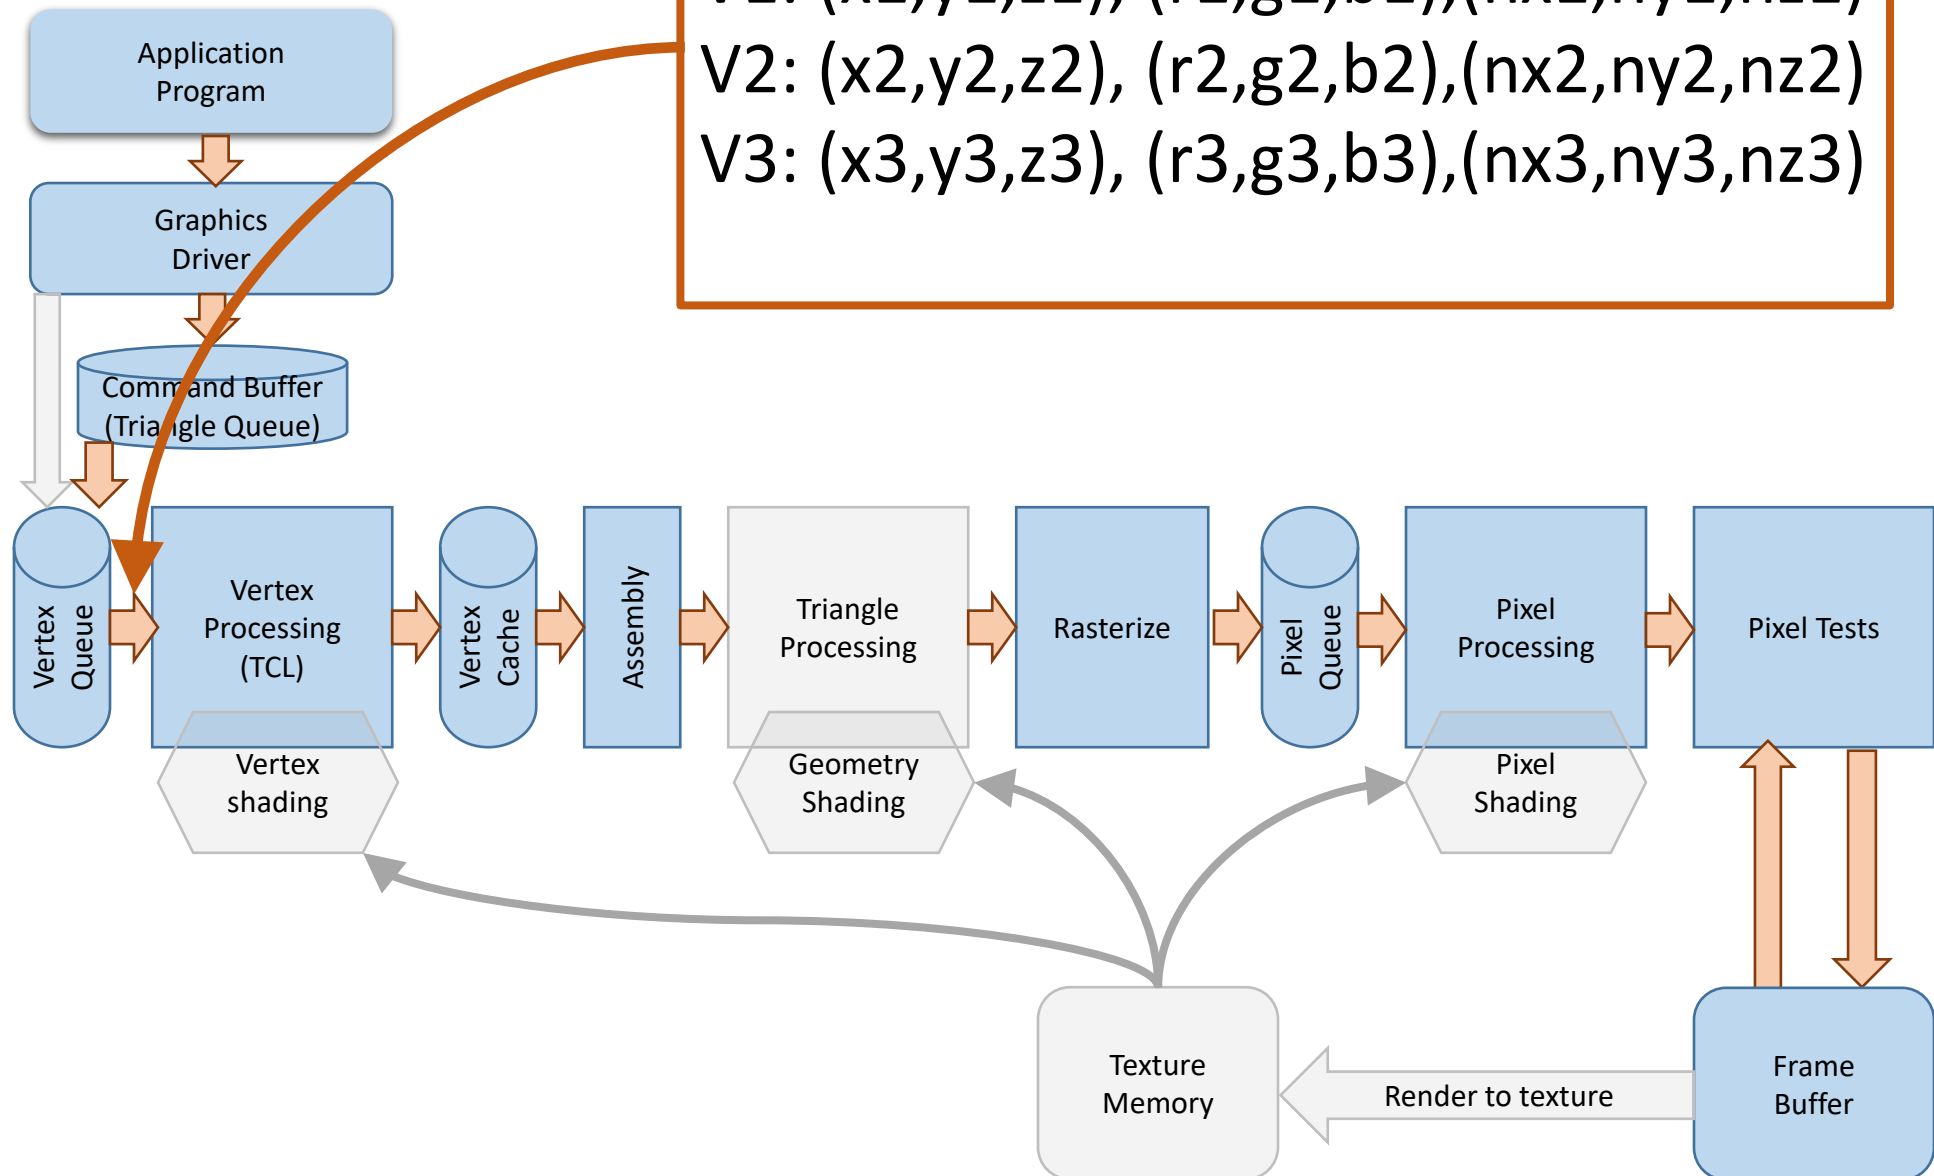
- Process vertices in parallel

## Buffer / Queue the **vertices**

V1: (x1,y1,z1), (r1,g1,b1),(nx1,ny1,nz1)

V2: (x2,y2,z2), (r2,g2,b2),(nx2,ny2,nz2)

V3: (x3,y3,z3), (r3,g3,b3),(nx3,ny3,nz3)

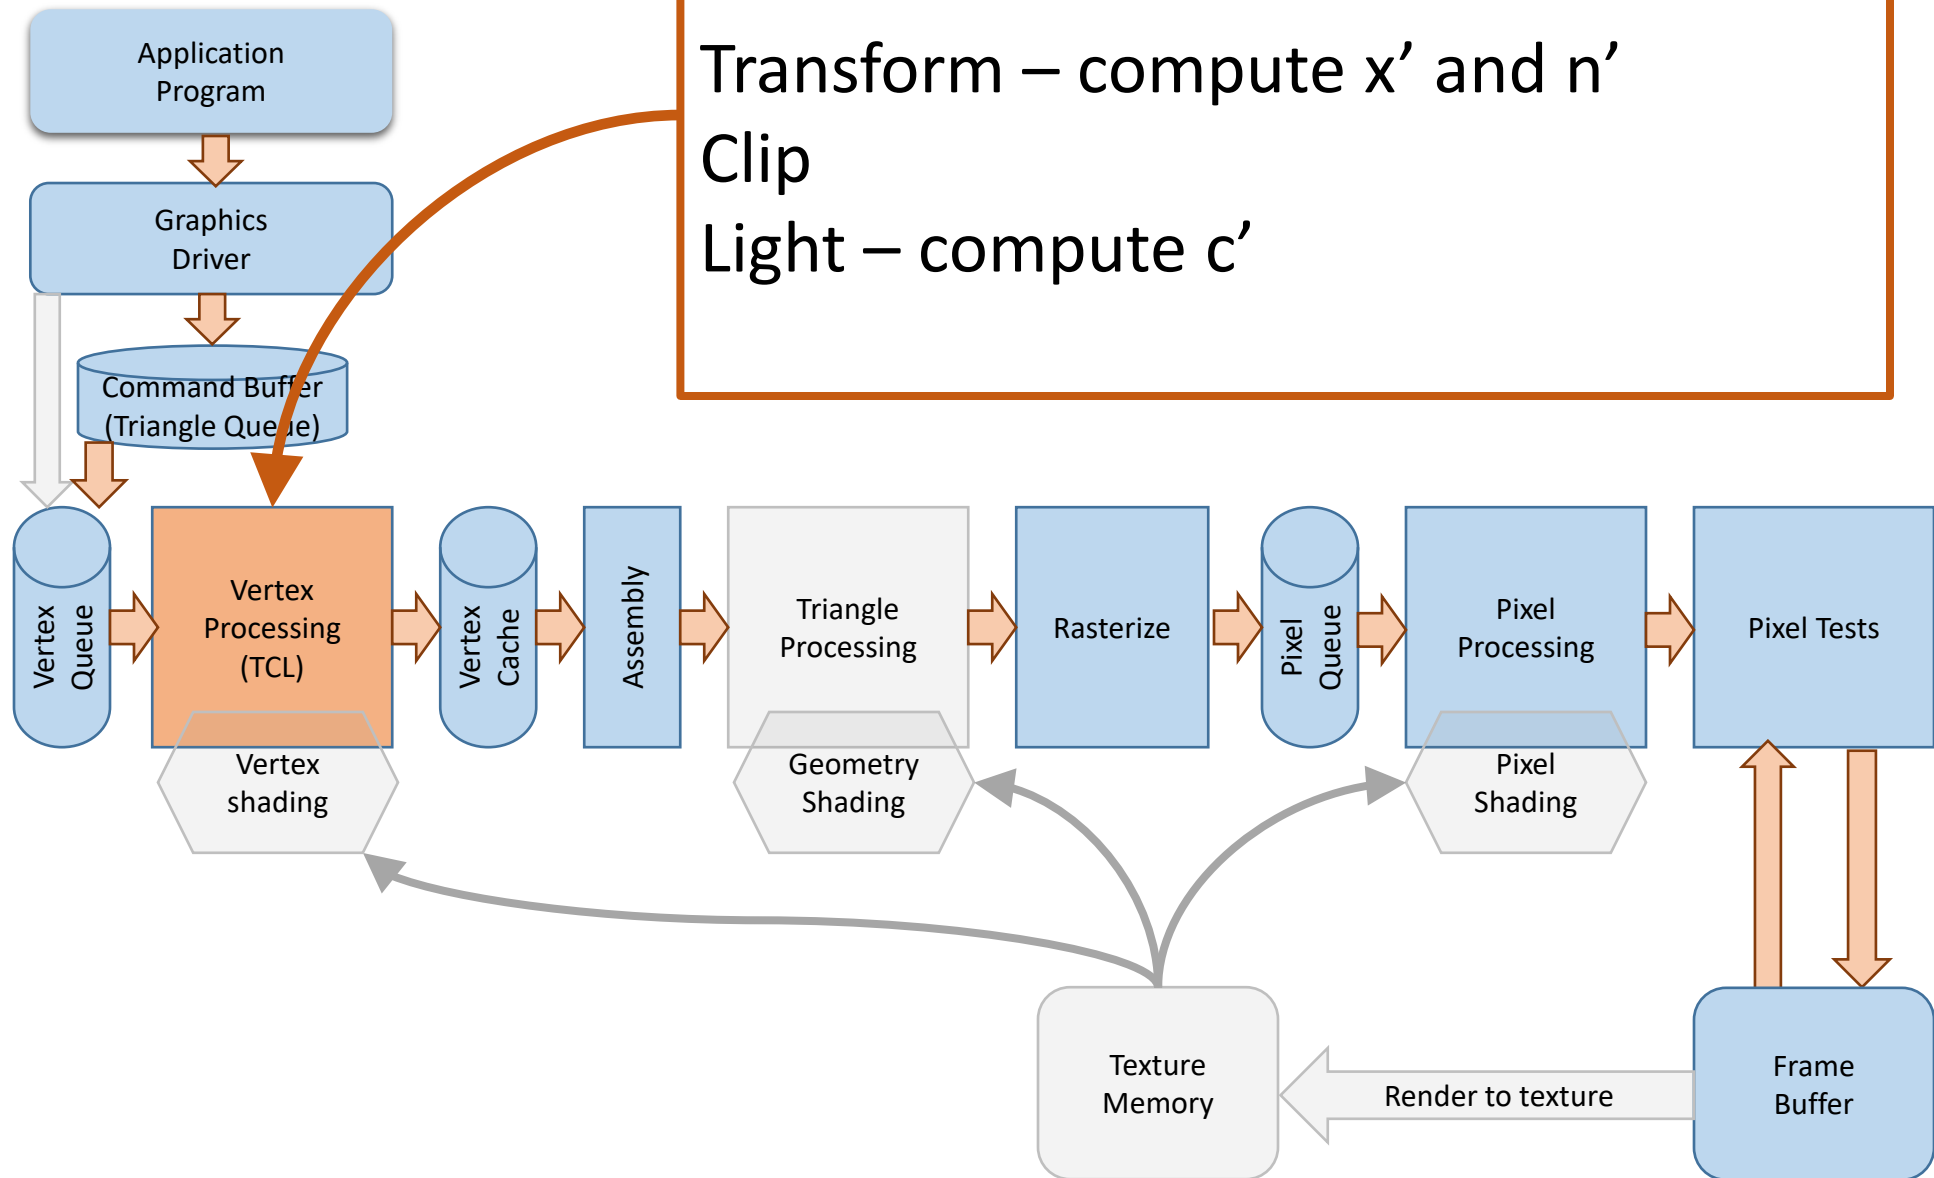


Process each vertex independently

Transform – compute  $x'$  and  $n'$

Clip

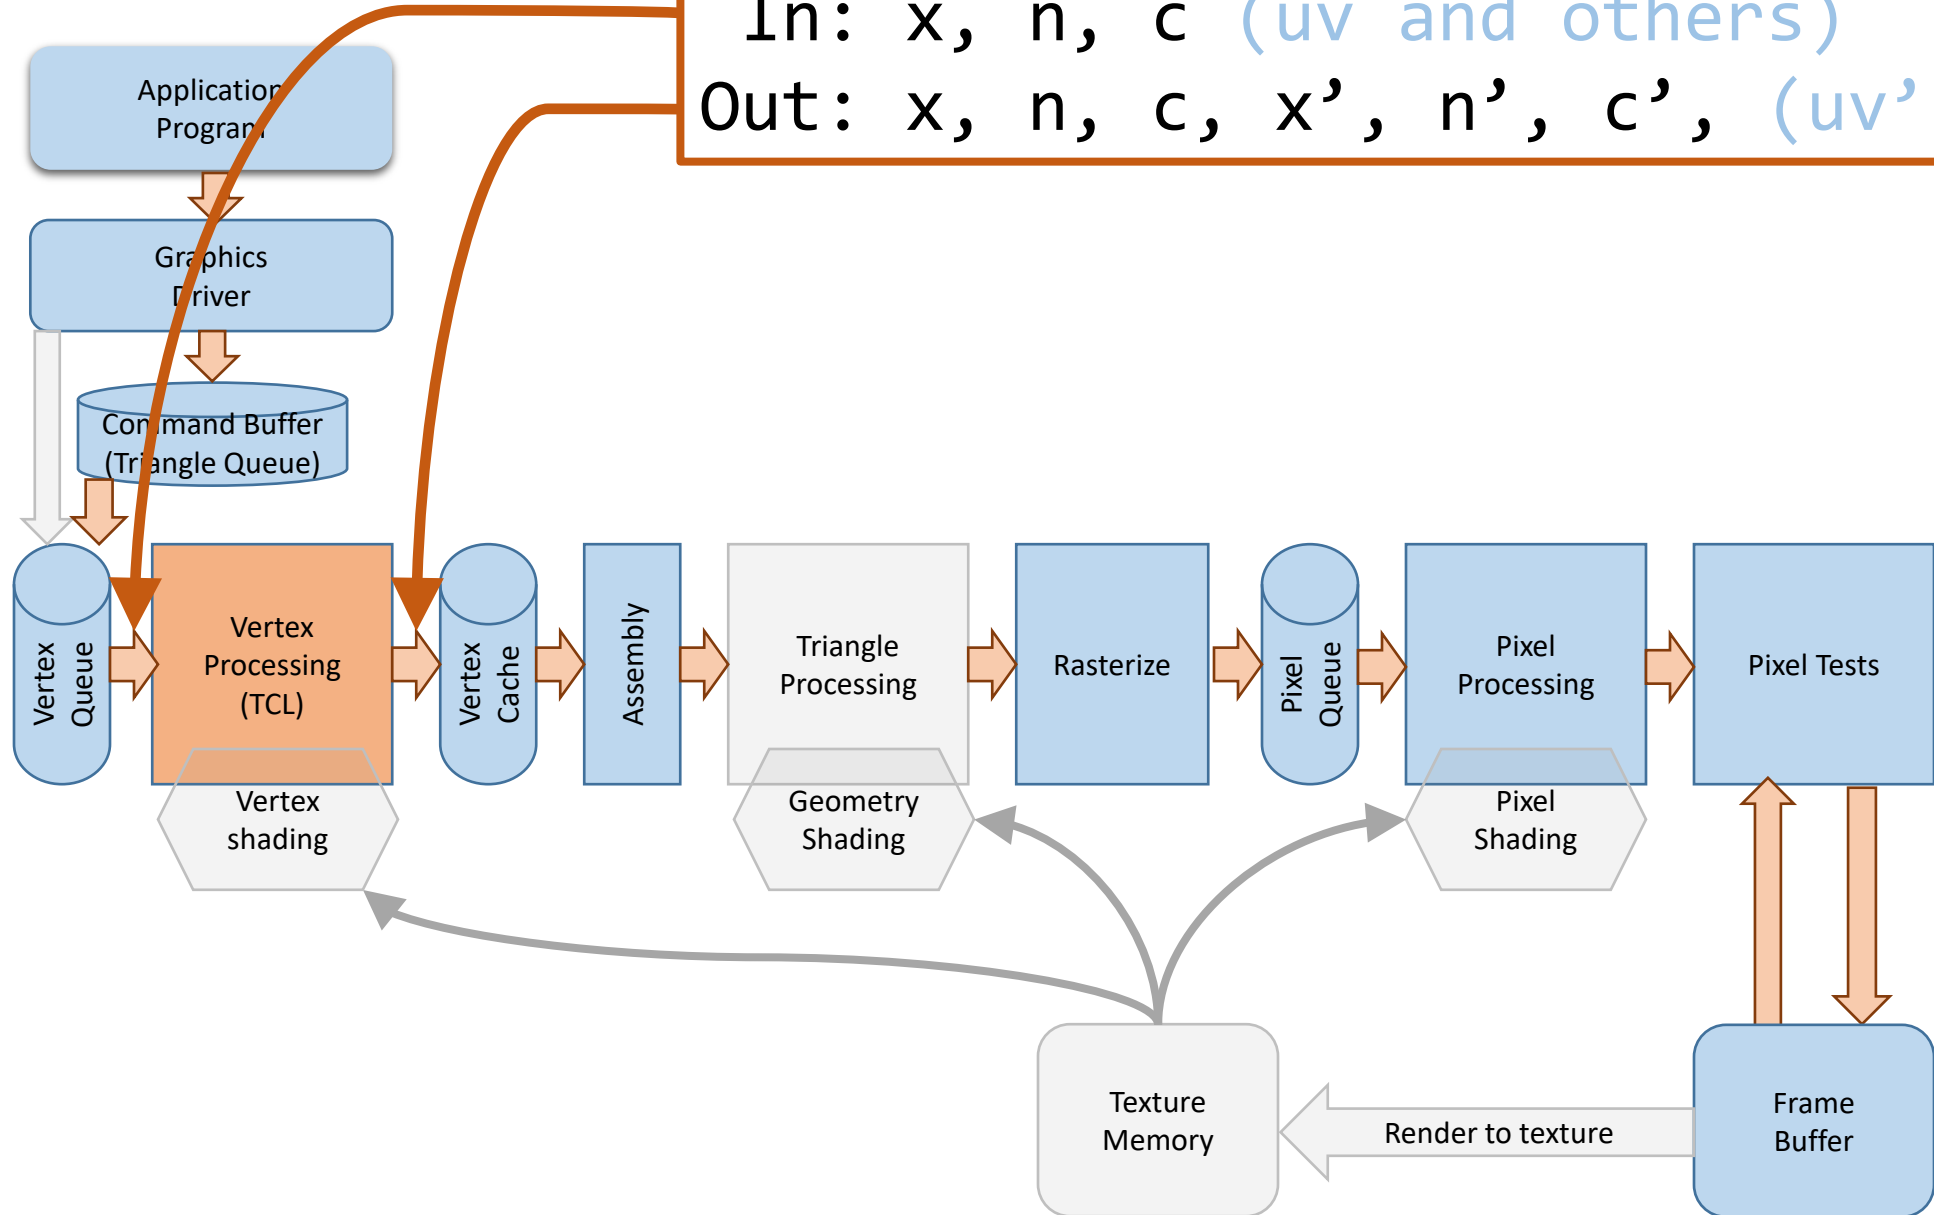
Light – compute  $c'$



Vertex in  $\rightarrow$  Vertex out

In:  $x, n, c$  (uv and others)

Out:  $x, n, c, x', n', c', (uv' \dots)$



# Vertex Processing

Just adds information to vertices

Computes transformation

screen space positions, normals

Computes “lighting”

new colors

(in the old days, clipping done here hence TCL)

(in the old days, lighting done per-vertex)



# Vertex Processing:

## Each vertex is independent

Inputs are:

- vertex information for **this vertex**

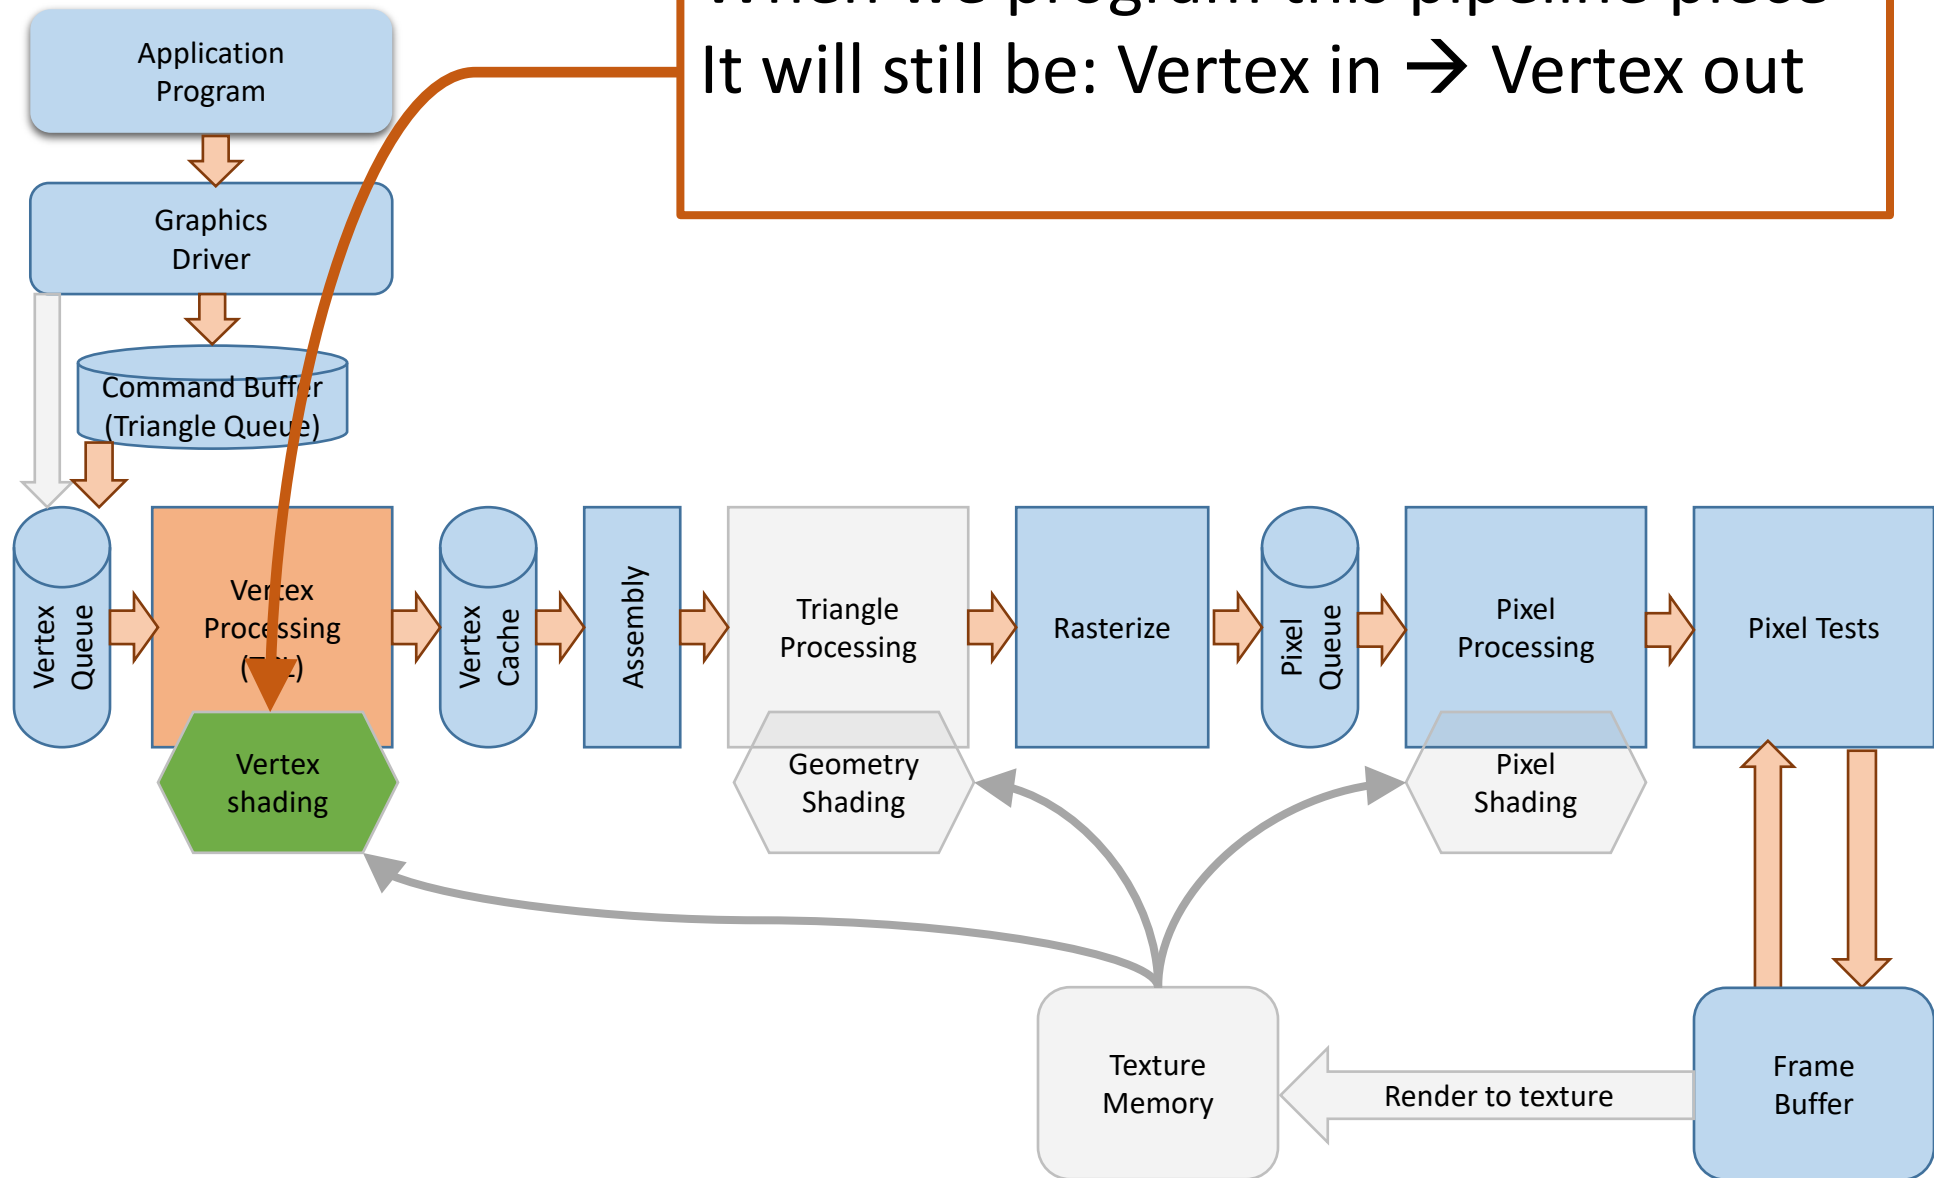
- any “global” information

- current transform, lighting, ...

Outputs are:

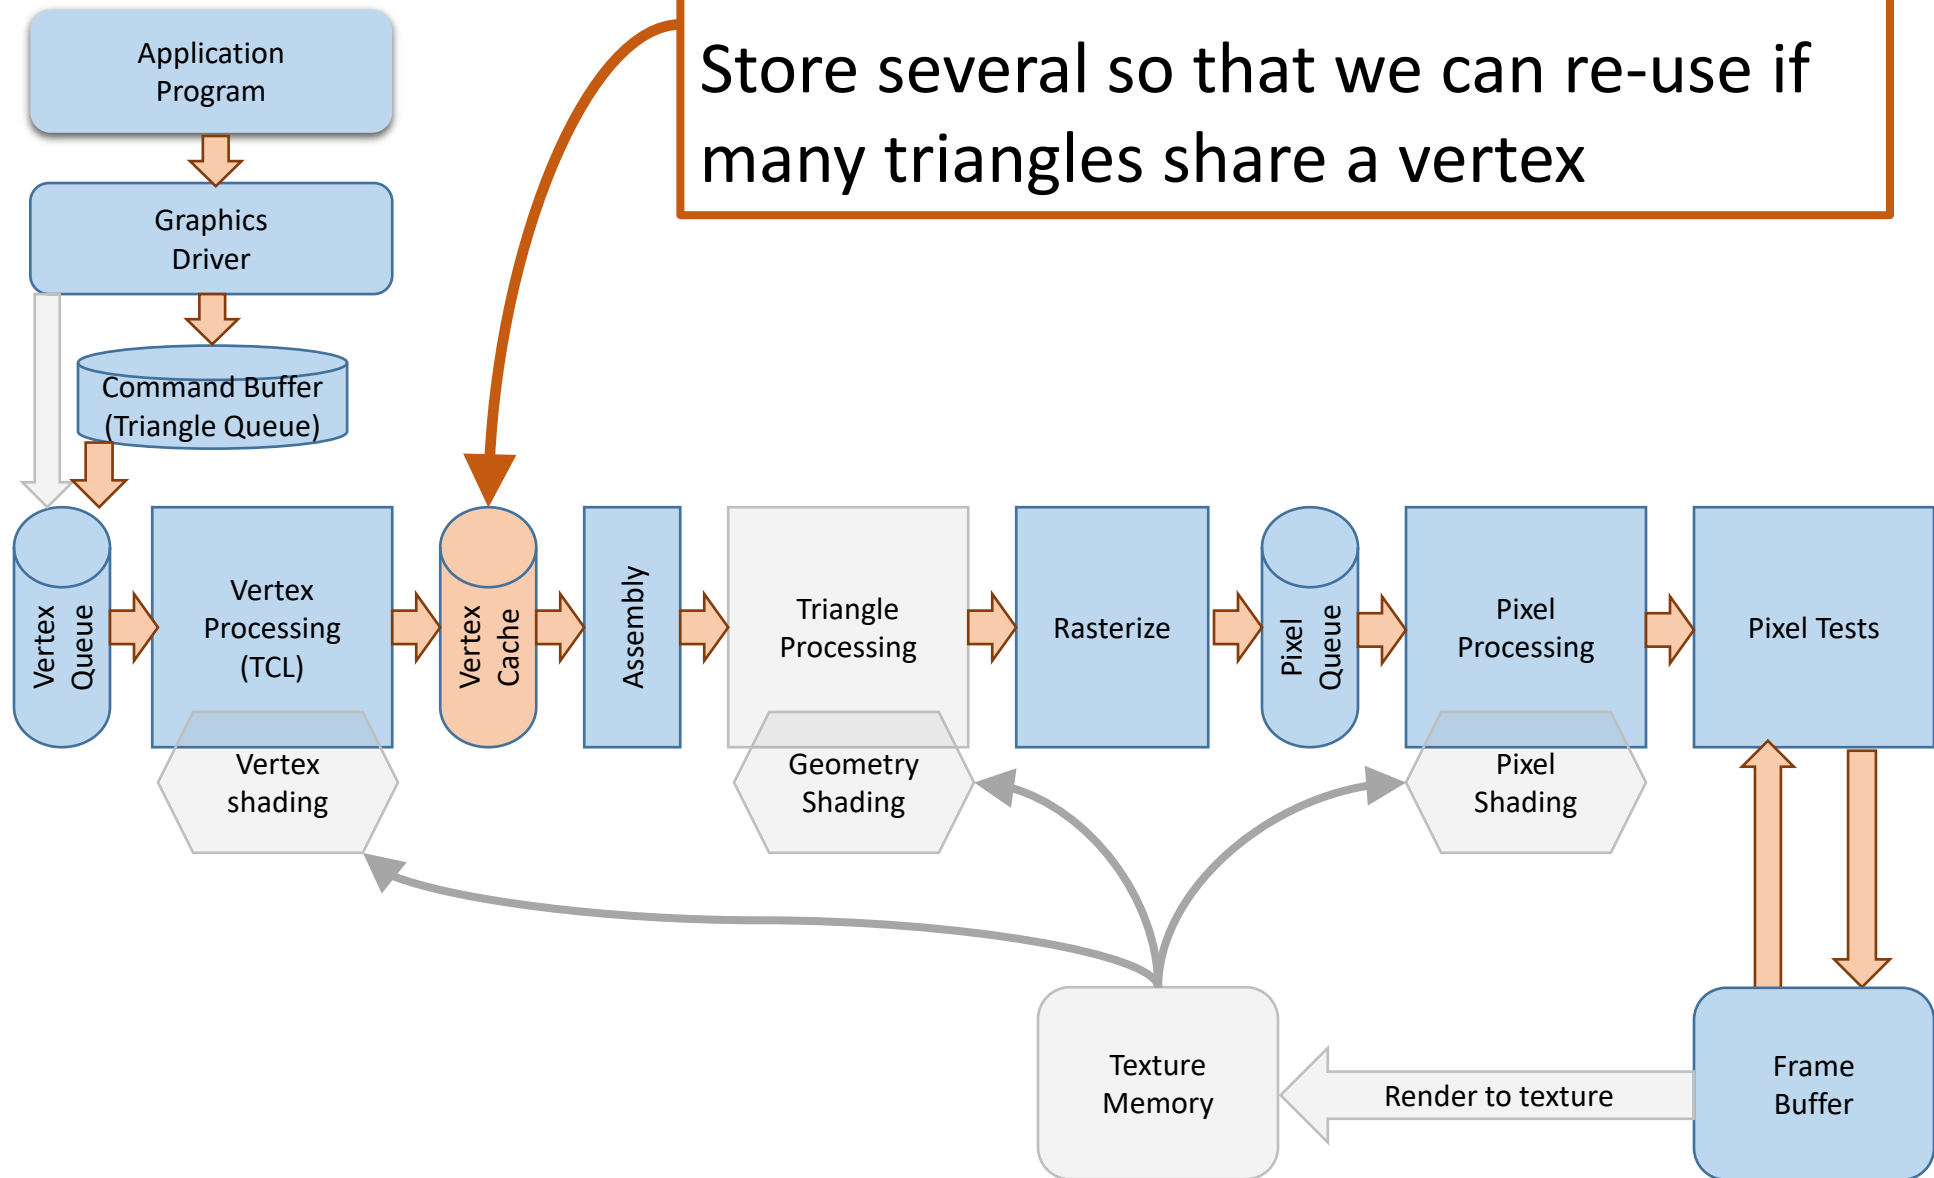
- vertex information for **this vertex**

Looking ahead...  
When we program this pipeline piece  
It will still be: Vertex in → Vertex out



Store processed vertices in a **cache**

Store several so that we can re-use if many triangles share a vertex



# Vertex Caching

Old days:

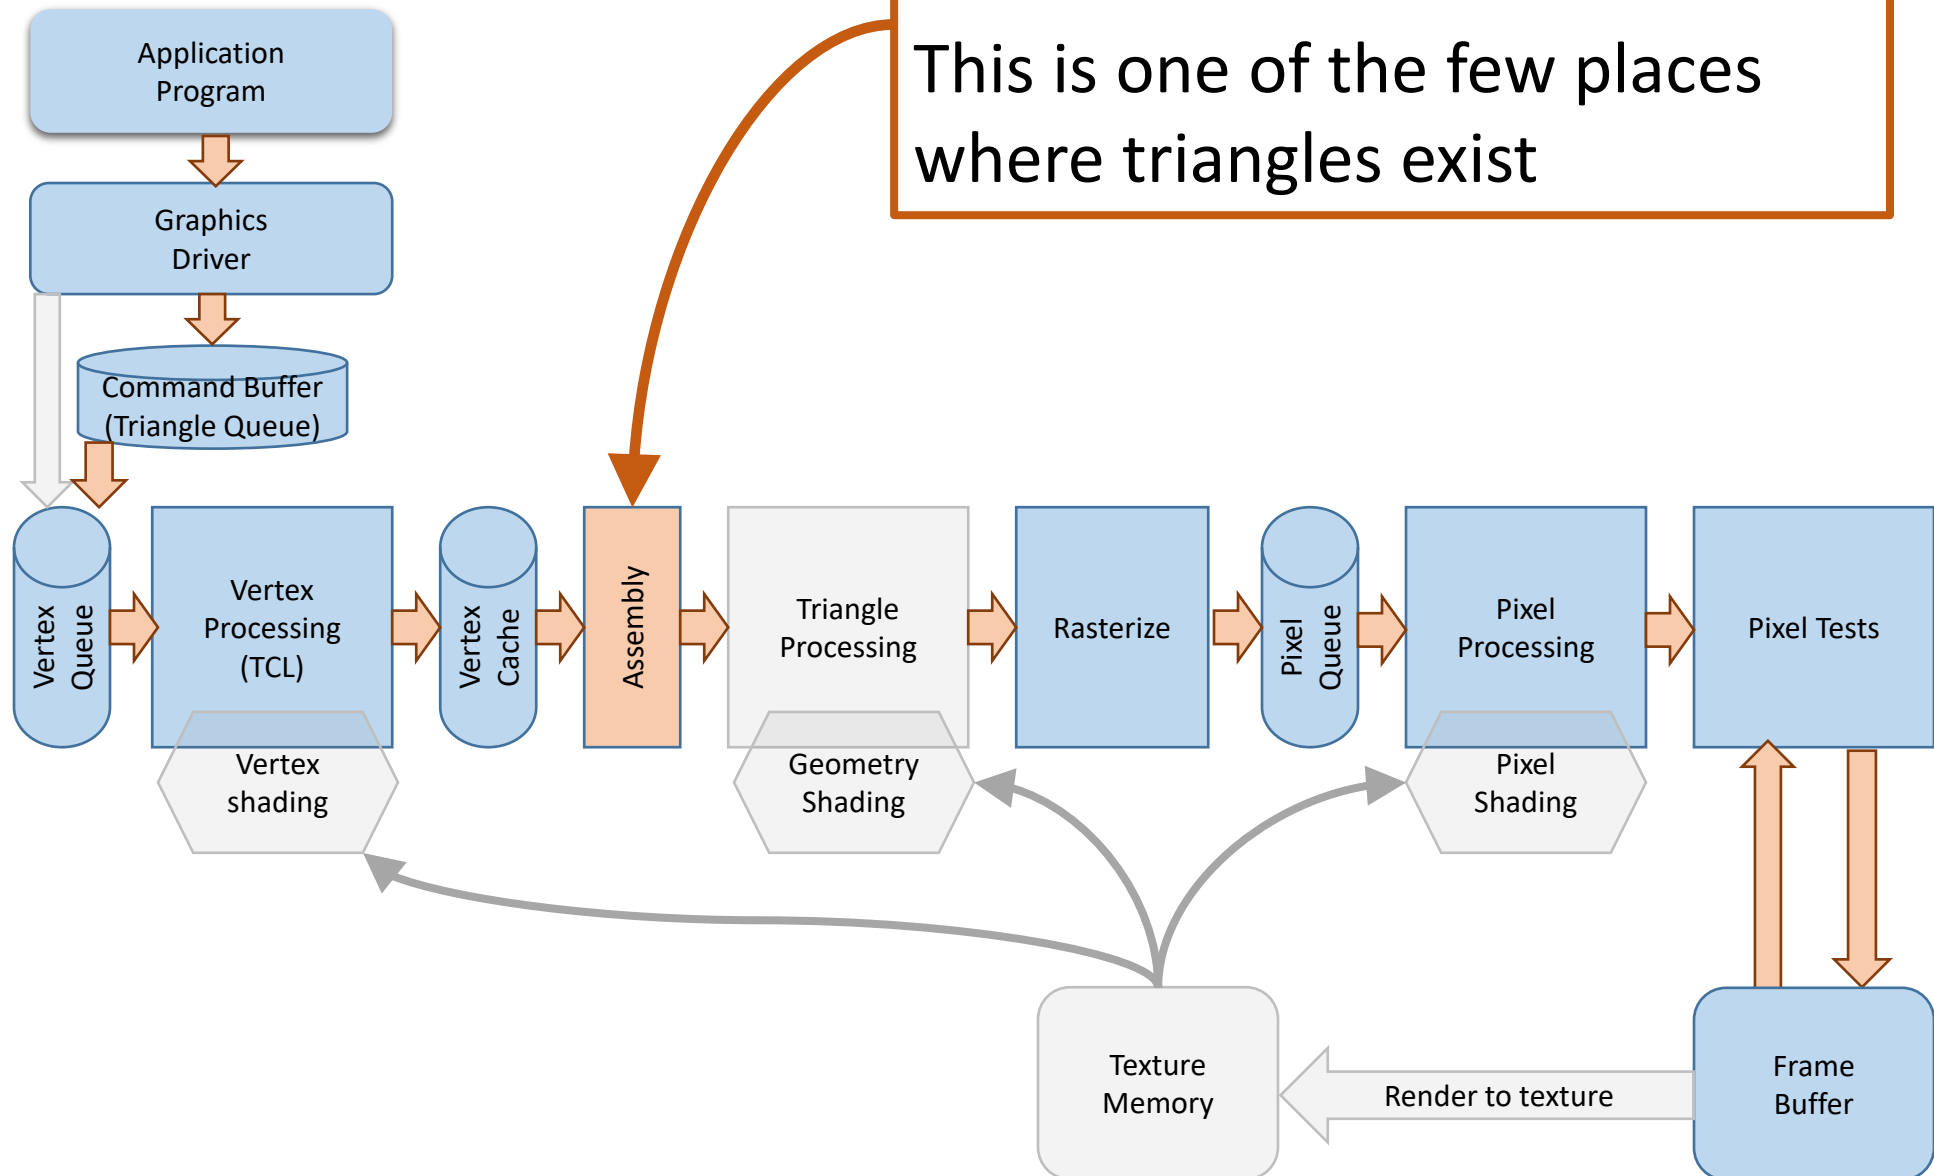
Big deal, important for performance

Now:

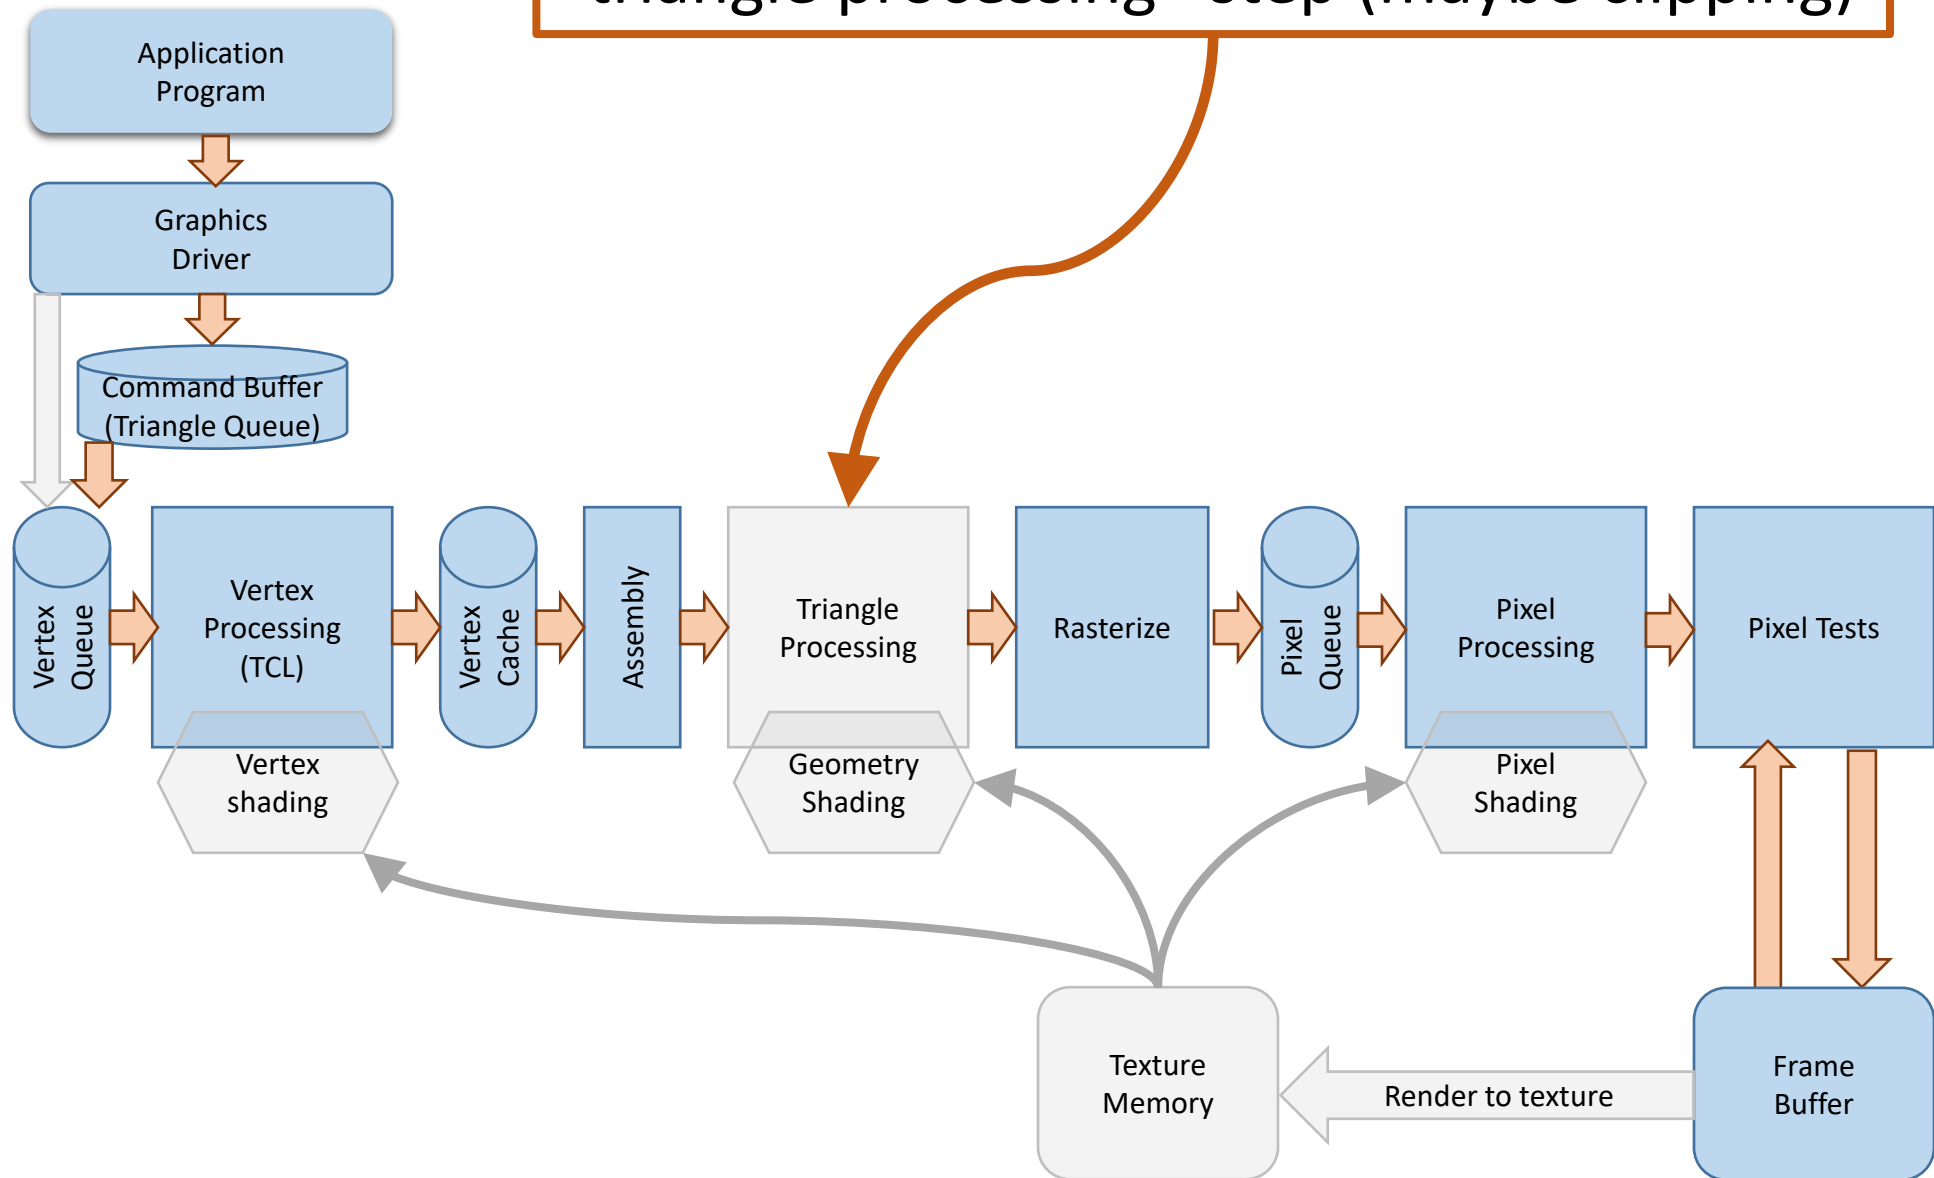
Not even sure that it's always done

Put triangles back together

This is one of the few places  
where triangles exist



In the fixed-function pipeline, there is no “triangle processing” step (maybe clipping)

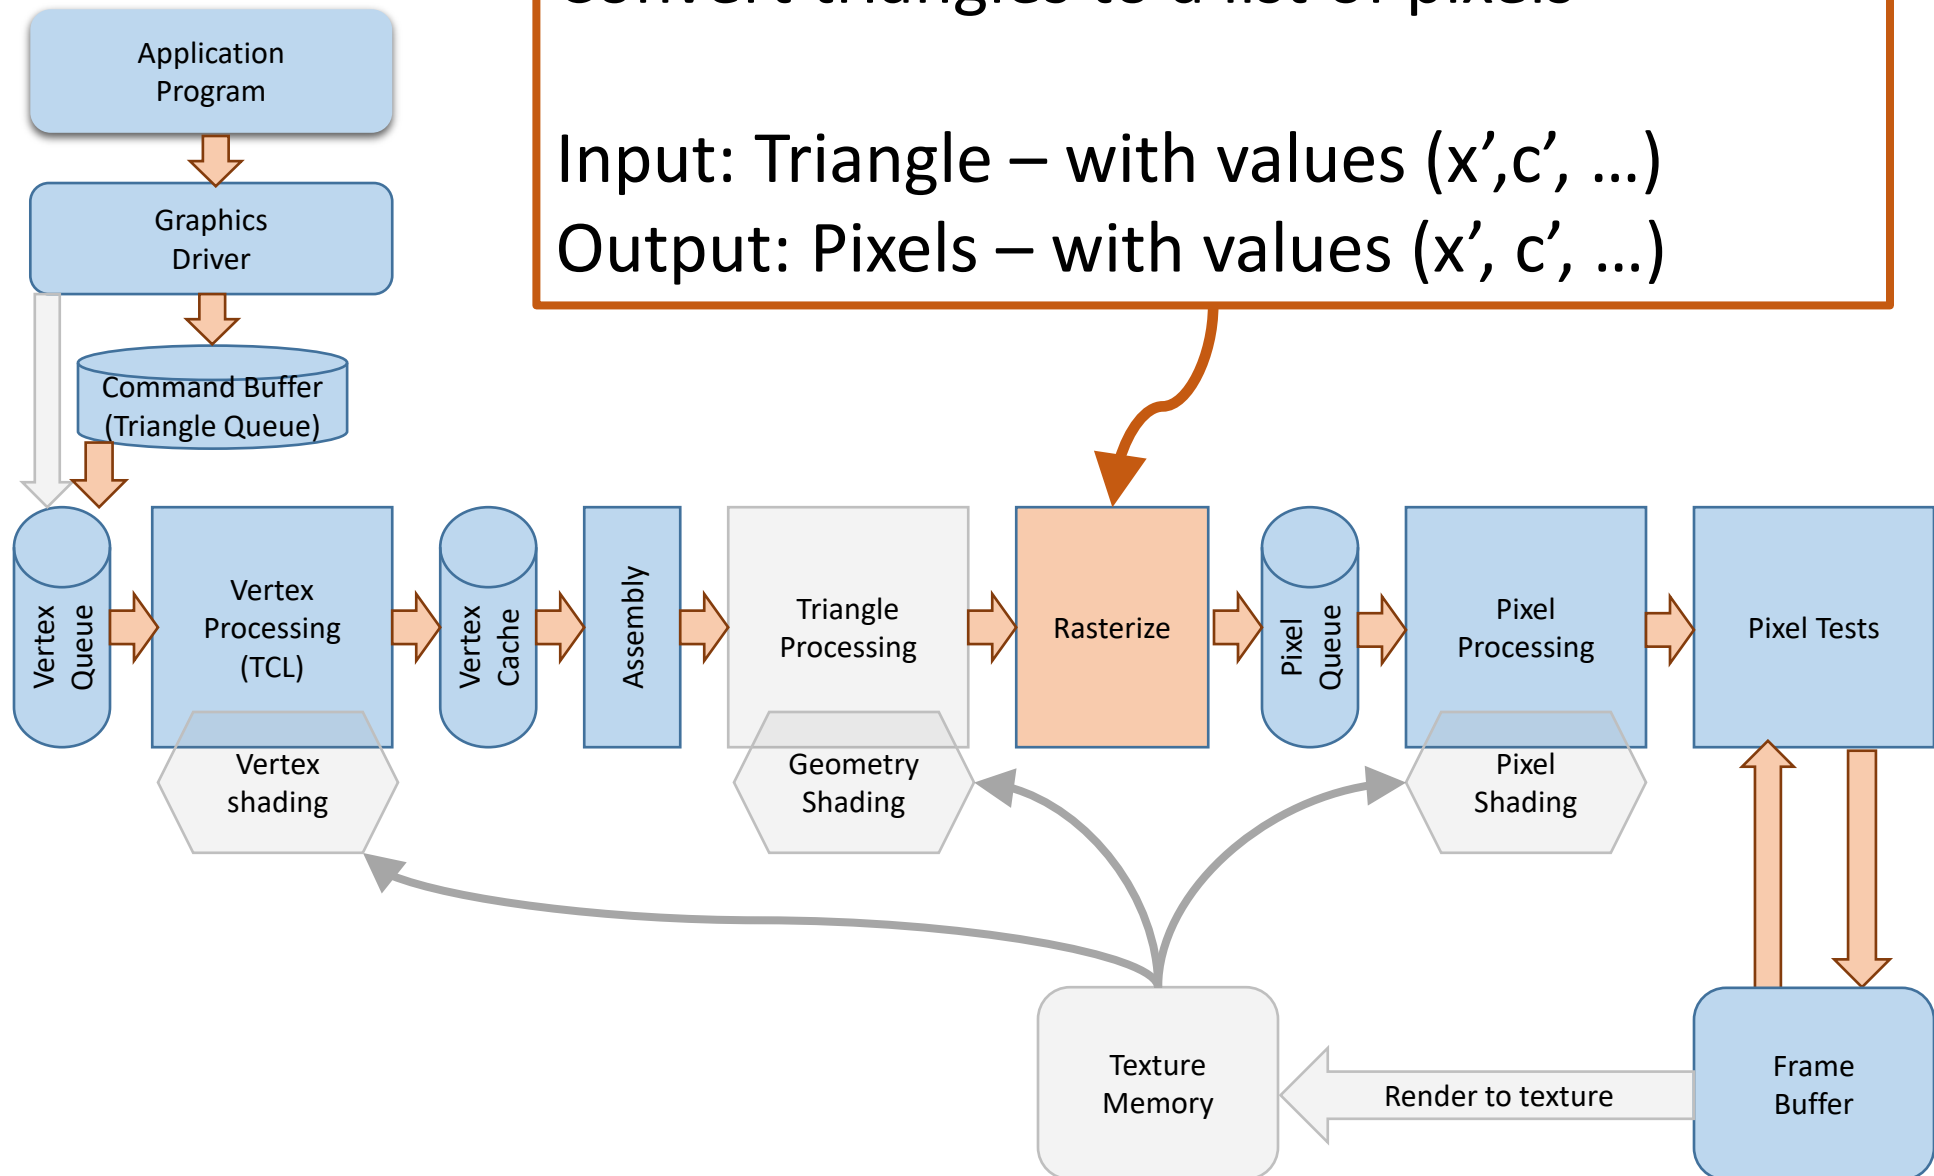


## Rasterizer:

Convert triangles to a list of pixels

Input: Triangle – with values ( $x'$ ,  $c'$ , ...)

Output: Pixels – with values ( $x'$ ,  $c'$ , ...)



# Pixels or Fragments

I am using the terms interchangeably  
(actually, today I am using pixel)

## Technically...

Pixel = a dot on the screen

Fragment = a dot on a triangle

might not become a pixel (fails z-test)

might only be part of a pixel



# Rasterization

Figure out which pixels a primitive “covers”

Turns primitives into pixels

Interpolate vertex values to interior pixels

# Barycentric Coordinates


Any point in the plane is a convex combination of the vertices of the triangle

$$P = \alpha A + \beta B + \gamma C$$

$$\alpha + \beta + \gamma = 1$$

Inside triangle

$$0 \leq \alpha, \beta, \gamma \leq 1$$



# Where do pixel values come from?

Each vertex has values

Each pixel comes from 3 vertices

Pixels interpolate their vertices' values

Barycentric interpolation

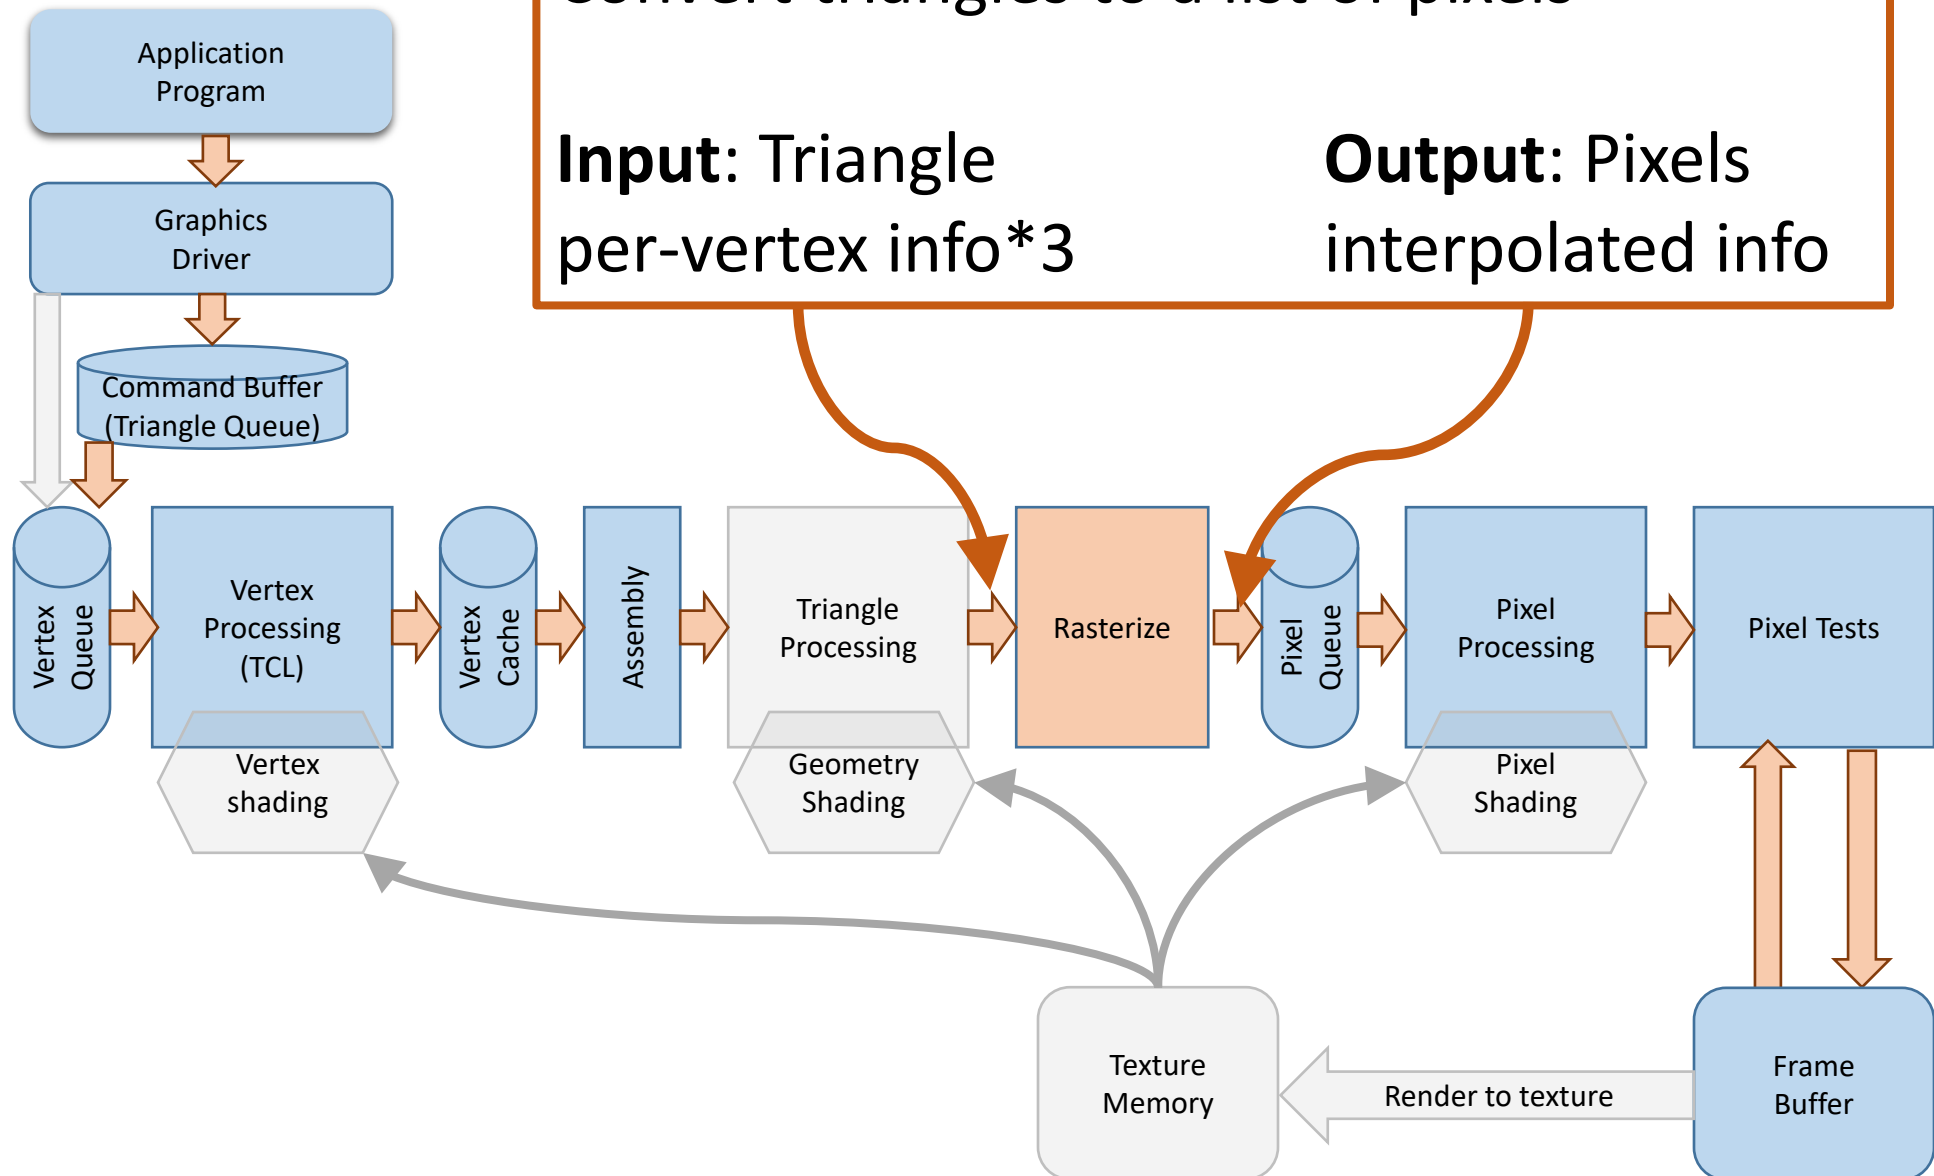
All values (in a pixel) are interpolated

Rasterizer:

Convert triangles to a list of pixels

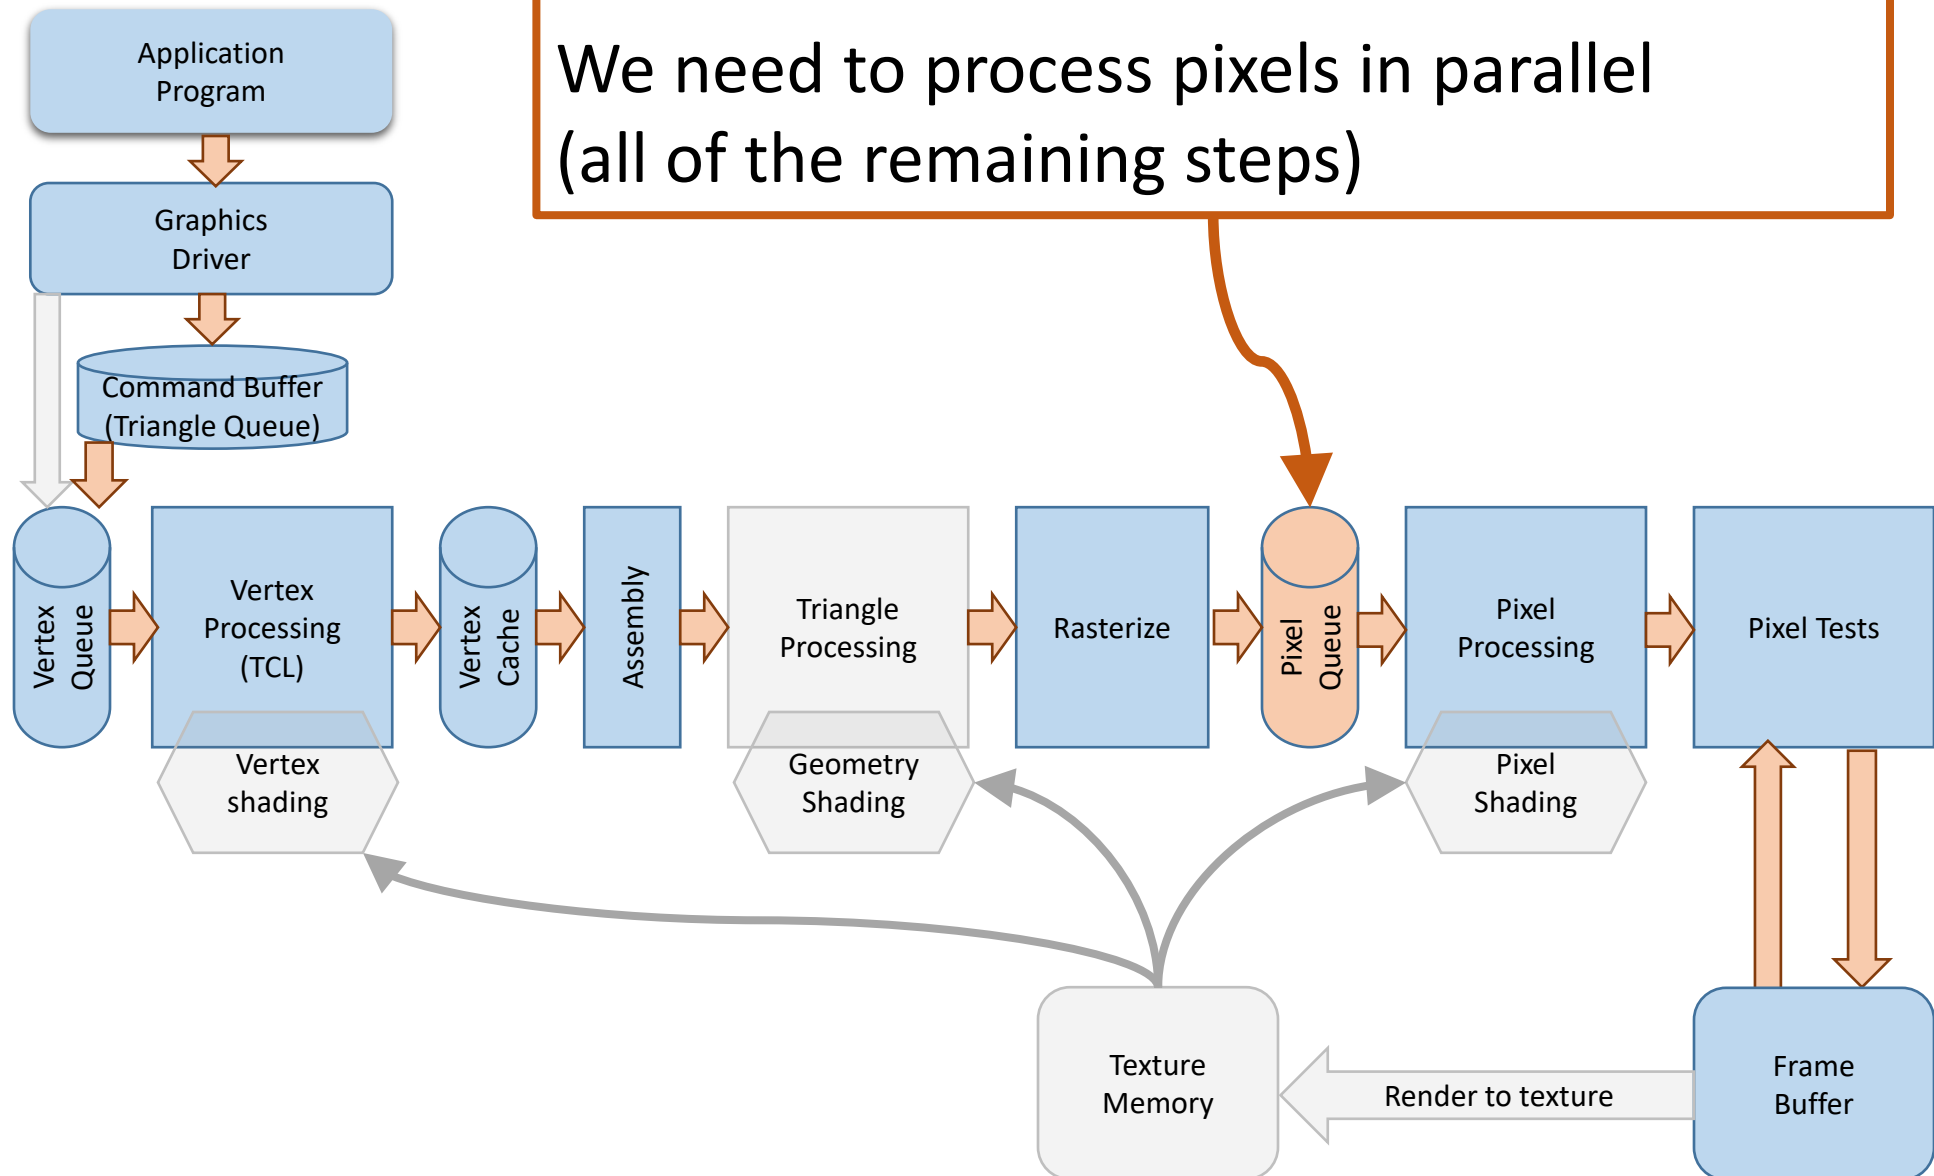
**Input:** Triangle  
per-vertex info\*3

**Output:** Pixels  
interpolated info



Each triangle can make lots of pixels

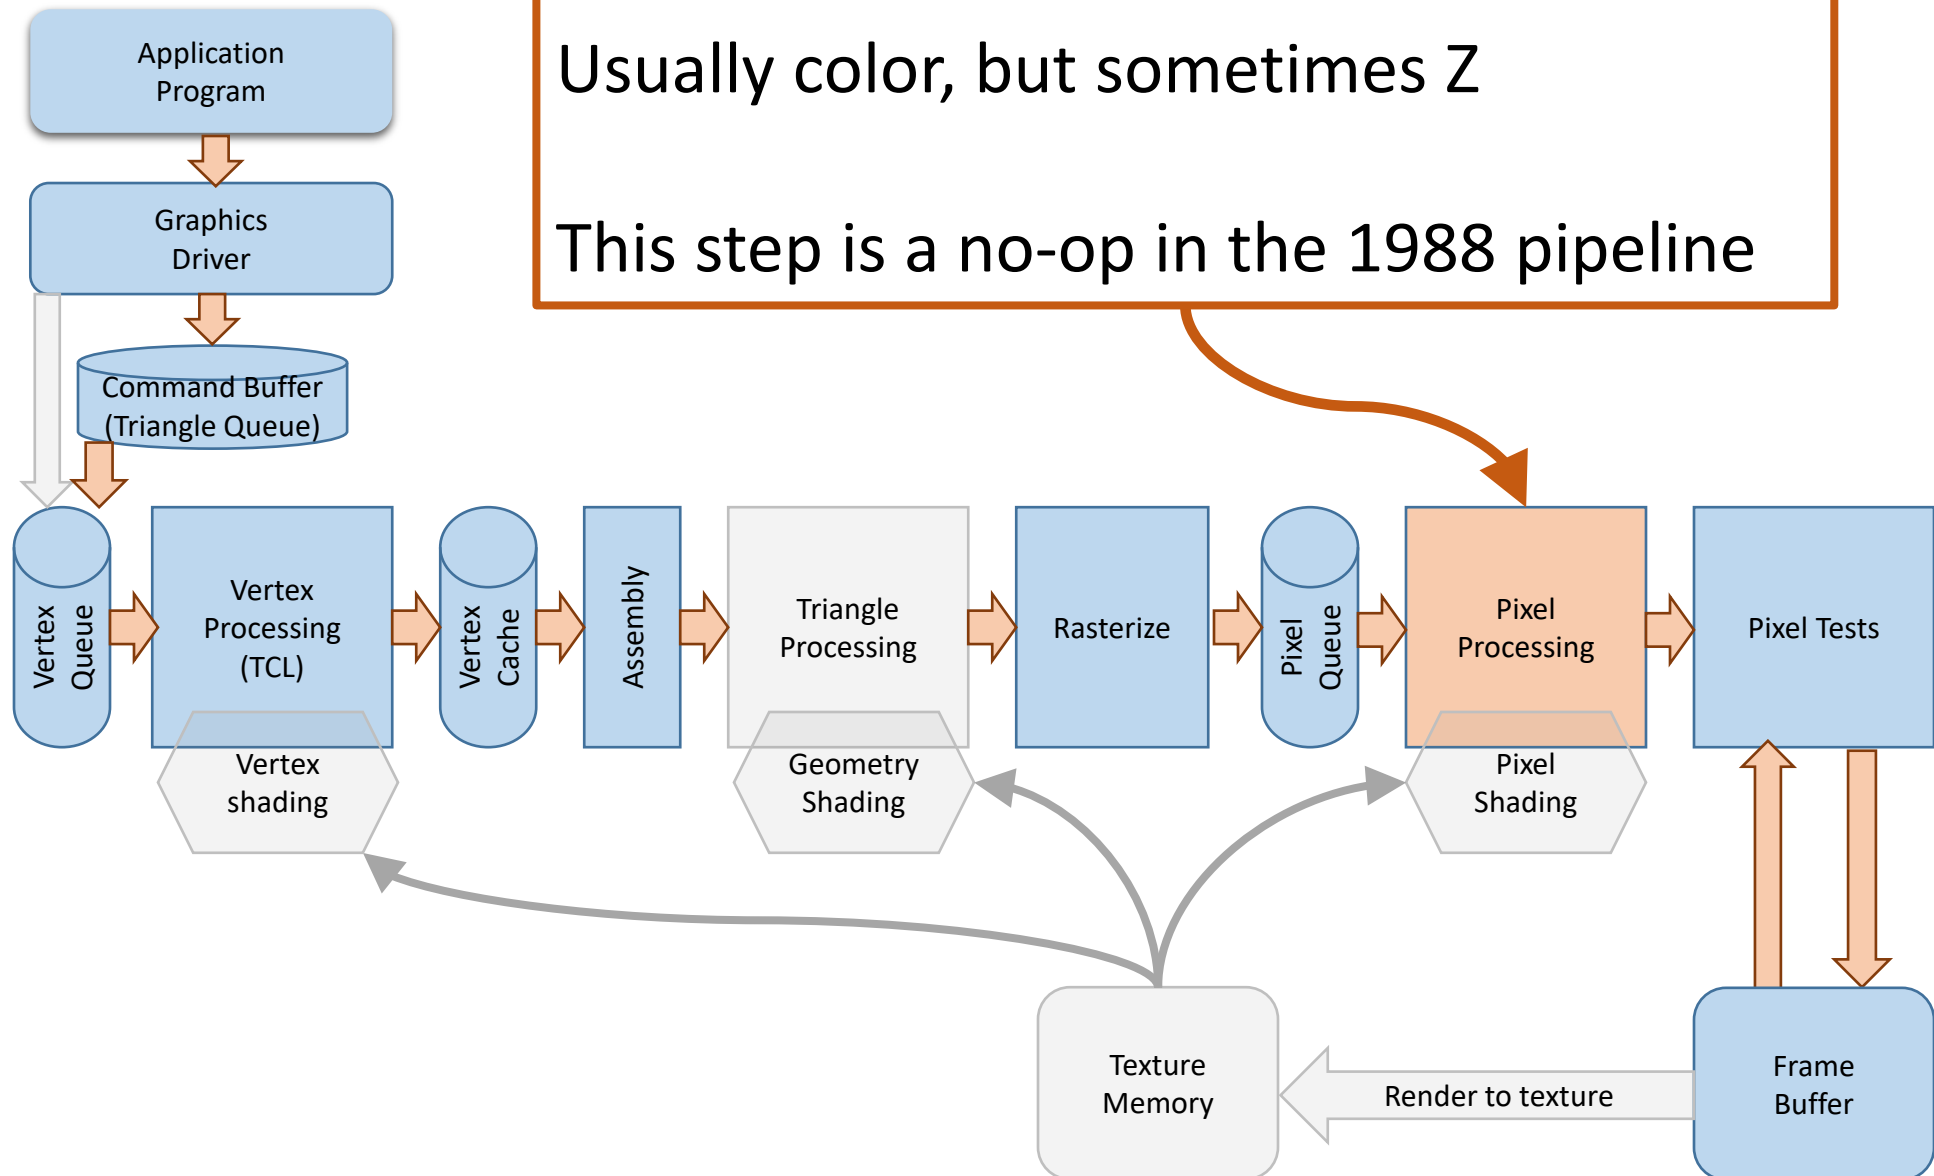
We need to process pixels in parallel  
(all of the remaining steps)



Process each pixel to get its final values

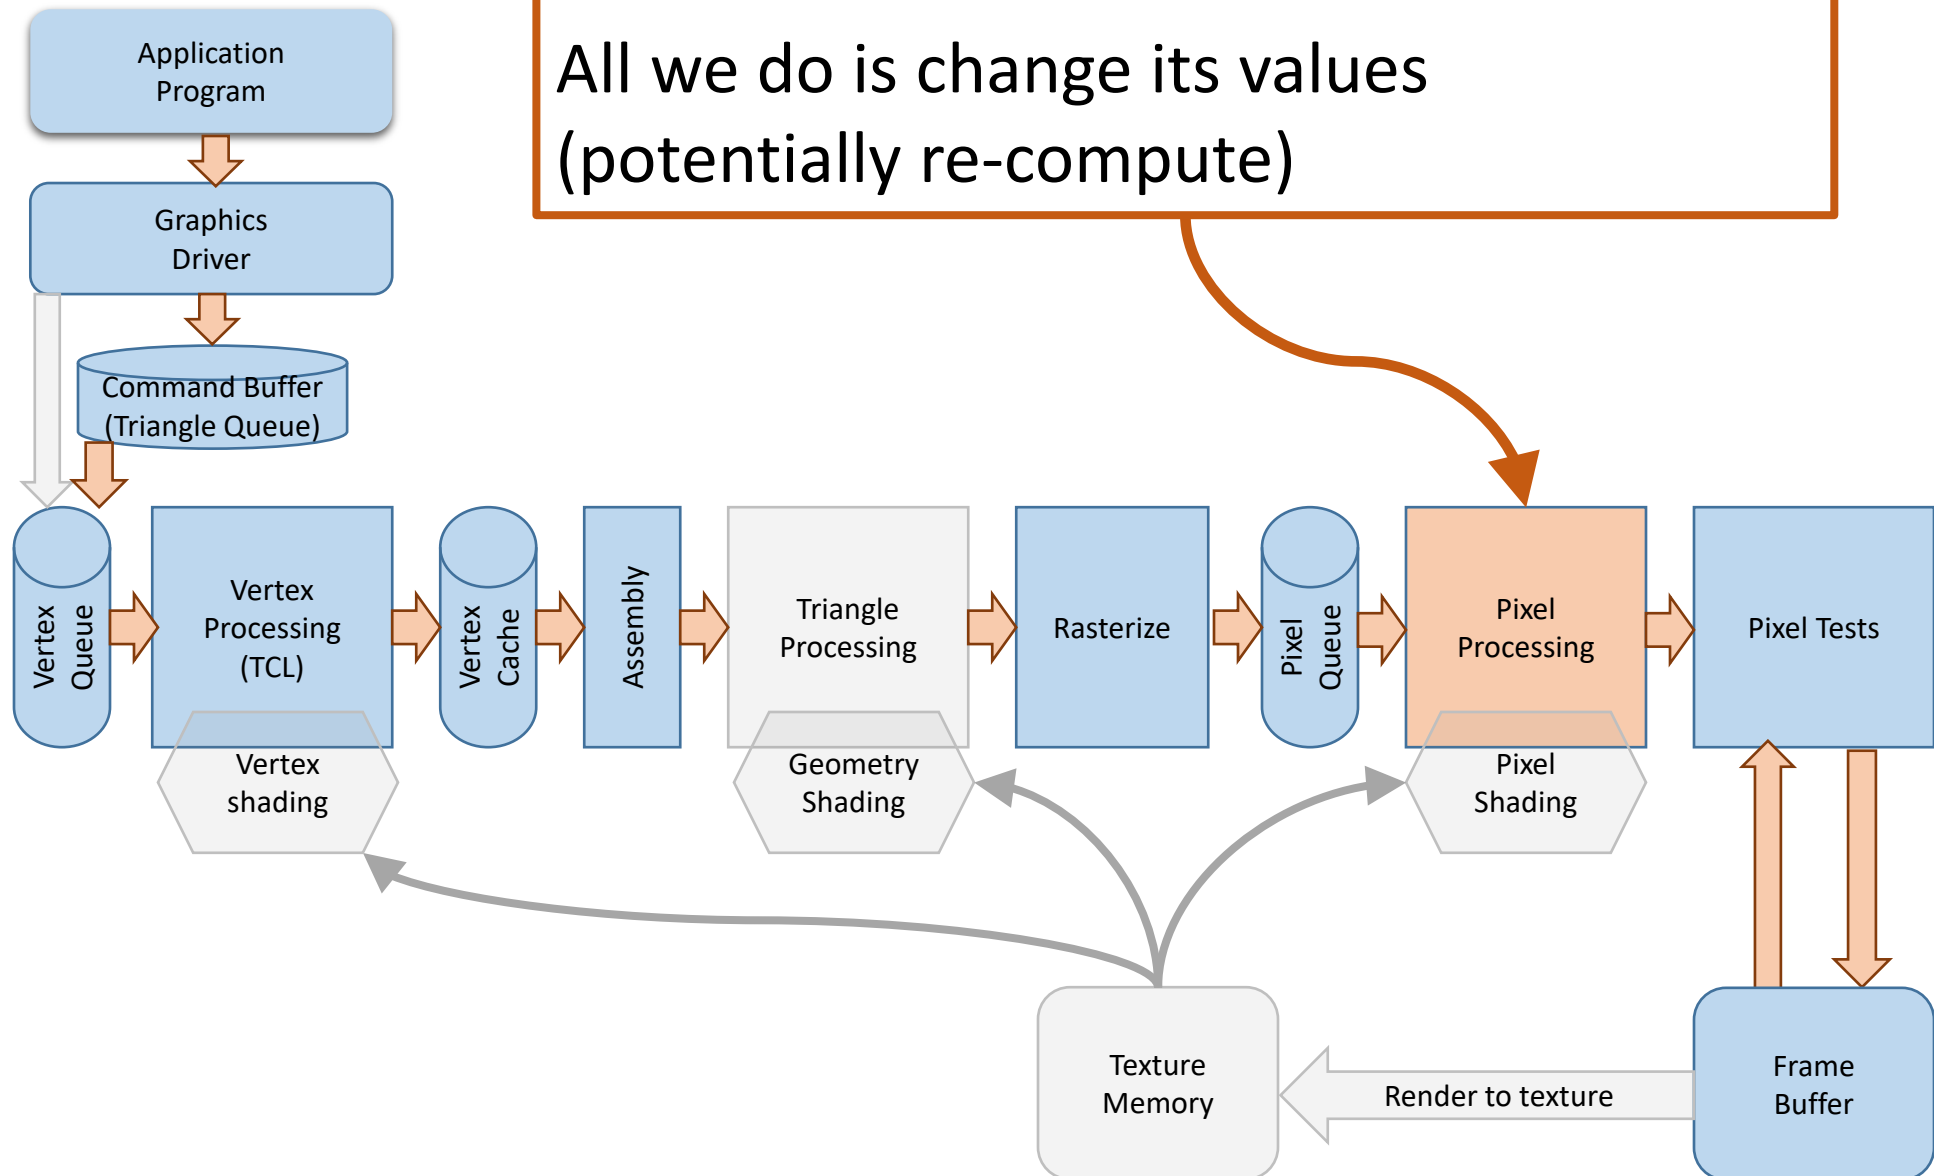
Usually color, but sometimes Z

This step is a no-op in the 1988 pipeline




Pixel in → Pixel out, each independent

All we do is change its values  
(potentially re-compute)



Coming attractions...

This step will get to be pretty exciting





# Pixel Processing Ground Rules

Pixels are independent

Pixel in  $\rightarrow$  Pixel out

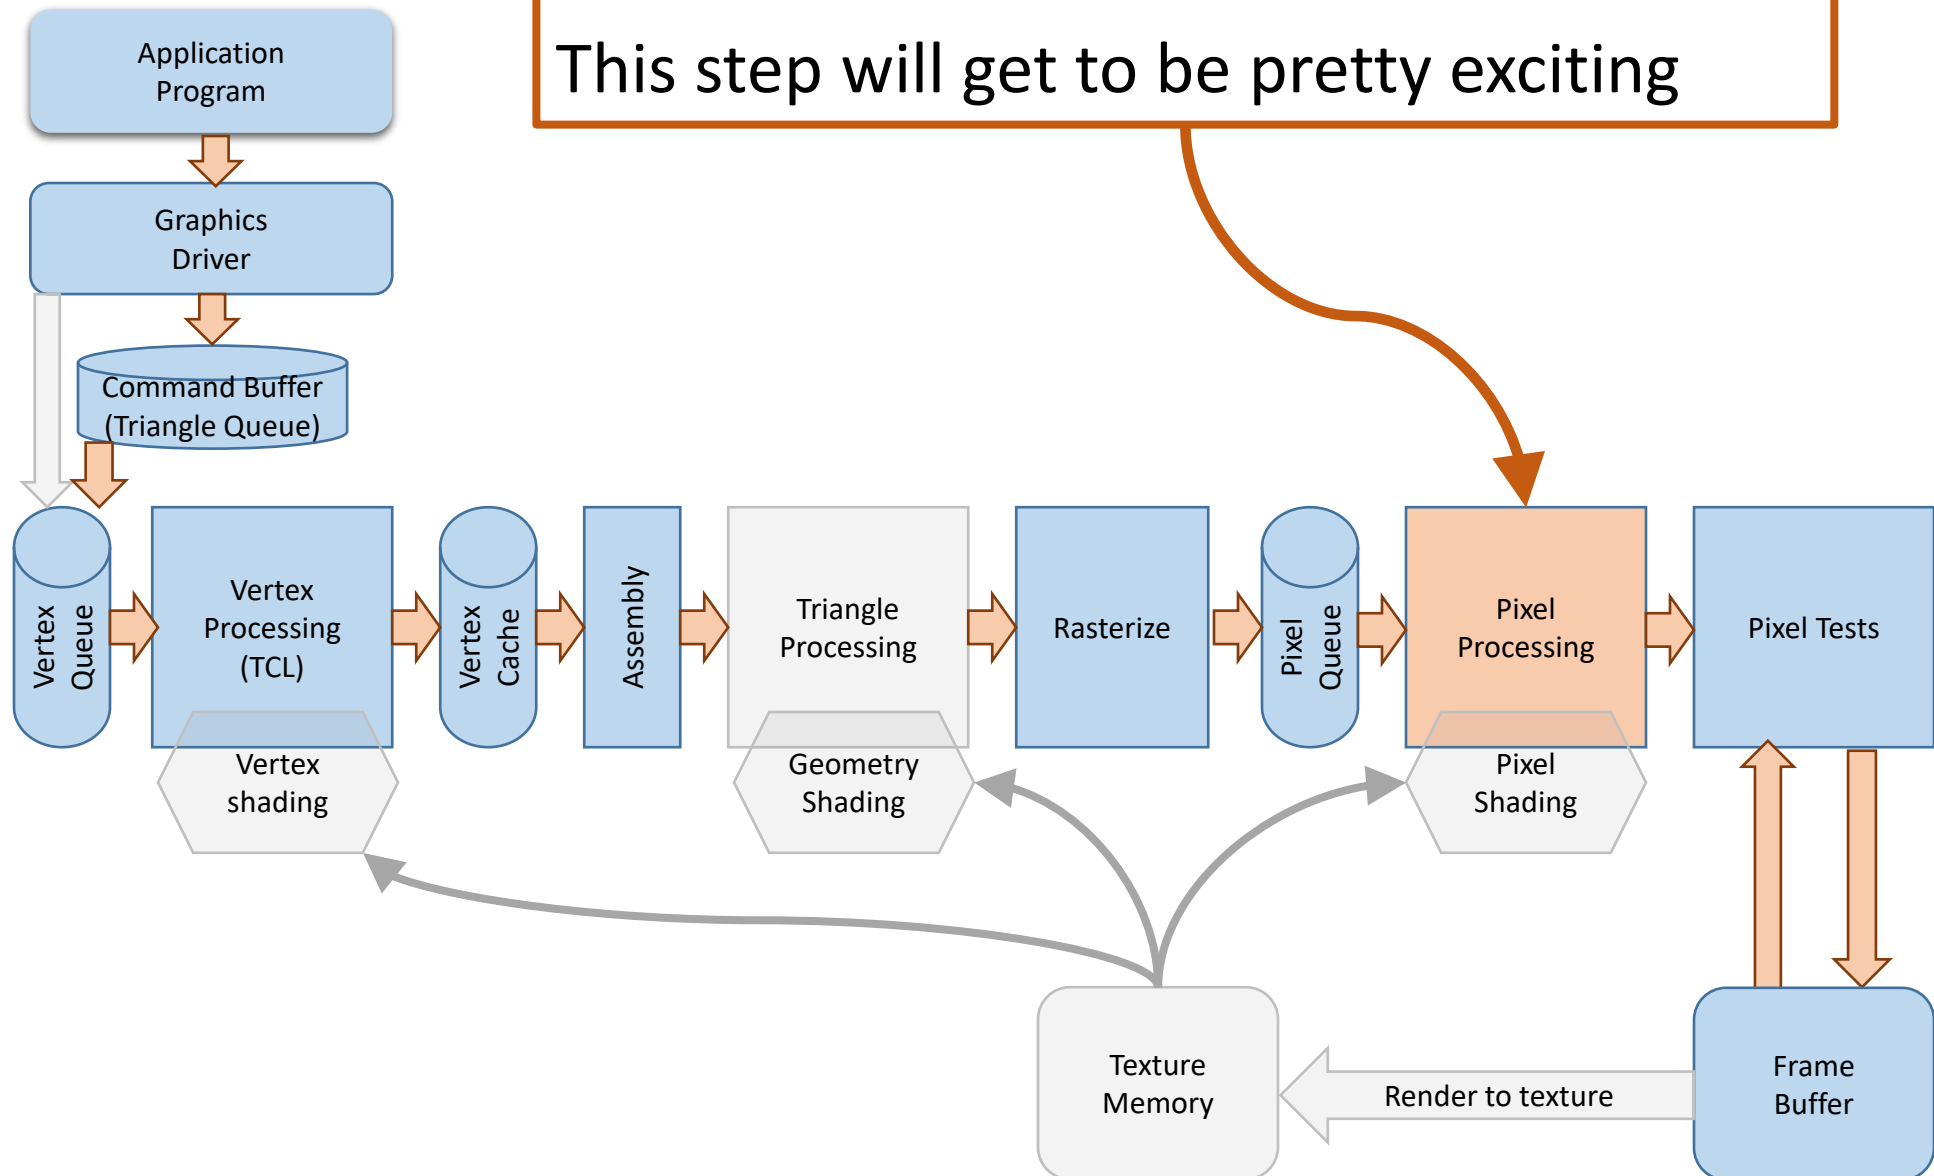
Changing its position (x,y) makes it a different pixel (so you can't)

Can change other values

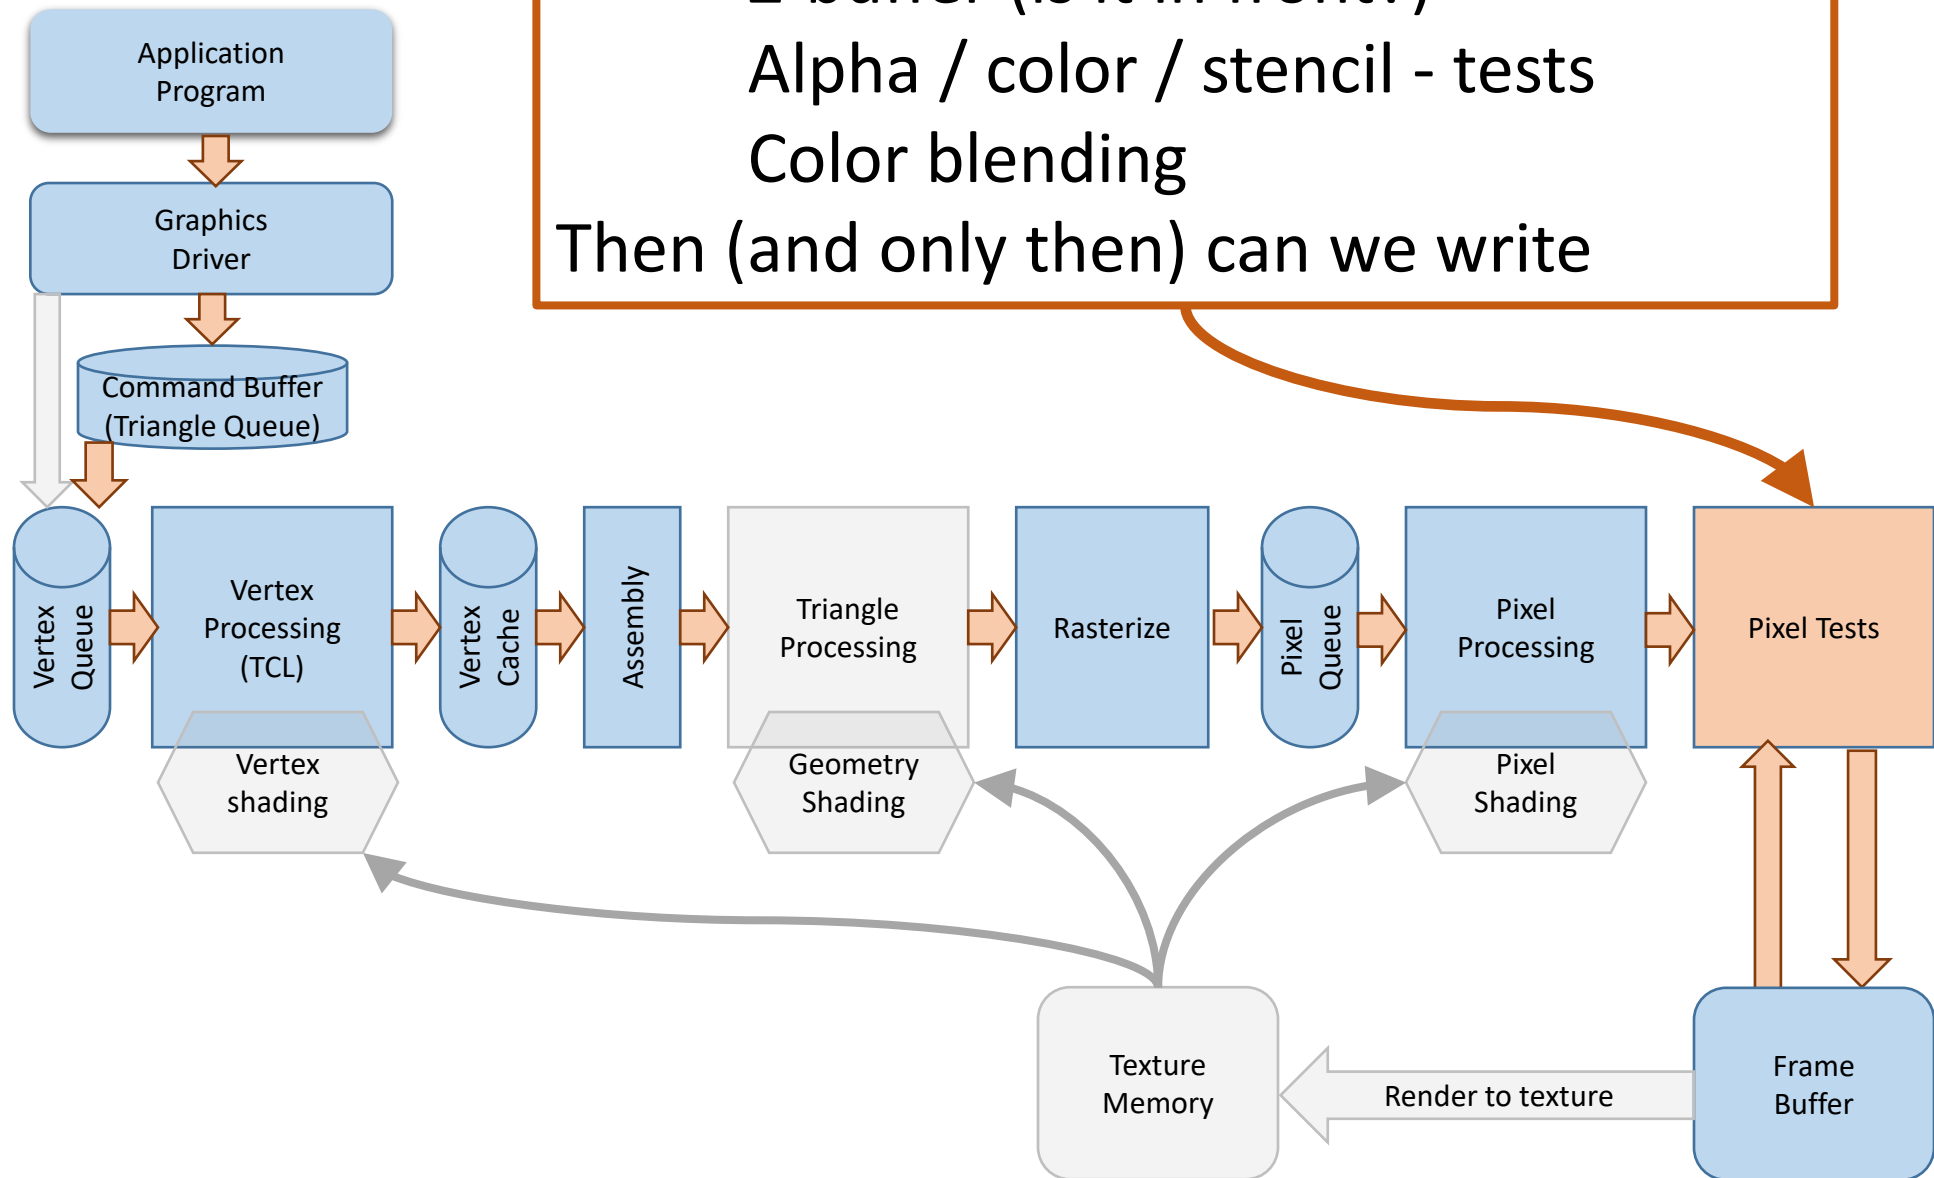
Or “reject”

Coming attractions...

This step will get to be pretty exciting



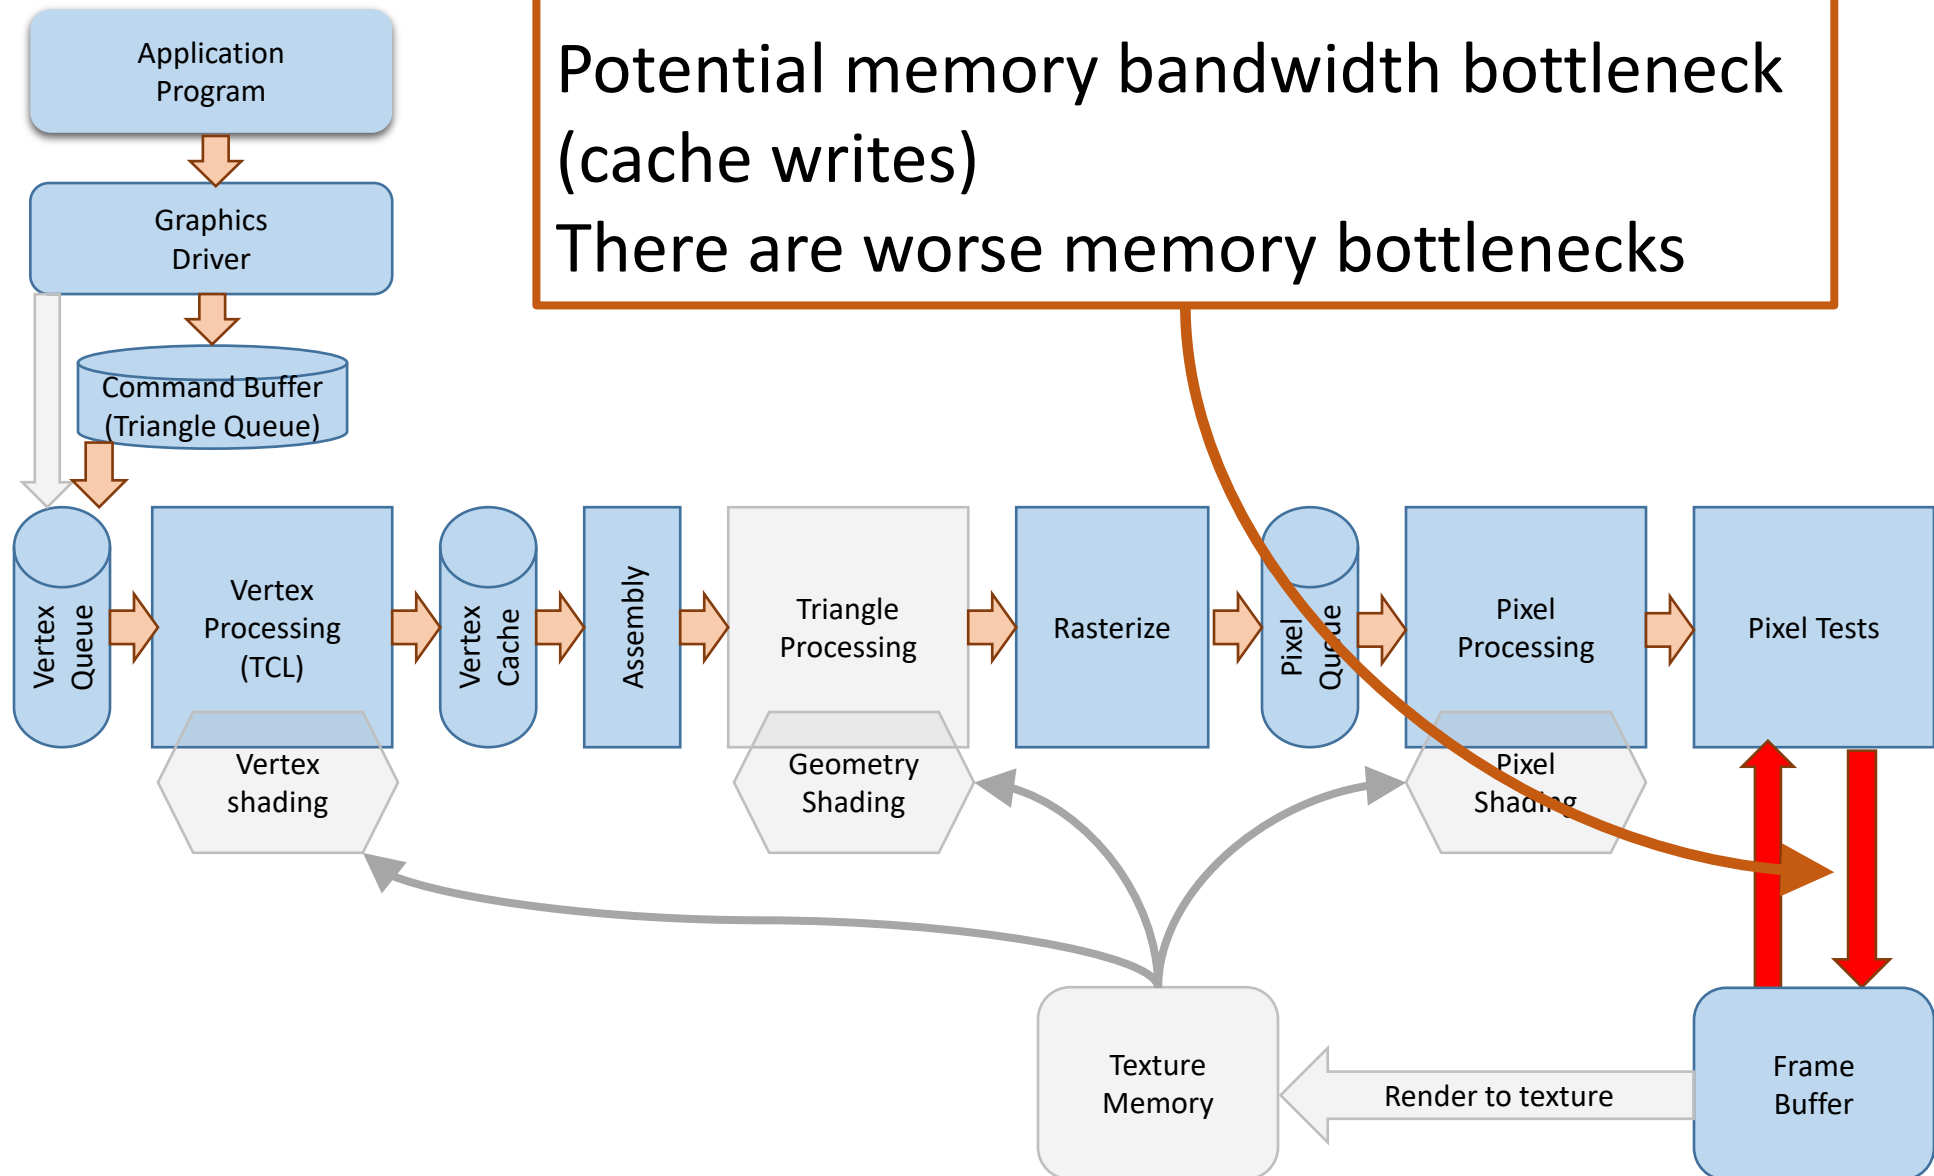
Consider the pixel and it's destination  
Z-buffer (is it in front?)  
Alpha / color / stencil - tests  
Color blending  
Then (and only then) can we write



Each pixel requires a read/write cycle

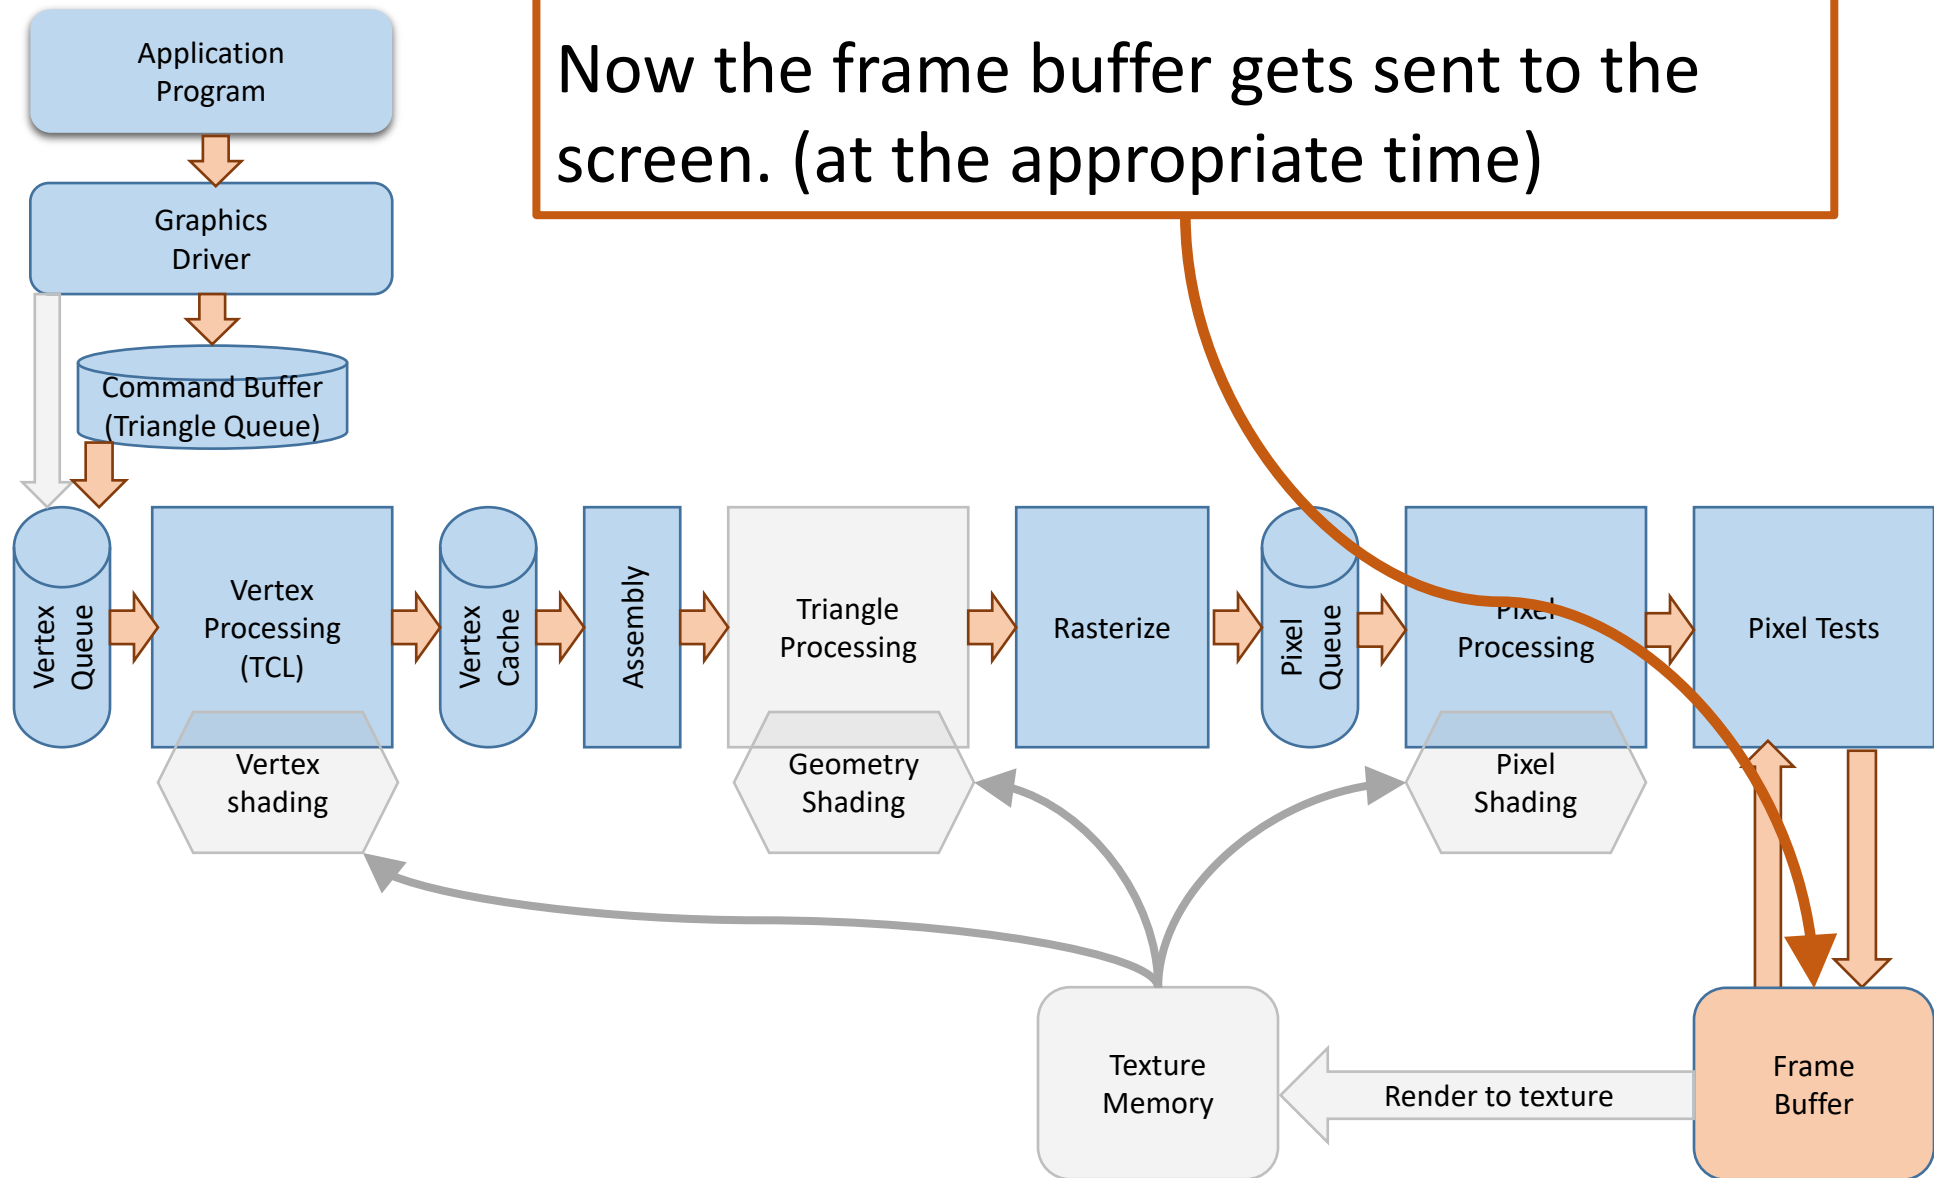
Potential memory bandwidth bottleneck  
(cache writes)

There are worse memory bottlenecks



We've made it!

Now the frame buffer gets sent to the screen. (at the appropriate time)

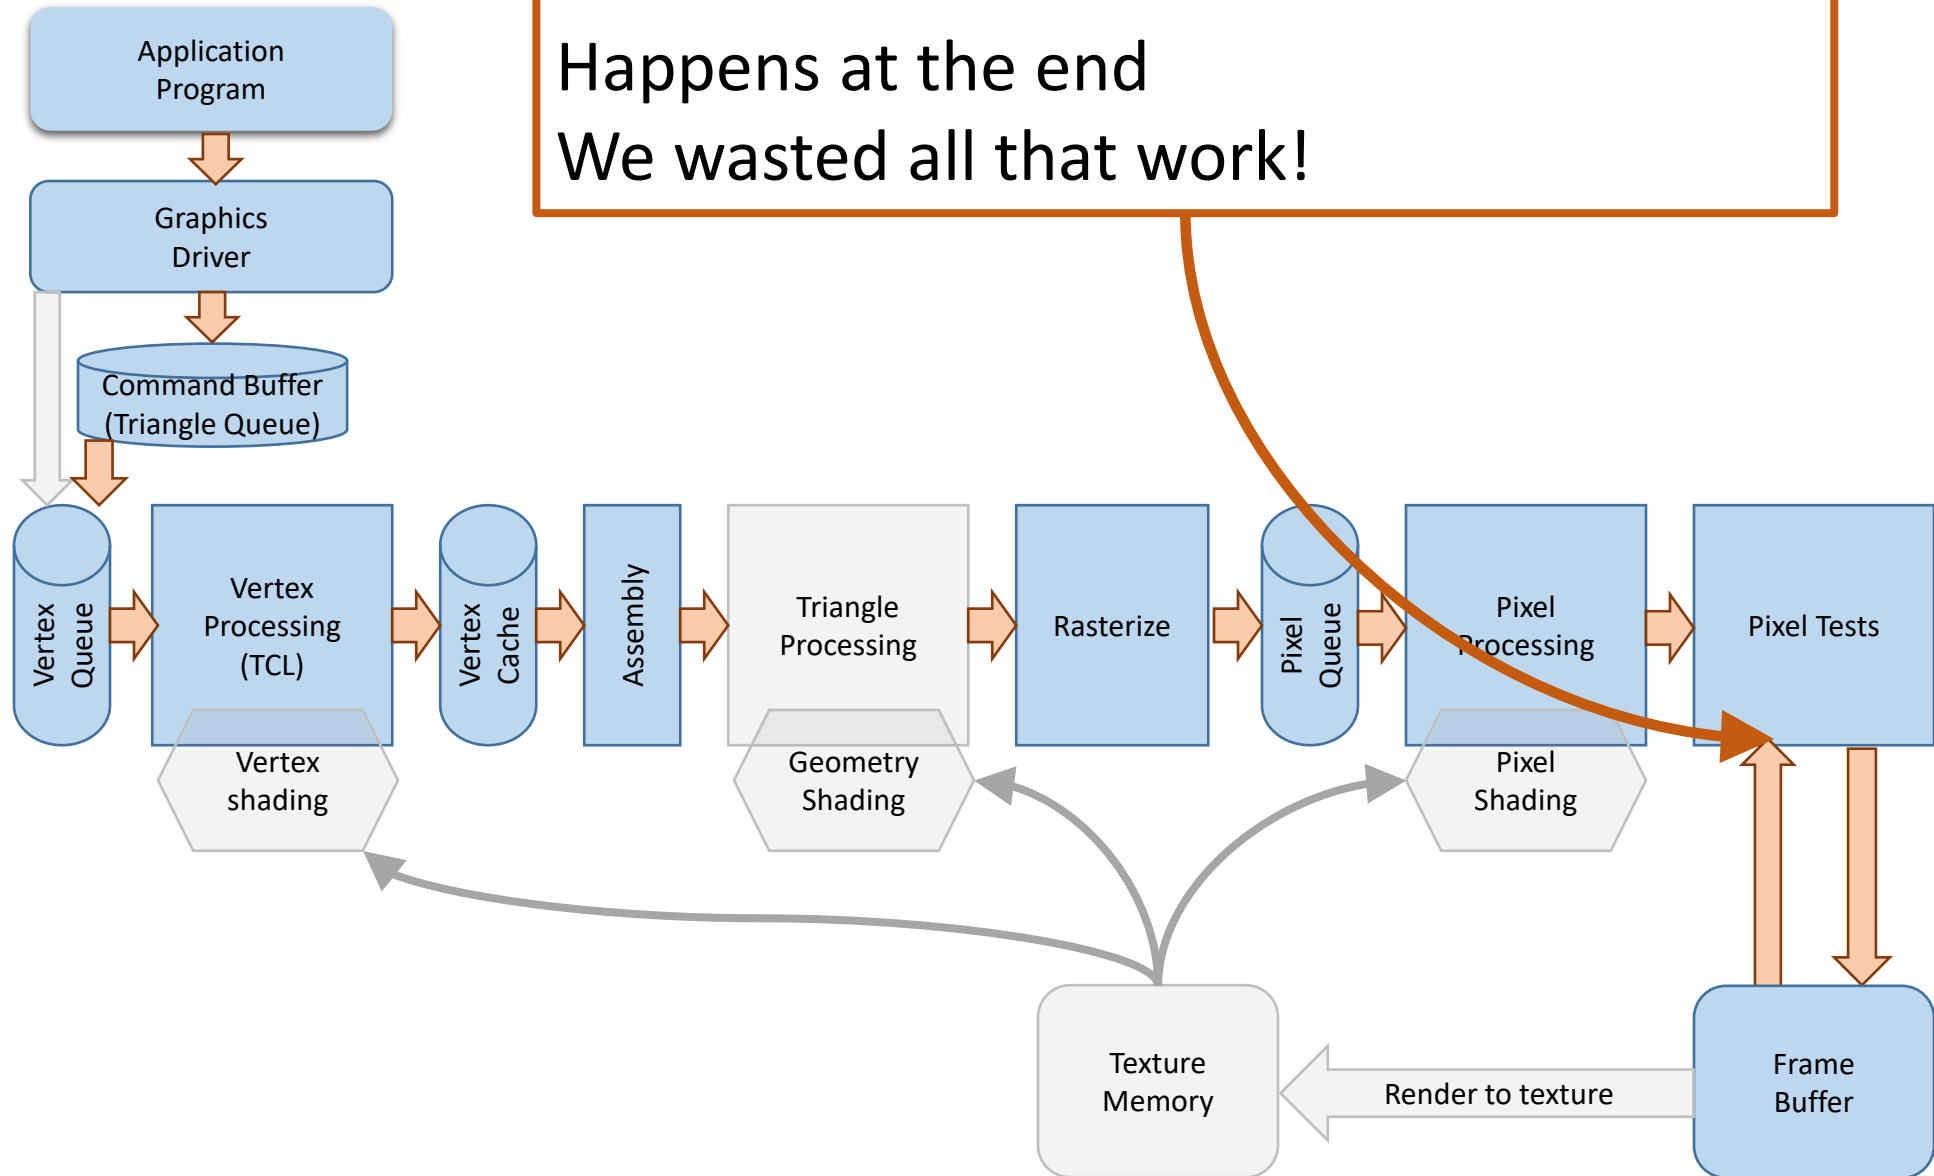


# What if we didn't make it...

Suppose the triangle's pixels are occluded  
Removed by the z-buffer

# Normal Z-Test

Happens at the end  
We wasted all that work!



# Summary: A Triangle's Journey

Graphics happens as a series of steps (Pipeline)

Transfer information about vertices from application

Process vertices

Assemble / Rasterize Triangles

Process fragments/pixels

Test/write pixel results



# What's next? Program these steps!

If you understand what the boxes do...

You can program them! (Shaders)

Efficient transfer of data from application

Program the Vertex Processor

Program the Fragment Processor