

# Lecture 22

# Shaders 2

---

# Now that we can we can write shaders...

---

Now we know how to write shaders...

What shaders do we write?

1. Lighting
2. Textures
3. Other tricks

# Lighting

---

THREE implements lighting very well (nice shaders)

Generally, we don't need to implement it ourselves

- if we want to understand lighting
- if we want to do something non-standard
- if we want to combine it with some other shading

# Where to implement lighting?

---

## Vertex Shader

Compute lighting at every vertex

Interpolate colors across triangle

Per-Vertex Lighting

Gouraud Shading

Used in the old days

## Fragment Shader

Interpolate normals and surface colors

Compute lighting at every pixel

Per-pixel (fragment) shading

Hardware is fast enough to do it

# What Lighting to Implement?

---

We skipped talking about lighting...

# Lighting in the Real-World

---

# Lighting: The General Idea

---

- Light starts at a source
- Interacts with objects
- Arrives at a sensor

# Lighting: the simple local model

---

A photon arrives from some direction - where does it go?

- absorbed?
- perfect mirror bounce?
- micro-geometry
- subsurface effects
- depends on **material and direction**

# Lighting: Local vs. Global

---

## Local Lighting

what happens at **specific point**  
each point considered individually  
assume light comes from somewhere  
(direct from source)

## Global Lighting

considers how points interact  
entire scene considered together  
considers how light **transports**

# Local Lighting

---

we have to start somewhere...

1. consider each point individually
2. assume light comes from sources (not other objects)
3. only consider how material reflects light

this means no global effects:

- no shadows
- no reflections
- no spill / bleed

we'll get these back later

# A Simple Local Lighting Model

---

- Assume an "ideal" material
- Make a model
- Tweak the model
- Combine with other models to mix effects

# "The" Lighting Model (Historic)

---

- Emission (things give off light)
- Ambient (light from nowhere)
- Specular (direct reflection)
- Diffuse (rough reflection)

## Less Historic

- ditch the hack pieces (emission, ambient)
- use less idealized models for specular and diffuse

# Preliminaries

---

- Consider a point - with local geometry
  - surface normals
- Sum over all lights and reflectance types
  - close to how lighting really works
- Colors by doing things component-wise
  - can work in RGB or other color spaces

# Diffuse Materials

---

Consider a very rough surface - chalk

# Lambertian Materials

---

- scatters light in all directions
- doesn't matter what direction you look from

# The direction of the light does matter

---

Consider a direction, and a small piece of surface

# Diffuse Reflection

---

$$r_{\text{diffuse}} = \hat{\mathbf{n}} \cdot \hat{\mathbf{l}}$$

where:

- $r_{\text{diffuse}}$  = amount of diffuse reflection
- $\hat{\mathbf{n}}$  - unit surface normal
- $\hat{\mathbf{l}}$  - unit vector to light source

# using this...

---

$$\text{color} = (\hat{\mathbf{n}} \cdot \hat{\mathbf{l}}) \quad \mathbf{c}_{\text{light}} \quad \mathbf{c}_d$$

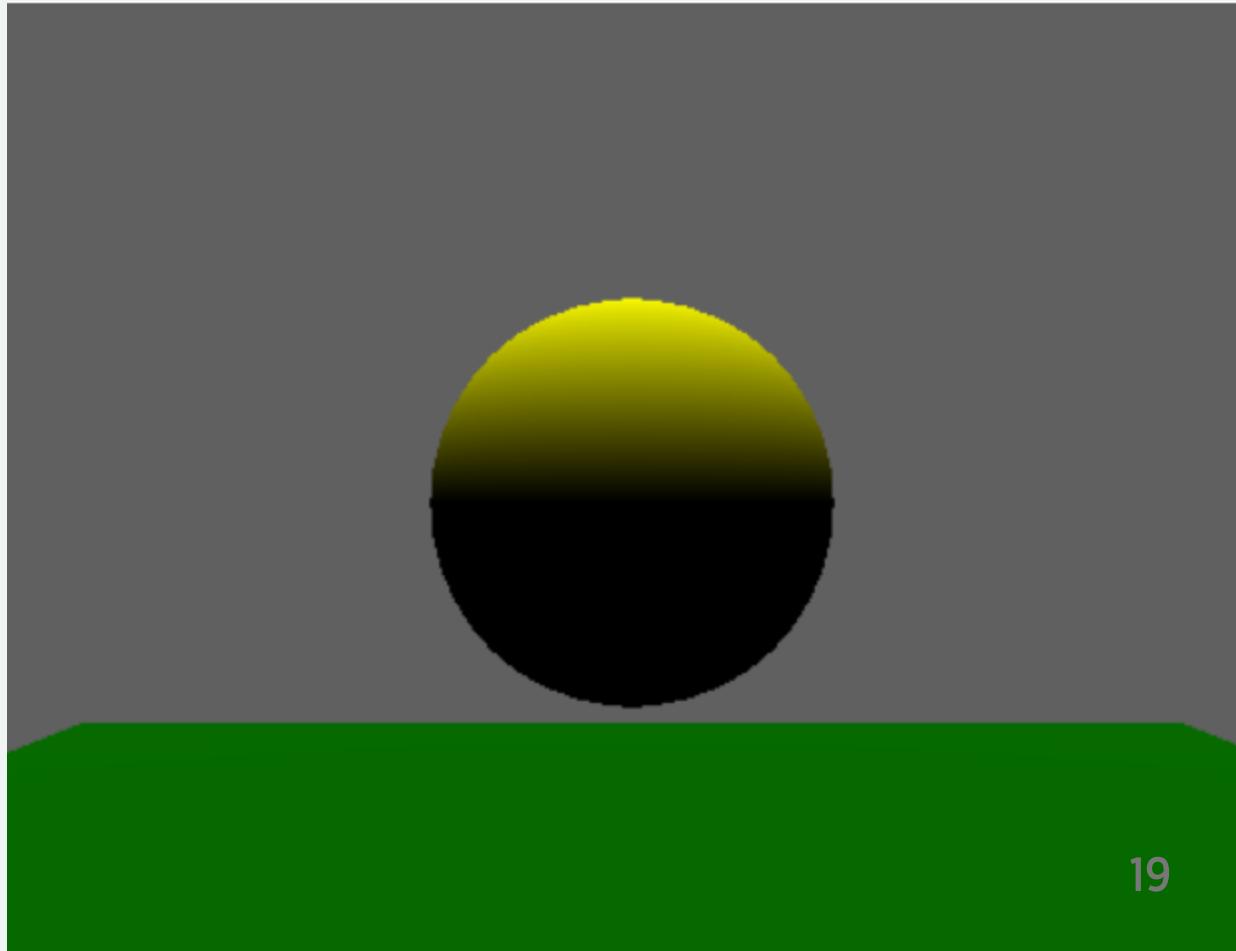
where

- $\hat{\mathbf{n}}$  - unit surface normal
- $\hat{\mathbf{l}}$  - unit vector to light source
- $\mathbf{c}_{\text{light}}$  - color/intensity of light
- $\mathbf{c}_d$  - color of the material (diffuse reflectance)

# looking at a sphere...

---

why a sphere is a good "test probe"



# Shiny Things

---

We are reflecting the lights (for now)

**Specular** reflection

# A Perfect Mirror

---

- Light direction and eye position matter
- The eye needs to be in the exact correct position
- $\hat{e}$  (eye vector) and  $\hat{r}$  (reflection vector)

# A Realistic (Imperfect) Mirror

---

**Phong model - it's a hack!**

Gradual fall off as we get away from the optimal direction

$$\hat{e} \cdot \hat{r}$$

$$\text{keep} > 0$$

# Phong Model

---

Raise to the "power" of "shininess" (phong exponent)

$$r_{specular} = (\hat{\mathbf{e}} \cdot \hat{\mathbf{r}})^p$$

where

- $\hat{\mathbf{e}}$  is the eye vector
- $\hat{\mathbf{r}}$  is the reflection vector
- $p$  is the "shininess" (material property), phone exponent
- $r_{specular}$  is the amount of specular reflection

Shinier = more like a perfect mirror

# An Alternate Way to Compute

---

$$r_{specular} = (\hat{\mathbf{n}} \cdot \hat{\mathbf{h}})^p$$

where

- $\hat{\mathbf{n}}$  is the normal vector
- $\hat{\mathbf{h}}$  is the half-way vector (between  $\hat{\mathbf{l}}$  and  $\hat{\mathbf{e}}$ )
- $p$  is the "shininess" (material property), phone exponent
- $r_{specular}$  is the amount of specular reflection

# Using This

---

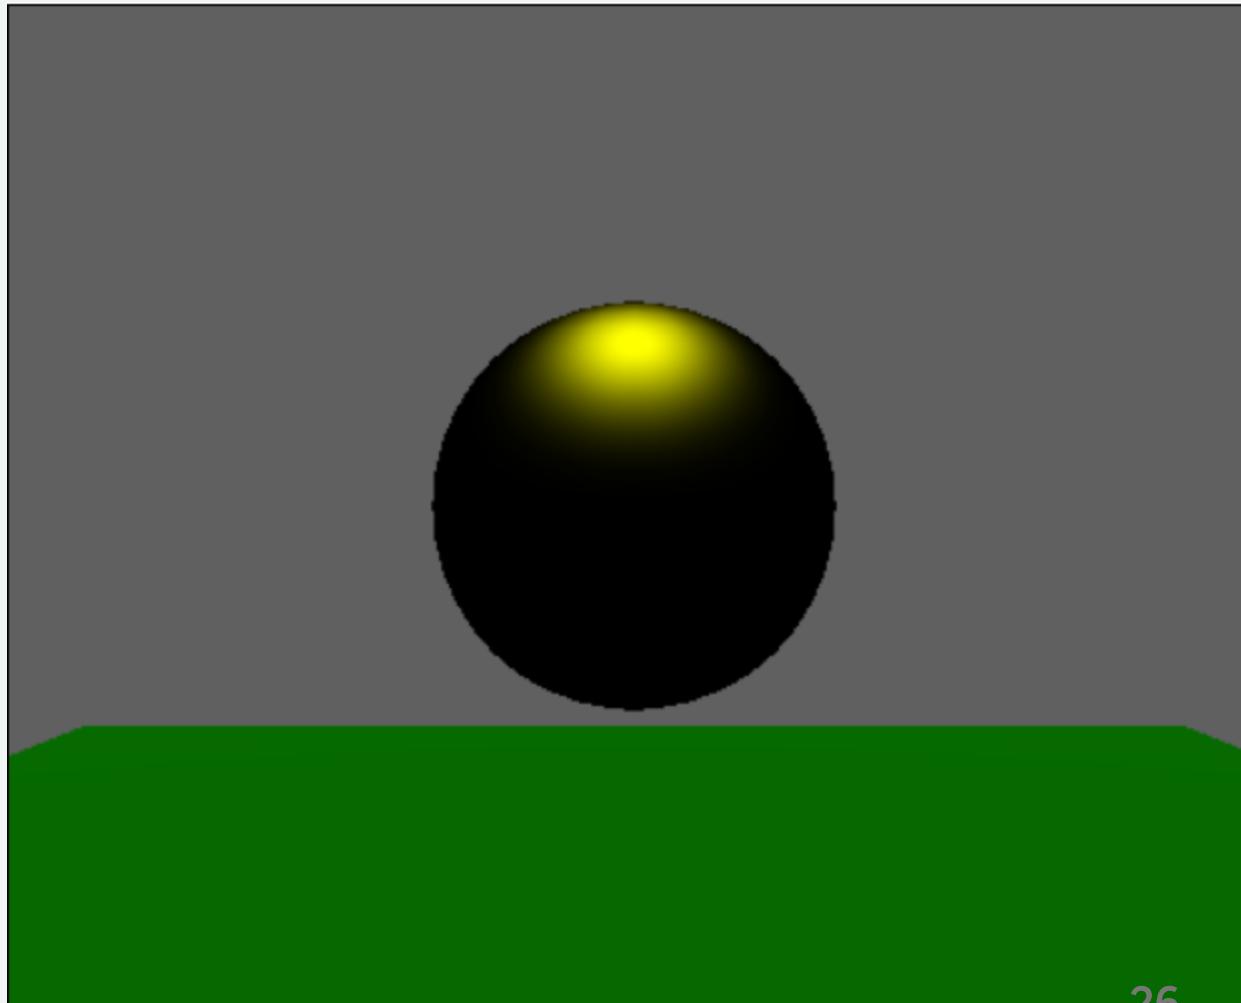
$$\text{color} = (\hat{\mathbf{n}} \cdot \hat{\mathbf{l}})^p \quad \mathbf{c}_{\text{light}} \quad \mathbf{c}_s$$

where

- $\hat{\mathbf{n}}$  - unit surface normal
- $\hat{\mathbf{l}}$  - unit vector to light source
- $\mathbf{c}_{\text{light}}$  - color/intensity of light
- $\mathbf{c}_s$  - color of the material (specular reflectance)
- $p$  - shininess of the material

# Demo

---



# Putting it together...

---

Sum of different lighting components:

- emissive (objects give off light)
- ambient (generic light from all directions)
- diffuse (add for each light source)
- specular (add for each light source)

# A Material is defined by:

---

- an emissive color  $\mathbf{c}_e$
- an ambient color  $\mathbf{c}_a$
- a diffuse color  $\mathbf{c}_d$
- a specular color  $\mathbf{c}_s$
- a specular exponent (shininess)  $p$

# The Phong Lighting Model

---

$$\text{color} = \mathbf{c}_e + \mathbf{c}_a \mathbf{l}_a + \sum_{l \in lights} \left( (\hat{\mathbf{n}} \cdot \hat{\mathbf{l}}) \mathbf{c}_{\text{light}} \mathbf{c}_d + (\hat{\mathbf{n}} \cdot \hat{\mathbf{h}})^p \mathbf{c}_{\text{light}} \mathbf{c}_s \right)$$

which uses:

- the surface geometry  $\hat{\mathbf{n}}$
- the light direction  $\hat{\mathbf{l}}$
- the eye direction (combined with light direction into the half-vector) ( $\hat{\mathbf{h}}$ )
- the light color (intensity)  $\mathbf{l}_a$ ) and  $\mathbf{c}_{\text{light}}$
- the surface properties  $\mathbf{c}_e$ ,  $\mathbf{c}_a$ ,  $\mathbf{c}_d$ ,  $\mathbf{c}_s$  and  $p$

# Using this in THREE.js

---

Use the `MeshPhongMaterial`

- specify 3 colors (no separate ambient)
  - `color` is diffuse and ambient color
  - `specular` is specular color
  - `emissive` is emissive color

# Designing Materials

---

*warning: Phong is a hack, choosing parameters that look good is an art*

Amount of Shininess:

- how bright is the shininess ( $c_s$ , relative to  $c_d$ )
- how focused is the shininess ( $p$ )

Color of Shininess

- what color is the shininess (plastic has white highlights, metal is colored)

# Alternative Workflow

---

Could choose a different workflow:

- **metalness** - behaves like metal or plastic
  - metal specular keeps color, plastic, specular is white
- **roughness** - dull or shiny
  - amount of specularity, size of specularity

Or...

- use a less hacky model of local lighting!

# The lighting to implement (Phong)

---

$$\text{color} = \mathbf{c}_e + \mathbf{c}_a \mathbf{l}_a + \sum_{l \in lights} \left( (\hat{\mathbf{n}} \cdot \hat{\mathbf{l}}) \mathbf{c}_{\text{light}} \mathbf{c}_d + (\hat{\mathbf{n}} \cdot \hat{\mathbf{h}})^p \mathbf{c}_{\text{light}} \mathbf{c}_s \right)$$

which uses:

- the surface geometry  $\hat{\mathbf{n}}$
- the light direction  $\hat{\mathbf{l}}$
- the eye direction (combined with light direction into the half-vector) ( $\hat{\mathbf{h}}$ )
- the light color (intensity)  $\mathbf{l}_a$ ) and  $\mathbf{c}_{\text{light}}$
- the surface properties  $\mathbf{c}_e$ ,  $\mathbf{c}_a$ ,  $\mathbf{c}_d$ ,  $\mathbf{c}_s$  and  $p$

# A Material is defined by:

---

- an emissive color  $\mathbf{c}_e$
- an ambient color  $\mathbf{c}_a$
- a diffuse color  $\mathbf{c}_d$
- a specular color  $\mathbf{c}_s$
- a specular exponent (shininess)  $p$

Get color from:

- material (as a uniform)
- the vertices (interpolate attributes)
- a texture map (later)

# Diffuse

---

$$\text{color} = (\hat{\mathbf{n}} \cdot \hat{\mathbf{l}}) \quad \mathbf{c}_{\text{light}} \quad \mathbf{c}_d$$

where

- $\hat{\mathbf{n}}$  - unit surface normal
- $\hat{\mathbf{l}}$  - unit vector to light source
- $\mathbf{c}_{\text{light}}$  - color/intensity of light
- $\mathbf{c}_d$  - color of the material (diffuse reflectance)

# Diffuse and Ambient

---

$$\text{color} = \left( (\hat{\mathbf{n}} \cdot \hat{\mathbf{l}}) \mathbf{c}_{\text{light}} + \mathbf{c}_{\text{ambient}} \right) \mathbf{c}_d$$

where

- $\hat{\mathbf{n}}$  - unit surface normal
- $\hat{\mathbf{l}}$  - unit vector to light source
- $\mathbf{c}_{\text{light}}$  - color/intensity of light
- $\mathbf{c}_{\text{ambient}}$  - color/intensity of ambient light
- $\mathbf{c}_d$  - color of the material (diffuse reflectance)

# Where do those values come from?

---

$\hat{n}$  - unit surface normal

$\hat{l}$  - unit vector to light source

$c_{\text{light}}$  - color/intensity of light

$c_{\text{ambient}}$  - color of ambient light

$c_d$  - **color of material**

## Surface color

uniform for material?

varying (copy from attribute)

read from texture

# Where do those values come from?

---

$\hat{n}$  - unit surface normal

$\hat{l}$  - unit vector to light source

$c_{\text{light}}$  - **color/intensity of light**

$c_{\text{ambient}}$  - **color of ambient light**

$c_d$  - color of material

## Light Properties

Constants?

Uniforms?

Get from THREE's objects  
(tricky)

# The Geometric Properties

---

$\hat{n}$  - unit surface normal

$\hat{l}$  - unit vector to light source

$c_{\text{light}}$  - color/intensity of light

$c_{\text{ambient}}$  - color of ambient light

$c_d$  - color of material

**What coordinate system?**

Transform vertices

Interpolate normal vectors

# The Vertex Shader

---

```
attribute vec3 position;           // these are provided by THREE
attribute vec3 normal;            // you don't need to declare them
uniform mat3 normalMatrix;
uniform mat4 modelViewMatrix;
uniform mat4 projectionMatrix;

varying vec3 fNormal;
varying vec3 fPosition;

void main()
{
    gl_Position = projectionMatrix * modelViewMatrix * vec4(position, 1.0);

    fNormal = normalize(normalMatrix * normal);
    fPosition = modelViewMatrix * vec4(position, 1.0);
}
```

# Transformations of Normals

---

Transforming an object changes the normals

Normals are transformed by the inverse transpose (adjoint)

# What coordinate system?

---

World coordinate system?

Camera coordinate system?

Typically: lighting in camera coordinates

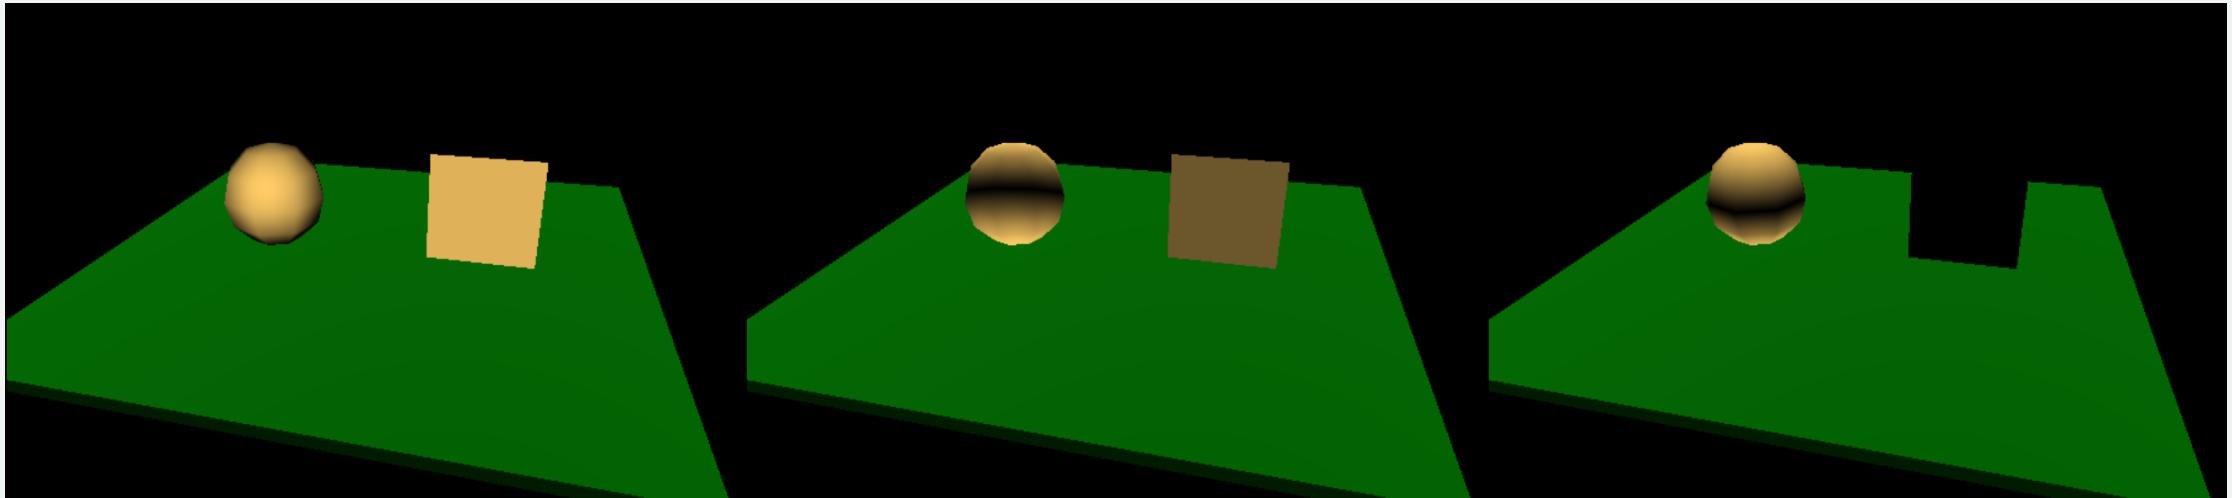
# Does it make a difference?

---

**Z-Axis Camera**

**Y-Axis Camera**

**Y-Axis World**



# What does THREE do?

---

from the documentation...

```
// = inverse transpose of modelViewMatrix  
uniform mat3 normalMatrix;
```

THREE's lights are transformed into view coordinates

# Light Geometry: Directional Light

---

```
const vec3 lhat = vec3(0,1,0);
```

Y Axis in same coordinate system as the normals (view)

# Light Geometry 2: Point Light

```
const vec3 lightPos = vec3(10,10,0);
varying vec3 fPosition;

main() {
    ...
    vec3 lhat = normalize(fPosition - lightPos);
```

but light position must be in camera coordinates

# Light Geometry2b: Point Light

```
uniform vec3 cameraPosition;  
const vec3 lightPos = vec3(10,10,0);  
varying vec3 fPosition;  
  
main() {  
    ...  
    vec3 lhat = normalize(wPosition - lightPos);
```

if we had wPosition as a varying as:

```
wPosition = modelMatrix * vec4(position, 1.0);
```

# Diffuse Lighting Shader...

```
varying vec3 v_normal;
const vec3 lightDir = vec3(0,0,1);
const vec3 baseColor = vec3(1,.8,.4);
const vec3 lightColor = vec3(.5,.5,.5);
const vec3 ambientColor = vec3(.1,.1,.2);

void main()
{
    /* we need to renormalize the normal since it was interpolated */
    vec3 nhat = normalize(v_normal);
    /* deal with two sided lighting */
    float light = abs(dot(nhat, lightDir));
    // or not...
    float light = clamp(dot(nhat, lightDir),0.0,1.0);

    /* brighten the base color */
    gl_FragColor = vec4((light*lightColor + ambientColor) * baseColor,1);
}
```

# Beyond Diffuse

---

Implement the equations of the model...

# What does THREE actually do?

---

1. `ShaderMaterial` - whatever we tell it to do
2. `MeshPhongMaterial` - basically the model we discussed  
see `lights_phong_pars_fragment.glsl`
3. `MeshStandardMaterial` - less clear what it does  
might be Phong with improved workflow
4. `MeshPhysicalMaterial` - a more complex model

# What if we want to use THREE's lights

---

We need to get them from THREE

- uniforms lib

We need to loop over however many lights there are

We need to do all the different kinds of lights

This is tricky - and not that instructive

Best bet: use the code from THREE's shaders

# Textures

---

1. Image-Based Textures
2. Procedural Textures

# Image-Based Textures in Shaders

---

Texture Maps are part of "Texture Objects"

- texture map (image)
- mip-map (images of different sizes)
- parameters (filter type, wrapping, ...)

GLSL Type `sampler2D` (for a 2D image)

# Getting the color from a texture

---

```
varying vec2 v_uv;  
  
// get the texture from the program  
uniform sampler2D texture;  
  
void main()  
{  
    gl_FragColor = texture2D(texture, v_uv);  
}
```

# **How does it know how to filter?**

---

Mip-Mapping set up in JavaScript (THREE's default)

## **Vertex Shader**

It doesn't know how to sample - uses the texture itself

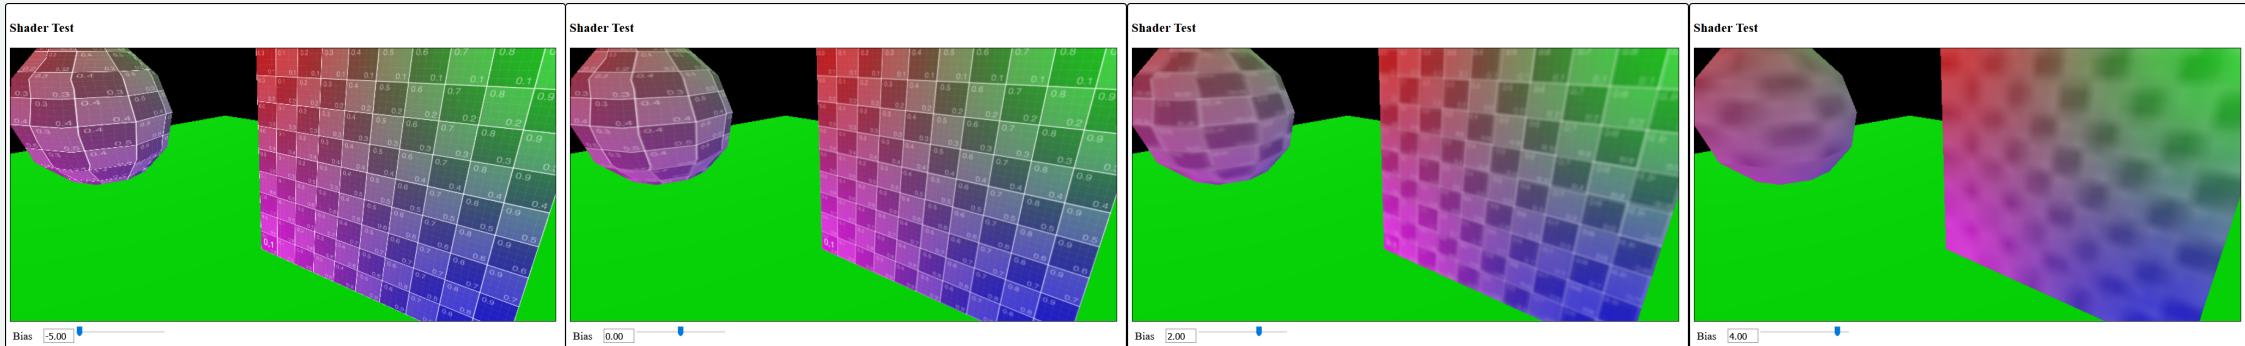
## **Fragment Shader**

It knows how far the next pixel is and uses that to estimate

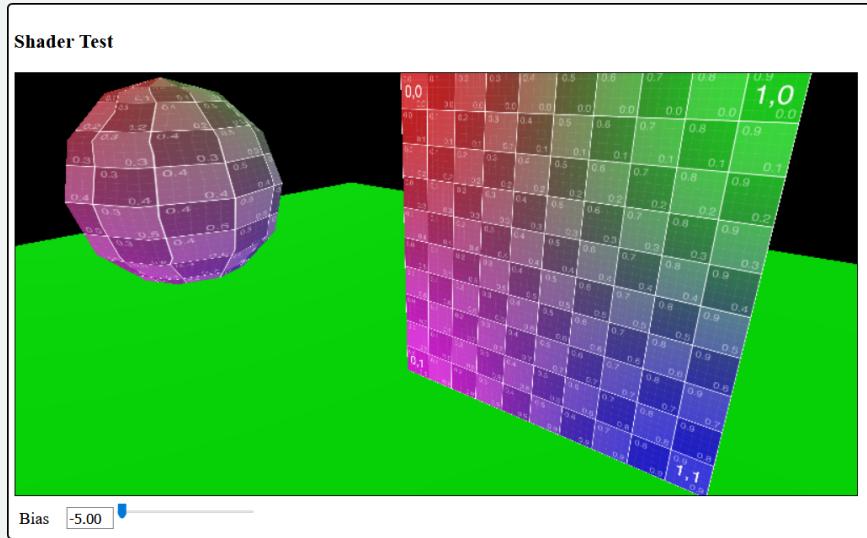
# Adjusting Filtering

we can **bias** the mip-map level (amount of blur)

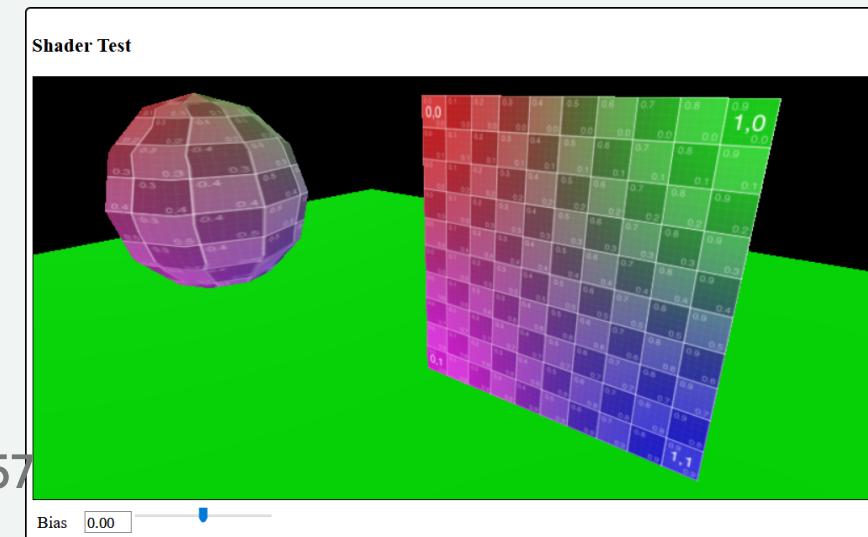
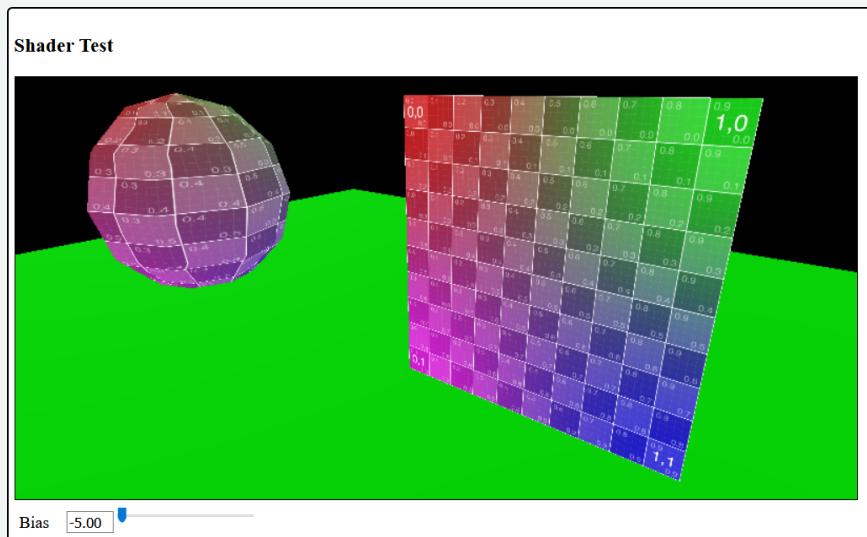
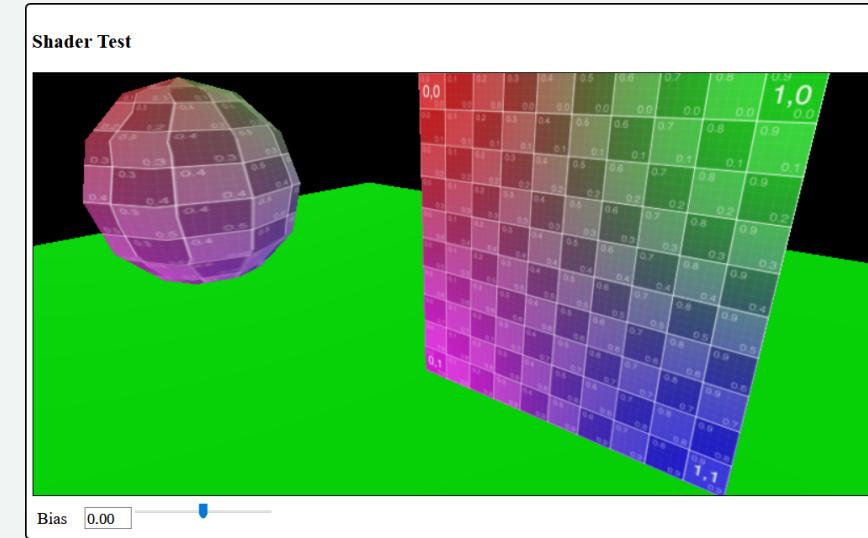
```
gl_FragColor = texture2D(texture, v_uv, bias);
```



# no filtering



# filtering



# Shader Techniques and Tricks

---

# Thinking About Shaders

---

The same little program run over different data

Historically, SIMD (single-instruction, multiple data)

Each processor runs in lock-step (same program counter)

```
if (x>.5) {  
    // something big  
} else {  
    // something small  
}
```

```
if (x>.5) {  
    // something big  
} else {  
    // something small  
}
```

```
if (x>.5) {  
    // something big  
} else {  
    // something small  
}
```

```
if (x>.5) {  
    // something big  
} else {  
    // something small  
}
```

All pixels execute both paths!

It doesn't work this way anymore - but it motivates how we think about shaders

# Use Built-Ins, not Conditionals

---

Functions for common tests:

- clamp
- step
- sign
- min
- max

Designed to be more efficient  
Convenient

# Use Built-Ins, not Expressions

---

- mix

# A Shader... Stripes

---



# Parameters - change as needed



Parameters for the number of stripes and the width  
Or change other things (colors)...

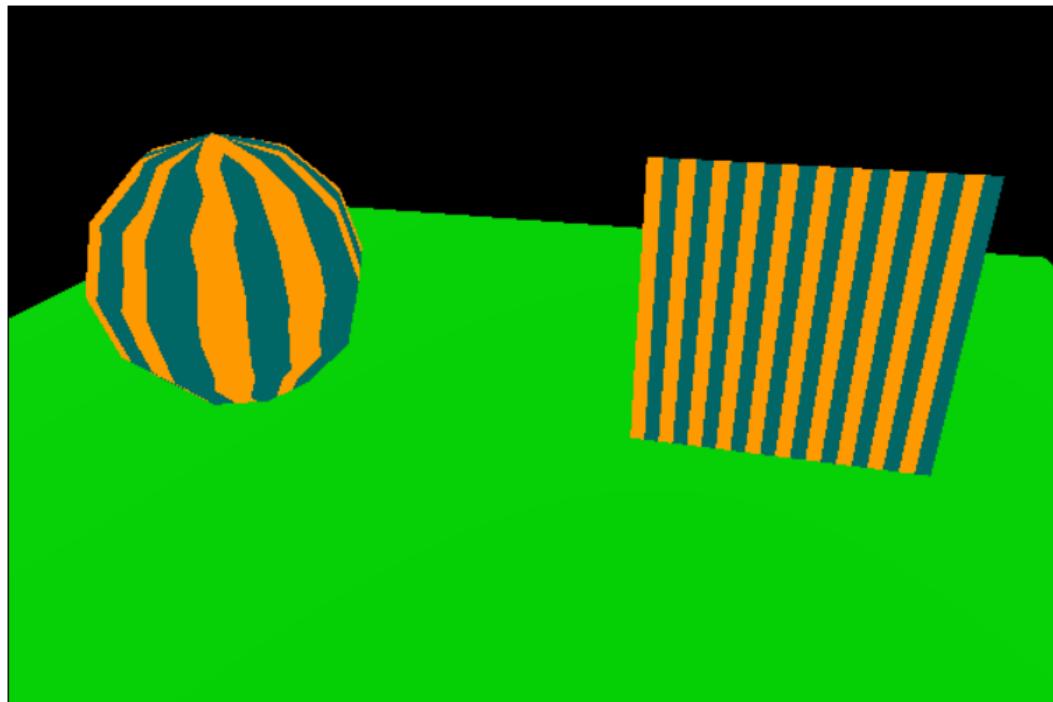
```
varying vec2 v_uv;

uniform vec3 color1;
uniform vec3 color2;
uniform float sw;
uniform float stripes;

void main()
{
    // broken into multiple lines to be easier to read
    float su = fract(v_uv.x * stripes);
    float st = step(sw,su);
    vec3 color = mix(color1, color2, st);
    gl_FragColor = vec4(color,1);
}
```

# Note the jaggies

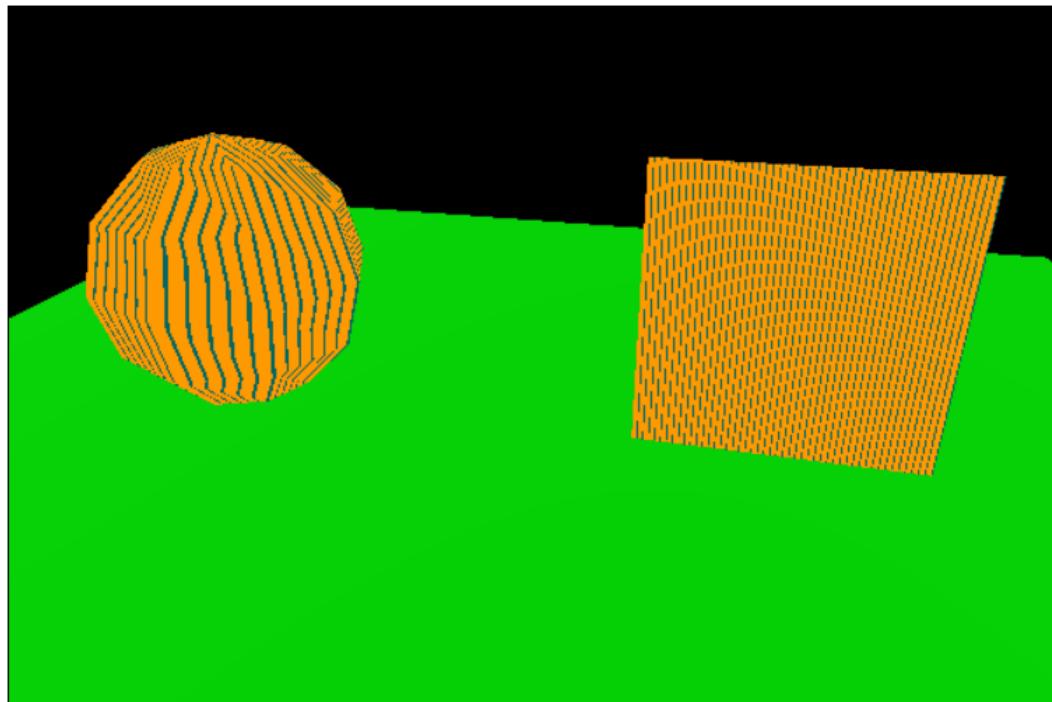
Shader Stripes



Stripes

Width

Shader Stripes



Stripes

Width

# Anti-Aliasing in Shaders

---

The problem - we are measuring at a specific place for each fragment

# Intuition: Sharp Edges are bad

Sharp Edge:

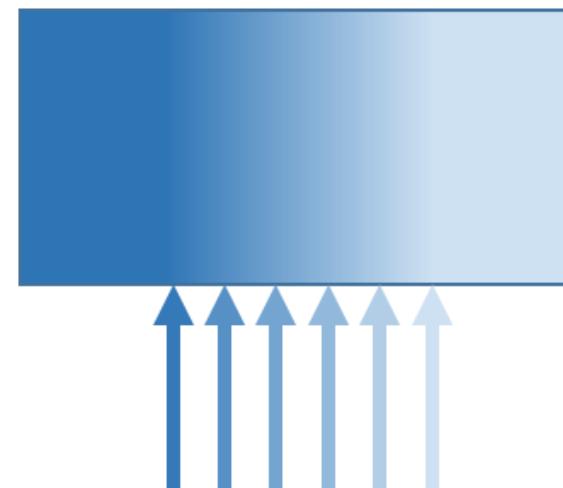
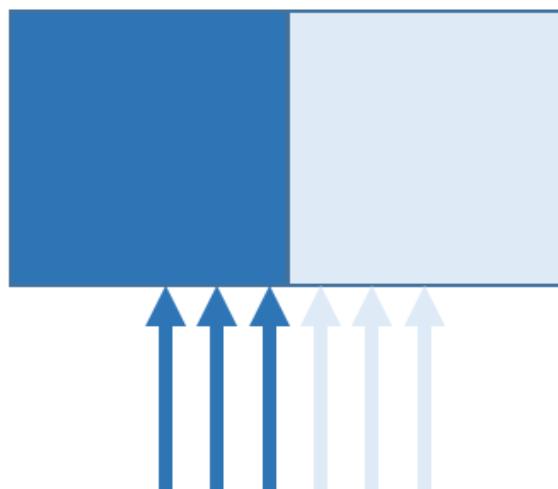
Small change in position

Big change in value

Smoother “Edge:”

Small change in position

Doesn't matter (that much)



# Where do sharp edges come from?

---

# Step vs. SmoothStep

---

## step

```
float st = step(sw,su);
```

## smoothstep

```
uniform float blur;  
float st = smoothstep(sw-blur,sw+blur,su);
```

# Warning - this is one sided!

---

Step is only for the 0-1 transition

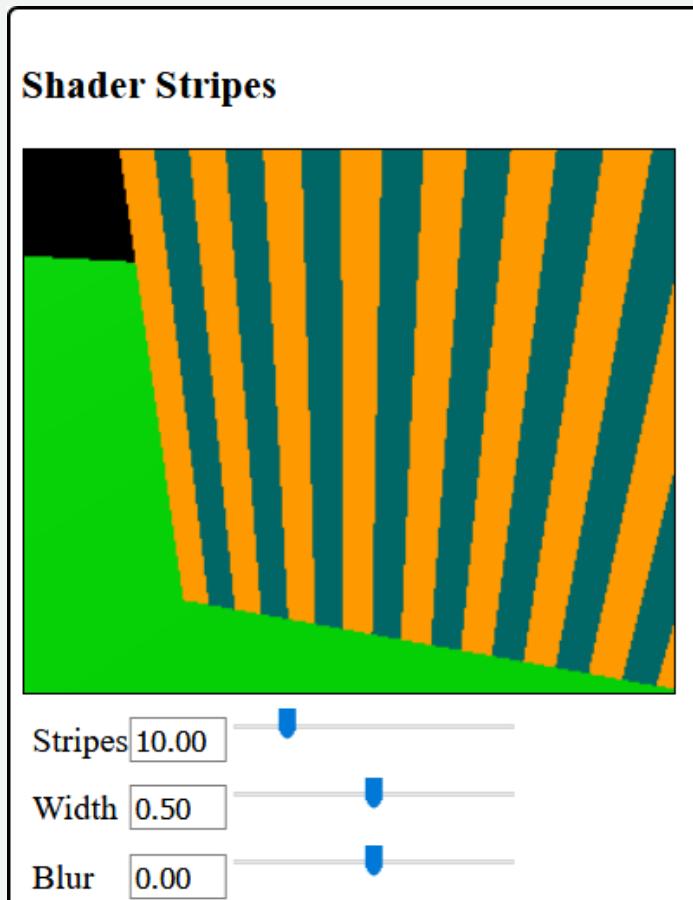
The 1-0 transition happens at the repeat

Need to deal with it separately

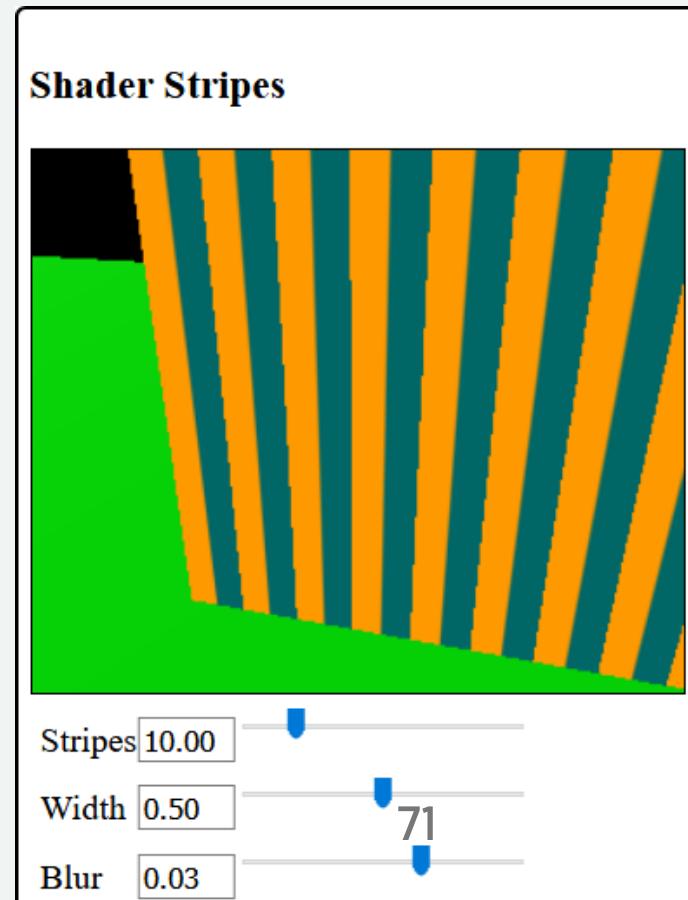
# Does this help?

---

**Width 0 (no blur)**



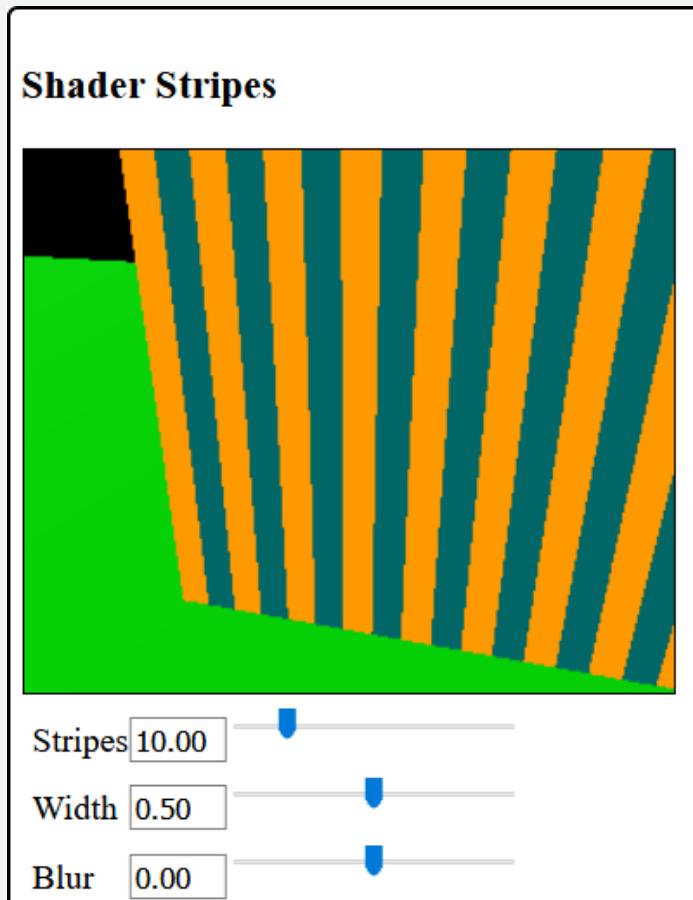
**Width 3 (blur)**



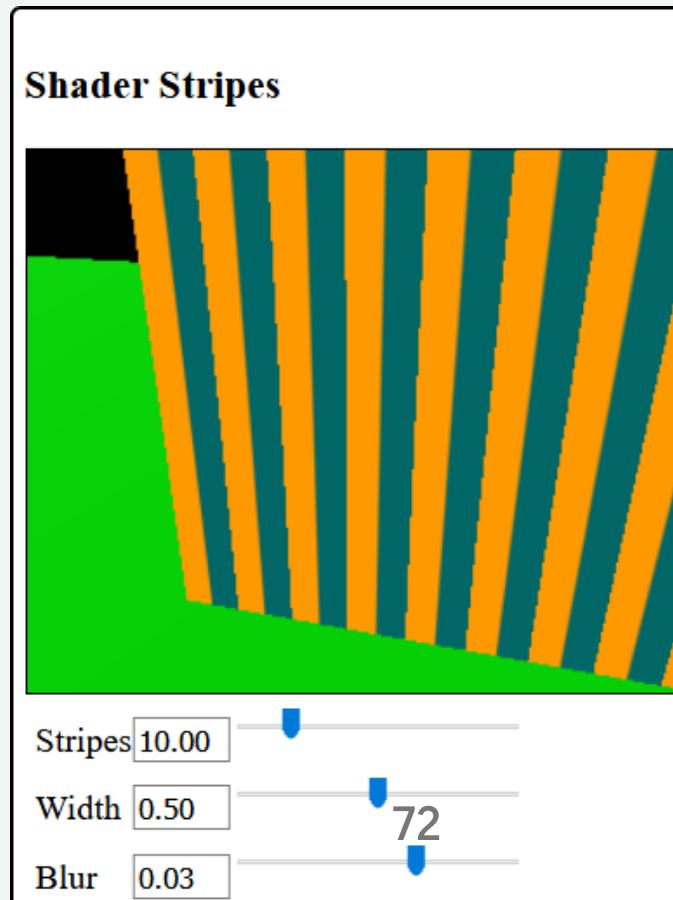
# Too much of a good thing?

---

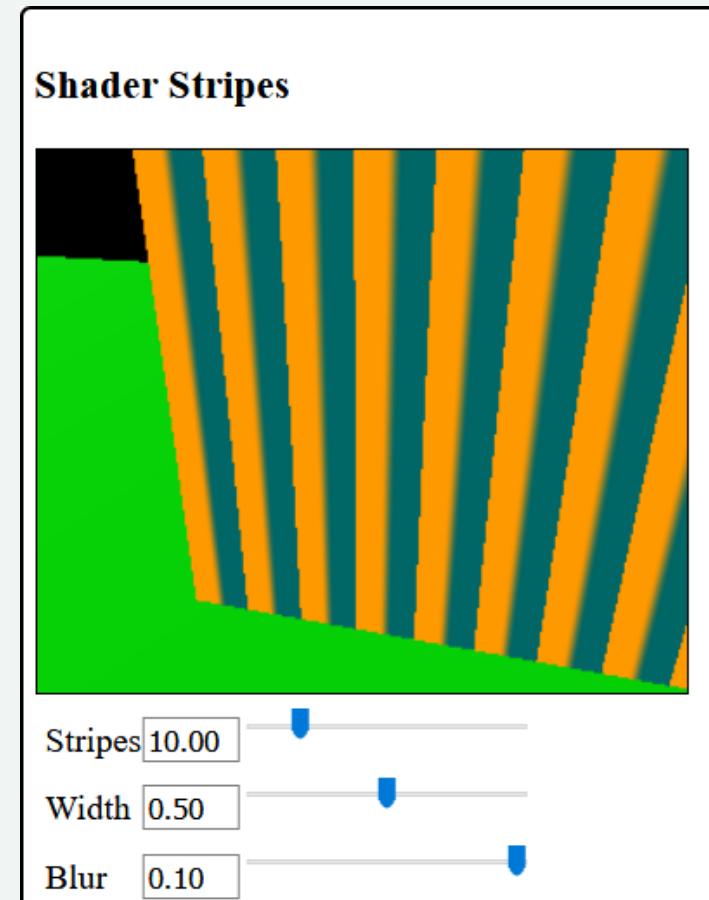
**Width 0 (no blur)**



**Width 3 (blur)**



**Width 10 (blurry)**



# What value for the width?

---

Want the blur to be "one pixel wide"

But what is that in u values?

GLSL will figure it out for us!

- `dFdx` - derivative of function with respect to x
- `fwidth` - derivative of function with respect to x and y

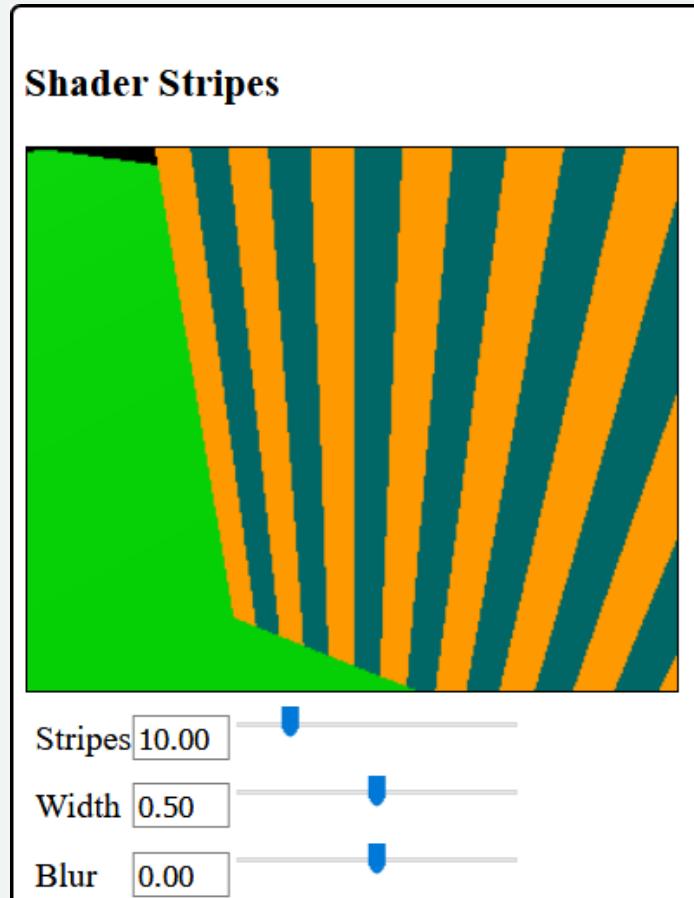
These are **extensions** to GLSL-ES (but three loads them for us)

- we need to enable them in the shader

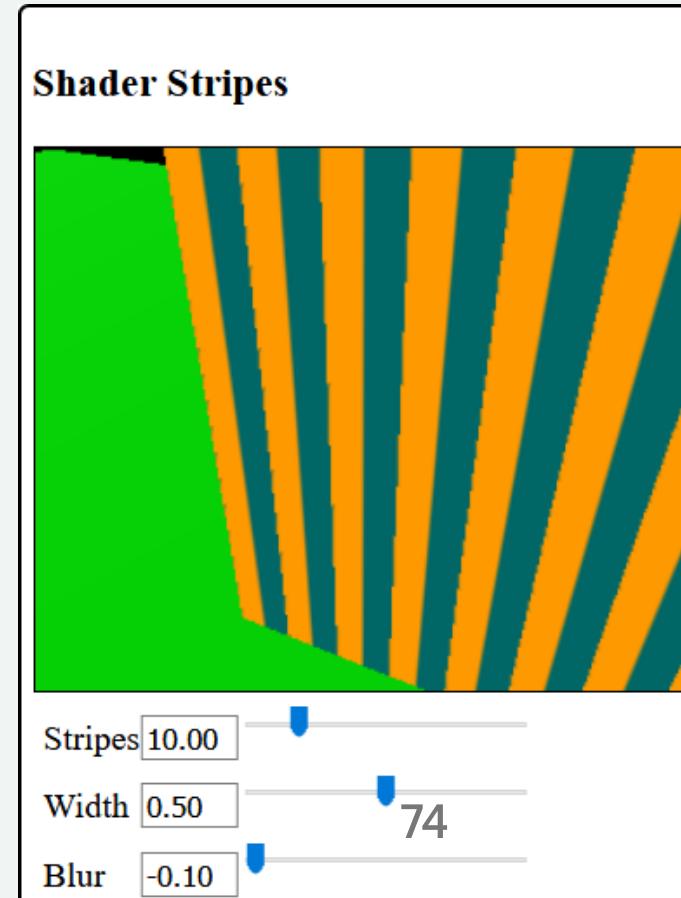
# fwidth: just right!

---

No Blur



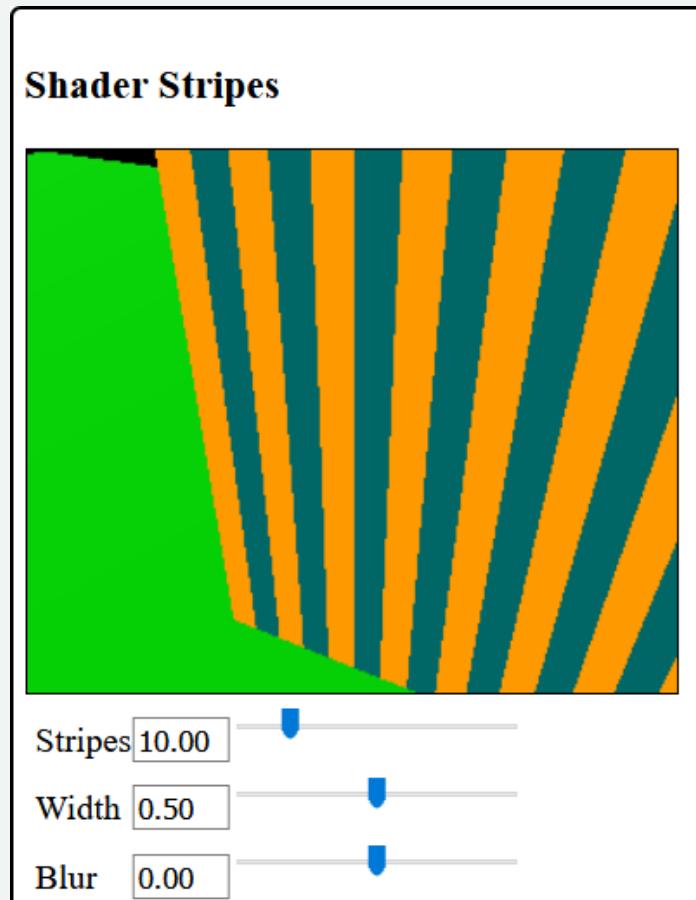
FWidth Blur



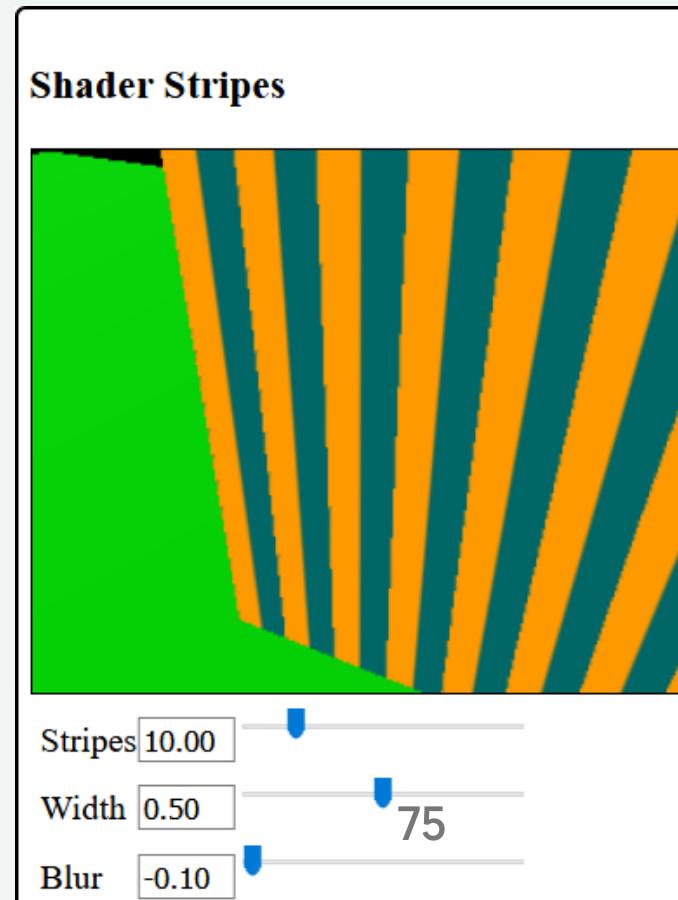
# fwidth: just right!

---

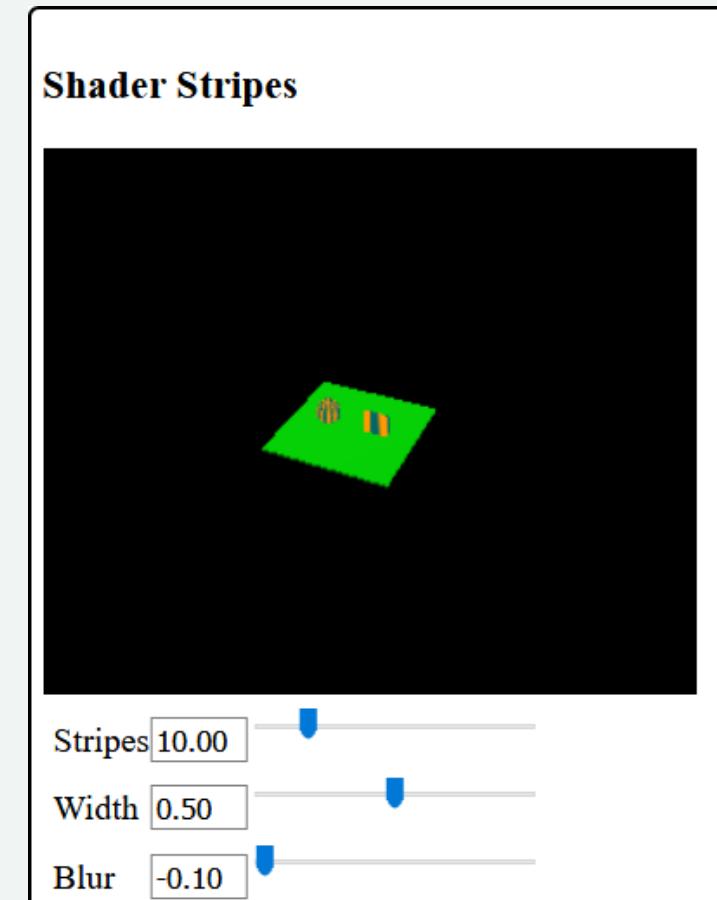
No Blur



FWidth Blur



Zoom out?



# **Smoothstep handles texture Magnification**

## **Not texture Minification**

We still need some way to do filtering

Hard to do in a general way for procedures

- over-sampling

Much easier to do for images